

Future-Proof Your Family

Future-Proof Your Family

*How to Protect Your Job, Your Income and the People You
Love in the Age of AI*

TODD LEGO

Future-Proof Your Family: How to Protect Your Job, Your Income and the People You Love in the Age of AI

Copyright © 2026 Todd Lego
All rights reserved.

No part of this book may be reproduced, stored or transmitted in any form or by any means without the prior written permission of the author, except for brief quotations in reviews and articles.

This book provides general educational information only. It is not financial, legal, medical or psychological advice, and it does not take your personal circumstances into account. For decisions about money, law, health or safety, speak with a qualified professional.

Artificial intelligence changes quickly. The information in this book was checked against the sources listed in the Notes as at October 2026. Some figures will have changed by the time you read it, so verify anything important before you rely on it. Product, company and organisation names are used for identification only and do not imply endorsement.

First edition, 2026

Contents

Why I Wrote This	vii
How to Read This Book	ix
PART ONE · IT HAS ALREADY STARTED	
1 The Hiring Freeze Nobody Announced	1
2 What AI Actually Is.....	8
3 The Doubling Clock	14
PART TWO · THE NEXT FIVE YEARS	
4 2027–2028: The Agent Years.....	20
5 2029–2031: Three Futures.....	26
PART THREE · YOUR JOB	
6 The Task Test	32
7 Office, Admin, Finance, Sales and Support	41
8 Health, Education, Law and the Licensed Professions.....	47
9 Trades, Transport, Factories and Warehouses	55
10 Creative and Tech Work	62
11 If You Do X, Start Learning Y	68
PART FOUR · MONEY	
12 The Income Tripod	74
13 Earning With AI: From \$0, \$100 and \$1,000.....	81
PART FIVE · BECOMING AI-NATIVE	
14 How to Actually Use AI.....	88
15 The 30-Day AI Reskill Challenge	95
PART SIX · PROTECTING YOUR FAMILY	
16 Scams, Voices and Fakes.....	100
17 Love, Loneliness and Machines	107
18 The Family AI Preparedness Plan	113
PART SEVEN · CHILDREN AND EDUCATION	
19 What Our Kids Should Learn Now	118
PART EIGHT · THE DARKER SIDE	
20 The Risks, Straight.....	126
PART NINE · THE OPPORTUNITY	
21 The Human Premium.....	134

22 Direct It 138

PART TEN · YOUR FIVE-YEAR PLAN

23 Your Plan 143

24 The AI Resilience Scorecard and a Last Word 151

Glossary 158

Resources 161

Notes 164

About the Author 174

Why I Wrote This

Two years ago I knew bugger-all about artificial intelligence.

I'd seen the headlines like everyone else. I'd heard people say it was going to change everything, and I'd heard just as many say it was a fad, a bubble, a toy for tech blokes in California. I didn't know who was right. Honestly, for a while I didn't think it had much to do with me.

But something in my bones told me this one was different. It didn't feel like 3D televisions or the latest gadget that everybody forgets about in eighteen months. It felt more like the early days of the internet, when most people couldn't see what the fuss was about and a few people quietly worked out that the world was about to be rearranged around it.

I figured I had two choices. I could bury my head in the sand, tell myself it was somebody else's problem and hope the world didn't change too much. Or I could get off my backside and learn.

I chose to learn.

I'm not a computer scientist. I'm not a programmer by training, and I'm not an academic. I came at this the way most of you will: as an ordinary person trying to work out what this stuff actually is and what it means for real life. For the past two years I've put a huge amount of my time into it. Reading the research. Using the tools every day. Building things with them, breaking things with them, getting things badly wrong and slowly getting better. Learning how computers and these new AI systems actually work, at least well enough to explain them to someone else without hiding behind jargon.

The deeper I went, the more one thing bothered me. The people building this technology talk about it constantly. The people investing in it talk about it constantly. But the people it's going to affect most, ordinary workers, parents, small business owners, kids leaving school, older people living on their own, mostly aren't part of the conversation. A lot of them are

unaware of how fast it's moving. Some are frightened by it. Some have written it off. And a lot of them simply don't know where to start.

I don't think that's good enough. People deserve a plain-English explanation of what's happening, what isn't, and what they can actually do about it. The more I learned, the more I felt an obligation to pass it on.

So this is the book I wish ordinary families already had.

I want to be upfront about a few things.

I'm not going to tell you I can predict the future. Nobody can, and anyone who claims to know exactly what 2031 looks like is selling something. What I can do is show you the best evidence available, tell you honestly how solid it is, and help you make good decisions anyway. You don't need to predict the future perfectly to prepare for it sensibly.

I'm also not going to scare you for the sake of it. There's enough of that around. Some of what's in these pages is genuinely confronting, and I won't water it down. But fear on its own is useless. Fear without a plan just keeps you up at night. Every chapter in this book ends with something you can actually do.

And I'm not going to sell you anything. There's no secret course at the back of this book and no "AI money machine". Just the most useful things I've learned, written down as clearly as I can manage.

If this book does its job, you'll finish it understanding what's really going on, knowing where you and your family stand, and holding a practical plan for the next five years. You'll know how to protect your job, how to protect your income, and how to protect the people you love from some new and nasty risks. And I hope you'll see the opportunities too, because there are a lot of them.

You can't control how fast AI moves. You can control how ready you are when it arrives.

Let's get you ready.

Todd Lego

How to Read This Book

Almost everything you've heard about artificial intelligence has been told to you by someone who wanted something. Tech companies want your subscription. Influencers want your clicks. The doom merchants want your fear, because fear keeps you watching. So most people end up stuck in the middle: half-convinced it's all hype, half-worried it isn't, and unsure which half to trust.

I can't fix that by telling you to calm down, and I won't fix it by telling you to panic. What I can do is show you what the evidence says and tell you plainly how strong it is.

So wherever it matters, especially when we're talking about the future, I label things the way a decent weather forecaster would.

HAPPENING NOW. It's documented and it's real. The Notes at the back tell you where the evidence comes from.

TRAJECTORY. Strong evidence points this way if current trends continue. Not certain, but worth planning around.

PLAUSIBLE SCENARIO. Credible people argue for it, but the evidence is thin or the experts genuinely disagree. Worth understanding, not worth losing sleep over.

WILD CARD. It could happen. Nobody can reliably predict whether or when. I mention these so you aren't blindsided, then move on.

When you see a number in this book, it came from somewhere real, and the Notes tell you where. A lot of the best data on AI and jobs, and on scams, comes from the United States, because that's where the most detailed figures are published. Where I've found good Australian or international figures, I've used them too. Dollar amounts are US dollars unless I say otherwise.

When I'm guessing, I'll say so. When the experts are arguing, I'll show you both sides of the argument.

You don't have to read the book in order. It's built so you can jump to the part that's keeping you up at night. But I'd ask everyone to read three chapters whatever else they skip: **Chapter 1**, because it shows what's already happening quietly; **Chapter 6**, because it gives you an honest way to look at your own job in about twenty minutes; and **Chapter 16**, because it could stop someone you love from being robbed.

Most chapters finish with a short section called **Do this now**. A few small actions, most of them under twenty minutes. If all you did with this book was work through those, you'd be further ahead than most people you know. That isn't a sales line. Most people aren't doing anything at all.

CHAPTER 1

The Hiring Freeze Nobody Announced

Nobody got fired.

That's the part people miss.

When we imagine machines taking jobs, we picture something loud. A factory floor going quiet. A boss with a cardboard box. A news crew outside the gates and a politician promising to fix it all by Christmas.

What's actually happening is quieter than that, and in some ways worse, because you can't see it.

Someone in accounts leaves to have a baby and nobody replaces her. Her work gets split between two colleagues and "the new system". A law firm that used to take on eight graduates takes three. A marketing agency stops advertising junior roles altogether. No announcement, no redundancies. A door that used to be open just quietly shuts.

You'd never notice from the inside. The people who still have jobs keep them. The office looks the same and the Christmas party still happens. The only people who notice are the ones standing outside the door, knocking, wondering why nobody answers.

WHAT THE PAY SLIPS SHOW

In August 2026 a team of economists at Stanford University, Erik Brynjolfsson, Bharat Chandar and Ruyu Chen, published an update to a study they'd first released a year earlier. They'd been given access to payroll records from ADP, one of the largest payroll processors in the United States. Not surveys

and not opinions. Real pay slips covering millions of workers: who was being paid, for what job, at what age.

They were looking for an answer to one question. Is AI already changing who gets work?

The first thing they found was that there's no economy-wide jobs collapse. Most people's work looked the way it always had. If you've been told that AI has already wiped out millions of jobs, that isn't what the pay slips say.

The second thing was that one group is falling behind. Young workers aged 22 to 25, in the jobs most exposed to AI, such as software development and customer service, were employed at levels about 19 per cent below where you'd expect if they'd kept pace with young people in less-exposed jobs. A year earlier the gap had been 15 per cent. It was growing.

Older, experienced workers in those same jobs showed no comparable gap. And it was happening through hiring, not firing. People weren't being pushed out. New people just weren't being let in. Wages hadn't moved much either; the change was in how many people were employed, not what they were paid.

The researchers were careful. They called these patterns, not proof. They couldn't say for certain that AI caused them, some of the shift started before AI tools were widespread, and the gap shrinks when you account for things like education. I respect that caution and I'll hold onto it throughout this book.

But notice *how* the change showed up. Not as a wave of redundancies. As a gap. As jobs that were never advertised.

That's what change looks like before it's obvious.

WHY TWO STORIES CAN BOTH BE TRUE

You'll hear two stories about AI and jobs, and they sound as though they can't both be right. One says it's a bloodbath and the entry-level office job is dying. The other says relax, the data shows nothing much is happening.

Both are pointing at real evidence.

The Budget Lab at Yale University has been tracking the whole American labour market since ChatGPT was released in late 2022. Its conclusion,

updated repeatedly through 2026, is that the overall mix of jobs in the economy hasn't shifted in any unusual way, and that whatever softness there's been in the job market is "probably not (yet)" AI. Pay attention to that little word in brackets.

In August 2026 the US unemployment rate sat at 4.1 per cent. That's not a crisis number. Your neighbour who drives a delivery van, your cousin who's a nurse, your mate who fits kitchens: their work looks roughly the same as it did three years ago.

Look closer at recent university graduates, though, and something odd appears. The Federal Reserve Bank of New York tracks graduates aged 22 to 27. Through the first half of 2026 their unemployment rate sat at around 5.6 to 5.7 per cent, higher than the rate for workers overall. Historically a degree has made you *less* likely to be out of work than everyone else, not more. And about 42 per cent of recent graduates were underemployed, working in jobs that don't need their degree.

There are plenty of possible reasons. A cooling economy hits new entrants first. Some degrees don't match what employers want. AI is only one suspect in the room. But it's in the room.

So here's how I'd put the two stories together. The big picture is stable. The bottom rung of the ladder is wobbling.

And the bottom rung matters more than people think. It's where you learn. It's where a nineteen-year-old turns into a twenty-nine-year-old who knows what they're doing. If fewer people get on the ladder, you don't feel it this year. You feel it in ten years, when there aren't enough experienced people to go around.

THE COMPANIES ARE STARTING TO SAY IT OUT LOUD

For most of 2023 and 2024, companies were careful not to say they were cutting jobs because of AI. It sounded heartless and it was bad PR.

That changed. Challenger, Gray & Christmas, an American outplacement firm, has tracked announced job cuts for decades. Every month it records not just how many jobs companies say they're cutting but *why* they say they're cutting them. In 2026, artificial intelligence became the most commonly cited reason, month after month. By the end of June, AI had

been cited in more than 100,000 announced cuts that year, close to double the total for the whole of 2025.

I need to be straight with you about what that number means, because this is exactly where scary content goes wrong. Those are *announced reasons*. Companies choose what to say. “We’re investing in AI efficiency” sounds a lot better to shareholders than “sales are down and we over-hired in 2021”. Some of those cuts would have happened anyway. The honest reading is that AI has become a reason companies are comfortable giving. Some of the time it’s the real reason. Some of the time it’s a convenient story.

Either way, it tells you where managers’ heads are at. In boardrooms the conversation has shifted from “should we use AI?” to “how many people do we need now that we have it?” That conversation is happening about your team whether you’re in the room or not.

WHY THE QUIET VERSION IS MORE DANGEROUS

A loud disruption gives you warning. A factory closure is announced months ahead. There’s news coverage and retraining programs, however inadequate, and people rally around each other. A quiet disruption gives you nothing.

Picture a 38-year-old claims processor at an insurance company. Eleven years in the job, good at it, never been in trouble. Over three years, four people on her team of twelve retire or move on, and none of them is replaced. The work still gets done, because the company brought in software that handles the simple claims automatically. The team is now eight people, everyone’s a bit busier, and there’s a new dashboard on the wall.

Nobody ever tells her that her job is at risk. Not this year, anyway. But she’s now one of eight people doing work that used to take twelve. At the next restructure the question won’t be whether to cut the team. It’ll be whether they need eight or whether five would do. And when she goes looking for another claims job, because her manager’s a nightmare or she wants to move closer to her mum, she’ll find the whole industry has been quietly doing the same maths. There are fewer claims jobs everywhere.

She’s a composite, by the way, built from the patterns in the data rather than from one real person. When I use a real person or company in this

book I'll name them, and the Notes will tell you where the information comes from. When I'm painting a picture to help you see something, I'll tell you that too.

That's the pattern this book is trying to get you ahead of. Not a pink slip. A slowly shrinking pond, noticed only when you try to swim somewhere else.

“BUT IT’S STILL MOSTLY HYPE, ISN’T IT?”

Some of it is. Absolutely.

In 2025, researchers at MIT looked at how companies were getting on with generative AI and reported that about 95 per cent of corporate AI pilot projects weren't producing any measurable profit. The technology was being bought, trialled and talked about in meetings, then sitting there not changing the bottom line. That study copped a lot of criticism for its methods, which leaned on a fairly small number of interviews and surveys. But the pattern rang true to anyone who has worked in a big organisation. Buying the tool is easy. Changing how people actually work is hard.

So yes, plenty of companies are wasting money on AI and plenty of promises are overblown. If you're a sceptic, you've got good evidence on your side.

But hold this alongside it: the hype being overblown doesn't mean the change isn't real.

The internet was wildly overhyped in 1999. Hundreds of companies went bust in 2000 and 2001, and the people who'd called it a bubble felt very clever. They were right about the bubble. They were wrong about the internet. Within fifteen years it had changed how almost everyone shops, banks, dates, works, finds a plumber and argues with relatives. The hype came first and fell over. The real change came second and stayed.

I think there's a decent chance we're watching the same film again. You don't have to believe the hype to prepare for the change.

WHAT THIS MEANS FOR YOU

If you're young and trying to land your first proper job, you're not imagining it. It's harder than it was for people a few years older than you, and part of the reason is that the entry-level tasks you'd once have been hired to do are the ones AI is best at. This book will show you how to get in by a different door.

If you're mid-career and experienced, the data says you're mostly fine for now. Your experience is your armour. But armour works best when you maintain it, and the people most secure in five years will be experienced workers who also know how to direct AI. That's a very learnable combination.

If you're a parent, the path you took (school, maybe uni, an entry-level job, working your way up) might not be the path your kids get. Not because it's gone, but because the first few rungs are changing shape. Part Seven deals with that.

If you run a business, you're probably already feeling the pull to hire fewer people and use more software. Maybe that's right. But companies that cut too deep, too fast have sometimes had to walk it back, and Chapter 7 looks at a famous example.

And if you work with your hands, the next few years look quite different for you, in some ways a lot better. Hold that thought until Chapter 9.

THE MOST IMPORTANT SENTENCE IN THIS CHAPTER

By the time something is obvious to everyone, the cheapest and easiest moment to prepare for it has already passed.

If the news ever fills up with stories about AI unemployment, every course will be full, every career coach will be booked, every job asking for AI skills will have five hundred applicants, and the people who started learning two years earlier will be the ones doing the hiring.

That moment isn't now. Now is earlier than that. You've got more time than you fear, but less than you'd like.

DO THIS NOW

1. Write down three things you actually did at work this week. Not your job title, the tasks. “Answered customer emails about late orders.” “Wrote a quote.” “Checked invoices against orders.”
2. Open a free AI assistant (ChatGPT, Claude, Gemini and Copilot all have free versions) and ask it to help with one of those tasks. Give it real detail, but leave out anyone’s private information.
3. Notice how you feel about the result. Relieved it was rubbish? Uneasy that it was good? Write one sentence about it. That feeling is useful information and we’ll come back to it.
4. Ask a young person you know how their job hunt is going, and really listen.

CHAPTER 2

What AI Actually Is

You don't need to understand how an engine works to drive a car. But you do need to know the brake from the accelerator, roughly how fast you're going, and what that warning light on the dashboard means.

That's the level we're aiming for here. Not engineering. Driving.

THE INTERN WHO'S READ EVERYTHING

The best way I've found to describe a modern AI assistant is this. Imagine the most well-read intern in history. They've read a huge chunk of everything ever written and put online: books, websites, manuals, forums, computer code, recipes, court cases and a few billion arguments. They can write in any style, explain almost anything, and they never get tired or grumpy.

But they've never been outside. They've never touched anything, never been responsible for anything, and never had to live with the consequences of being wrong.

That intern is incredibly useful. They'll draft your email in ten seconds, summarise a forty-page contract, explain your electricity bill, suggest six names for your business, or help your kid with fractions at nine o'clock at night when you've forgotten how.

That intern is also, sometimes, confidently and completely wrong. And they'll say the wrong thing in exactly the same calm, polished voice they use when they're right.

Hold both of those in your head at once and you understand AI better than a lot of people who talk about it on television.

HOW IT WORKS, IN TWO MINUTES

The AI tools most people use, such as ChatGPT, Claude, Gemini and Copilot, are built on something called a *large language model*. Here's what that means in plain English.

Engineers took an enormous amount of text and had a computer program study it over and over, getting better and better at one simple-sounding game: given everything written so far, what comes next?

Do that trillions of times on a big enough computer and something surprising happens. To get really good at predicting the next line of a physics textbook, the program has to learn something that looks a lot like physics. To predict the next step in a legal argument, it has to pick up something that works like legal reasoning. To predict the next line of computer code, it has to learn to code. Nobody sat down and taught it those things one at a time. They emerged from the game.

After that, the companies spend a great deal of effort training it to be helpful and safe: to answer the question rather than ramble, to refuse to help with dangerous things, and to admit when it isn't sure, which it still doesn't do nearly often enough.

That's the heart of it. It isn't a person. It isn't a search engine looking things up in a database, although many tools can now search the web as well. It doesn't "know" things the way you know your own phone number. It's more like an extraordinarily practised pattern-finisher, one that has become so good at finishing patterns that much of the time it's hard to tell the difference between that and understanding.

Whether there's any real understanding in there is something philosophers and scientists genuinely argue about. For your purposes it doesn't matter much. What matters is what it can do, what it can't, and how often it gets things wrong.

WHY IT MAKES THINGS UP

You'll hear the word *hallucination* a lot. It means the AI has produced something that sounds right but isn't: a fake quote, a book that doesn't exist, a court case that never happened, a statistic pulled out of thin air.

Go back to the game and you'll see why. Its deepest instinct is to produce text that *sounds like* what should come next. Most of the time, what sounds right is right. But when it doesn't know, it doesn't always go quiet. Sometimes it fills the gap with something plausible.

This has already embarrassed real professionals. In 2023 a New York judge fined two lawyers and their firm US\$5,000 after they filed a brief citing cases that ChatGPT had invented. In 2024 a Melbourne solicitor was referred to the Victorian legal regulator after handing a family court judge a list of authorities, produced by an AI research tool, that didn't exist. Neither of them set out to mislead anyone. They trusted the intern and didn't check.

Newer models make things up less often than older ones, and many can now check their answers against the web. But less isn't never. The golden rule, which you'll see again in Chapter 14, is simple: use AI to think faster, never to stop thinking. If the answer matters (money, health, the law, your kid's safety, anything you'll put your name to), check it.

FROM CHATBOT TO ROBOT: THE FOUR STEPS

This is the part most explanations skip, and it's the part that matters most for your job. AI isn't one thing. It's a ladder, and it's climbing. Think of a kitchen.

The recipe book. You ask a question and it gives an answer. "How do I make a white sauce?" It tells you; you do the cooking. This is a *chatbot*, and it's what most people think AI is, because it's what they tried in 2023.

The kitchen assistant. Now it can look at what you show it and work with your stuff. You photograph the inside of your fridge and it suggests dinner. You give it your shopping receipts and it works out your weekly food budget. It helps; you still do the doing. This is an *AI assistant*.

The cook. This is the big one. An *agent* is an AI that doesn't just answer but acts. You say, "Sort out dinner for Saturday, six people, one vegetarian, eighty bucks," and it finds recipes, builds a shopping list, places the order online, books the delivery and sends you a cooking timeline. It takes a goal, breaks it into steps and carries them out, using a computer the way a person would: clicking, typing, searching, filling in forms, checking its own work. In offices, agents are starting to do things like work through two hundred invoices, match them against purchase orders, flag the ones that don't add up and draft emails to the suppliers. Agents are where most of the change over the next five years is likely to come from, because they move AI from helping a person do a task to doing the task.

The cook with hands. Everything so far happens on a screen. The last step is putting that intelligence into a body that can move through the physical world: picking things up, walking through a warehouse, driving a car, folding washing. This step is real but much further behind. Driverless taxis are carrying hundreds of thousands of paying passengers a week in American cities (more on that in Chapter 9). But the humanoid robots you've seen in slick videos are mostly still in trials, and when Stanford University's 2026 AI Index looked at how robots performed on household tasks, they succeeded only about 12 per cent of the time.

The physical world is hard. Your hands are an astonishing piece of engineering. Keep that in mind, because it turns out to be one of the most important facts in this book.

HAPPENING NOW The first two steps are everyday reality. Agents are in real use and improving quickly. Robots are real in narrow areas such as driving and some factory work, but general-purpose robots are still early.

IT CAN SEE, HEAR AND TALK

You'll also hear the word *multimodal*. It just means the AI can work with more than text. The leading systems can now look at a photo of a rash, a car engine, a spreadsheet or a building site and talk about it. They can listen to a meeting and write up the notes. They can reply in a natural voice, in real time, with pauses and tone. And they can create images, video, music and voices.

That last one is where some of the biggest risks come from. If an AI can produce a realistic voice from a short recording, it can produce yours, or your daughter's. Chapter 16 deals with that properly, and you'll come out of it with a simple defence that works.

WHAT AI IS NOT

A few myths are worth clearing up, because they get in the way of good decisions.

It's not a person. It can sound caring, and it can be genuinely helpful when you're struggling, but it doesn't love you, miss you or worry about you after you close the app. That matters a great deal when we talk about AI companions in Chapter 17.

It's not all-knowing. It has blind spots, gaps and biases picked up from whatever it learned from. It can be out of date, and it can be wrong about local laws, local prices and local people.

It's not magic, and it's not nothing. The two most dangerous attitudes are "it can do anything" and "it's just autocomplete". Both lead to bad decisions. The truth sits somewhere uncomfortable in between: astonishingly capable at some things, surprisingly useless at others, and the line between the two keeps moving.

And it isn't one company's product. Several big companies, plenty of smaller ones and several countries are competing hard. The 2026 AI Index found the best American and Chinese models were within a few percentage points of each other on key tests in early 2026. That competition is one reason things are moving so fast.

WHY THIS MATTERS TO YOU

You don't need to know how any of this works under the bonnet. What you need is a working sense of what the intern can do, so you can answer three questions about your own life. Which parts of my work could the intern do? Which parts could the intern help me do faster? And which parts need someone who has been outside, touched things and can be held responsible?

That third category is going to become the most valuable thing you own. We'll build it out properly in Chapter 6.

DO THIS NOW

1. Have your first real conversation with an AI about your actual work. Try: "I work as a [your job]. Here's what a normal day looks like for me: [four or five sentences]. Which parts of my day could you help with, and which couldn't you help with at all? Be honest."
2. Catch it being wrong. Ask about something you know well (your town, your trade, your sport, a health condition you live with) and push until you find a mistake. Remember how confident it sounded.
3. Show it a photo of something from your day, such as a form, a dripping tap, a label or a page of homework, and ask it to explain what it sees.

CHAPTER 3

The Doubling Clock

There's an old riddle about a pond.

A single lily pad appears on the water. Every day the number of lily pads doubles, and after thirty days the pond is completely covered. On what day was the pond half covered?

Most people's gut says day fifteen. The answer is day twenty-nine.

Here's the bit that matters. On day twenty-five the pond was only about 3 per cent covered. If you'd walked past on day twenty-five, you'd have seen a few lily pads in one corner and thought there was nothing much going on. Four days later the water's gone.

Our brains aren't built for doubling. We think in straight lines, adding a bit each year, because most of the world works that way. But some things don't add, they multiply, and when something multiplies the early part looks boring right up until it doesn't.

This chapter is about whether AI is one of those things.

A STOPWATCH FOR MACHINES

For a long time people measured AI progress with exams. Can it pass the bar exam? Can it solve maths olympiad problems? Can it score well on a medical licensing test?

Those tests matter, but they have two problems. They measure whether AI can *answer questions*, not whether it can *do work*. And they get used up: once AI aces a test, the test stops telling you anything.

In 2025 a small independent research group in California called METR came up with a different measuring stick. I think it's the single most useful number in this whole debate, so I'll explain it carefully.

They took a large set of real tasks, mostly software and technical work of the kind people get paid to do, and timed skilled humans doing them. Some tasks took a person two minutes. Some took an hour. Some took a full working day or more. Then they gave the same tasks to AI systems and asked: how long a task, measured in human time, can the AI complete successfully about half the time?

That's the *time horizon*. Think of it as the size of a job you could hand to the AI and walk away from, with a coin-flip chance that it's done when you get back.

When METR lined up AI systems from 2019 onwards, the result was startling. In 2019 the best systems could handle tasks that took a skilled human a few seconds. By 2023 it was a few minutes. By early 2025, around an hour. In METR's January 2026 update, the best model it had measured was managing tasks of about five hours, and a month later it estimated that a newer model could handle tasks of roughly fourteen and a half.

Plotted on a graph, that makes a remarkably steady line. Across the whole period since 2019, the length of task AI could handle doubled roughly every seven months. Since 2024 the doubling has been faster, closer to every three or four months.

By mid-2026 METR was saying something striking: its test was running out of road. The newest models were scoring near the top of what its task set could reliably measure, around sixteen hours of human work.

HAPPENING NOW On the leading independent measure, the length of task AI can complete on its own (with about a 50 per cent success rate) has been doubling every four to seven months for several years.

WHAT DOUBLING FEELS LIKE

Make that concrete. Say an AI can do a one-hour task today, and the doubling keeps going at seven months. In seven months it can do two-hour tasks. In fourteen months, four hours: half a working day. In twenty-one months, a full day. In twenty-eight months, two days. In thirty-five months, nearly a working week. Just under three years from "an hour" to "most of a week", and quicker than that at the faster pace.

I want to be careful here, because this is exactly the kind of sum that gets shared online with “AND BY 2030 IT WILL DO A YEAR’S WORK IN A SECOND!!!” in capital letters. Trends don’t have to continue. Plenty of things that doubled for a while have stopped. The lily pond has edges.

So let’s look hard at the small print, because it’s real.

THE HONEST SMALL PRINT

Half the time isn’t reliably. The time horizon is the length of task the AI completes about half the time. Would you hire someone who got the job done half the time? For a lot of real work, including accounts, medicine and anything legal, you need it right almost every time. When METR measured at 80 per cent success instead of 50, the horizons were much shorter. In May 2026 it estimated that the model it could measure at sixteen hours or more at 50 per cent success was good for only about three hours at 80 per cent.

It’s mostly technical work. METR’s tasks are mostly software and research tasks, which is exactly where AI tends to be strongest. It doesn’t automatically follow that the same curve applies to managing a building site, calming an upset customer or teaching a room full of fourteen-year-olds.

The margins are huge. METR itself says its uncertainty is very wide. For the model it put at fourteen and a half hours, the plausible range ran from about six hours to more than ninety. Only a handful of its longest tasks had actually been timed with real people doing them. METR is honest about the limits of its own ruler, and I respect that.

A test isn’t a workplace. Finishing a well-defined task in a clean test environment is a different thing from doing it inside a real company with messy data, vague instructions, office politics and a boss who changes their mind.

So the fair summary is this: trust the trend, distrust any single number. The exact figures are shaky. The direction is not. Every independent measure I’ve looked at points the same way. These systems are getting able to do longer, more complicated chunks of work on their own, and they’re doing it quickly.

The International AI Safety Report, a major review published in February 2026, led by the Turing Award-winning computer scientist Yoshua Bengio and written by more than a hundred experts from over thirty coun-

tries, reached a similar conclusion. Capabilities are advancing quickly, and faster than the rules meant to manage them.

THE JAGGED EDGE

The second big idea in this chapter is the one that will save you from embarrassing mistakes. AI isn't evenly good.

In 2023 researchers from Harvard Business School and elsewhere ran a proper experiment with 758 consultants from Boston Consulting Group, some of the most highly trained professionals in business. Some got access to GPT-4, one of the leading AI models of the day, and some didn't. Then everyone was given realistic consulting tasks.

On tasks the AI was good at, the consultants using it completed about 12 per cent more tasks, finished about 25 per cent faster, and produced work rated around 40 per cent higher in quality. That's a huge jump, for elite professionals.

But the researchers had slipped in a trick. One task was chosen because it *looked* like something the AI would be good at, but actually needed careful judgement that the AI tended to get wrong. On that task, consultants using AI were 19 percentage points less likely to reach the right answer than those working alone. They trusted the machine, the machine was confidently wrong, and being smart and highly trained didn't save them.

The researchers called this the "jagged technological frontier". I call it the Jagged Edge. Picture a coastline rather than a straight wall. In some places the AI's abilities stretch way past what you'd expect. In others, right next door, they fall well short. And from where you're standing you can't see where the coastline runs, because the AI sounds exactly as sure of itself on both sides.

The danger isn't that AI is stupid. The danger is that it's brilliant and stupid at the same time and doesn't tell you which one you're getting.

That's why the people who'll do best aren't the ones who use AI the most. They're the ones who know where the edge is in their own field and check the work that falls near it. That knowledge comes from experience, which is one more reason experienced workers have an advantage, if they learn to use it.

THE PERCEPTION TRAP

One more study, because it's one of my favourites and it's a bit uncomfortable.

In early 2025 METR ran another experiment. It took sixteen experienced software developers working on big open-source projects they already knew inside out. For each of 246 tasks, a coin flip decided whether they could use AI tools or not. Before starting, the developers predicted AI would make them about 24 per cent faster. Afterwards they *believed* it had made them about 20 per cent faster. It had actually made them about 19 per cent slower.

METR now describes that result as historical. The tools have improved a lot since early 2025 and it has changed how it runs these studies. So I'm not telling you AI slows people down; in plenty of other studies it clearly speeds them up. I'm telling you something about us. We're bad at judging whether AI is actually helping. It feels productive. The words appear so fast, so fluently. That feeling isn't the same as the work being better or quicker. So measure. Time yourself with it and without it, and check the output.

WHERE THIS LEAVES US

The amount of work AI can do on its own is growing fast and has been for years; the best measure we have suggests it's doubling every several months. The exact numbers are uncertain and the tests don't cover everything, but the direction is consistent across many different measures. AI is jagged: great at some things, bad at others, and equally confident either way. And our own sense of whether it's helping is unreliable.

Which brings us back to the pond. If you look at AI today and think "it's useful, but it can't do my job", you might be completely right. Today. The better question is this: if the length of work it can handle keeps doubling a few more times, which parts of my job does it reach, and when?

That isn't a frightening question. It's a planning question, and planning questions have answers. Part Two is where we start answering it.

DO THIS NOW

1. Find the edge in your own world. Give an AI one task from your job that you're genuinely expert at, look for where it goes wrong, and write down exactly what it missed. That gap is part of your value.
2. Time it. Pick a small task you do often, such as an email, a quote or a summary. Do it once without AI and time yourself. Next time do it with AI, including checking and fixing, and time that too. Compare honestly.
3. Tell someone the lily pad riddle tonight and watch their face at "day twenty-nine". Explaining an idea is the fastest way to own it.

CHAPTER 4

2027–2028: The Agent Years

In this part of the book I’m going to do something most writers on this subject avoid. I’m going to describe the next few years as specifically as I honestly can, and every time I say something I’ll tell you how confident I am.

I’m doing it this way because vague warnings are useless. “AI will transform everything” doesn’t tell you whether to renew your lease, change your course, have a word with your boss or start something on the side. Specifics do. Even specifics that might turn out wrong are more useful than fog, as long as you know how likely they are to be wrong.

So here’s my best, clearly labelled attempt at 2027 and 2028. To be clear about what this is: it’s an informed reading of where current trends point, not a prophecy.

FROM “HELP ME” TO “DO IT”

If 2023 to 2025 were the chatbot years, when most people’s experience of AI was typing a question and getting an answer, I expect 2027 and 2028 to be remembered as the agent years. Remember the kitchen from Chapter 2. For a lot of office work, this is when AI moves from recipe book to cook.

Here’s one plausible picture of what that looks like.

It’s a Tuesday in 2028. A small accounting practice with six staff has an AI agent set up on its systems. Overnight the agent has downloaded bank feeds for forty clients, matched about nine in ten transactions automatically, flagged the rest with a note explaining what’s odd about each, drafted reminder emails to the six clients whose paperwork is late, and prepared first drafts of three quarterly tax returns. When the team comes in, their day isn’t

doing that work. It's checking, fixing, deciding, and talking to clients about it.

That's not science fiction. Pieces of it exist in 2026. What changes is that it becomes normal, reliable enough to trust with real work, and cheap enough for a six-person firm.

TRAJECTORY If current trends continue, AI agents that complete multi-hour chunks of computer-based work, with a person checking, will be common in offices by 2028, especially in finance, administration, customer service, software, marketing and research. How reliable they are will vary a great deal from task to task.

Why do I put this down as a trajectory rather than a mere possibility? Three reasons. The Doubling Clock from Chapter 3 is already there on technical tasks. The big technology companies are building exactly this and selling it to businesses right now. And the business incentive is enormous: work that runs overnight, never takes a sickie and costs a fraction of a salary.

Why not higher than that? Because, as Chapter 1 showed, most organisations are slow and clumsy at changing how they actually work. The technology will be ready before a lot of workplaces are.

WHAT THIS COULD MEAN FOR JOBS

The entry-level squeeze is likely to tighten. The Stanford payroll research showed young workers in AI-exposed jobs falling behind through 2025 and 2026, mainly through reduced hiring. If agents get better at exactly the tasks junior staff used to do (first drafts, data cleaning, basic research, simple code, routine customer questions), I'd expect that gap to keep widening.

Teams will probably get smaller rather than empty. The common pattern won't be "the whole department's gone". It'll be "the department used to be fifteen people, it's now nine, and the nine are expected to use AI to do what fifteen did". Headcount shrinks through retirements, resignations and roles that are never refilled.

Some big, public layoffs will happen, and some of them will be walked back. A few companies will cut hard and announce it proudly. Some of those will discover, as the Swedish payments company Klarna did (Chapter

7), that customers still want a human when things get complicated, and that quality drops when you cut too deep.

Experienced people who can direct AI will become worth more. Someone with ten years in a field who can manage three agents and check their work is suddenly producing the output of a small team, and employers will pay for that. The payroll data already shows experienced workers in exposed jobs holding up much better than young ones.

And most jobs will still exist and look broadly familiar. This gets forgotten. In 2028 nurses will still nurse, sparkies will still wire houses, teachers will still teach and chefs will still cook. Most jobs will change around the edges, in the admin, paperwork and scheduling, while the core stays human.

WHERE PEOPLE ARE LIKELY TO NOTICE IT FIRST

Based on where the evidence points, here's roughly the order in which I'd expect people to notice change over 2027 and 2028. Software development is already changing; AI writes a large share of routine code, and the 2026 AI Index found employment of software developers aged 22 to 25 had fallen by nearly 20 per cent since 2024, while senior developers who can design systems and review AI-written code remain in demand. Customer service and call centres are close behind, with AI handling more of the first line and people handling the angry, the complicated, the vulnerable and the valuable. Then comes administration, bookkeeping, data entry, claims and other back-office processing, which is rule-based computer work that agents suit well. After that, marketing and copywriting, where volume content gets cheap but strategy and genuine customer understanding hold their value. Then paralegal, research and junior analyst work. Translation and transcription are already heavily affected.

What these share is that the work happens on a computer, it's mostly about processing information, and the quality can be checked afterwards.

WHAT PROBABLY WON'T CHANGE MUCH BY 2028

Honestly? A lot. Hands-on trades such as plumbing, electrical work, carpentry, air-conditioning and roofing: AI helps with quotes and paperwork,

but it isn't crawling under your house. Care work, including aged care, disability support, childcare and nursing, where the human, physical and emotional core stays. Hospitality. Emergency services and policing. And most licensed roles that carry legal responsibility, such as the doctor, the pharmacist or the engineer who signs off a building. AI assists; a person is still accountable.

Robots will come for some of this eventually. But on the evidence available in 2026, general-purpose robots are still largely in trials, and driverless vehicles are spreading city by city rather than everywhere at once. Chapter 9 goes into the detail.

THE DARKER SIDE OF THE AGENT YEARS

I won't pretend this period is only about jobs and productivity.

Every tool that helps a small business write a thousand personalised emails also helps a scammer write a thousand personalised scams. In 2025 Americans reported nearly US\$21 billion in losses to online crime to the FBI, and people over sixty reported US\$7.75 billion, up 59 per cent in a single year. For the first time, the FBI's annual report included a section specifically on complaints involving AI. Australians reported A\$2.18 billion in scam losses in 2025, and older Australians bore a disproportionate share. As voices and videos become cheaper and better to fake, I'd expect voice-clone scams aimed at families ("Mum, I've been in an accident, I need money") to become far more common.

Trust in what you see online is likely to keep falling. When anyone can generate a realistic video of anyone saying anything, "I saw it with my own eyes" stops meaning much. That has big consequences for news, elections and family arguments.

And the rules will arrive slowly and patchily. In Europe, the EU AI Act's transparency rules, requiring people to be told when they're dealing with an AI and AI-generated content to be labelled, began applying in August 2026, while the stricter rules for "high-risk" uses in hiring, lending and education were pushed back to December 2027 for most systems. The United States has a patchwork of state laws, such as California's rules on AI companion chatbots, which took effect in January 2026. Australia chose in December

2025 not to bring in a standalone AI law for now, relying instead on existing laws and a new advisory AI Safety Institute.

My honest view is that rules will help at the edges. They won't stop the underlying change, and they'll mostly arrive after the problems they're responding to.

TRAJECTORY AI-enabled fraud grows substantially, and regulation stays slow, patchy and reactive. Chapter 16 gives you the defence.

SKILLS LIKELY TO BE WORTH MORE

If I had to bet on what employers will pay more for in this window, I'd start with the ability to direct AI: breaking a job into pieces an agent can do, briefing it well and checking what comes back (Chapter 22). Close behind is judgement in your own field, knowing where the Jagged Edge is and what looks right but isn't. Then the human things: client and customer relationships, trust, staying calm, reading people. Then the unglamorous job of making AI actually work inside a business, setting it up, fitting it into real processes and training people, because remember all those pilots that weren't producing results. Then cybersecurity and fraud awareness. Then the physical and technical trades, particularly electrical work and anything connected to power, cooling and construction for the data-centre boom. And finally, being the person who can be held responsible.

WHAT TO DO BEFORE 2027 ARRIVES

This is the whole point of the chapter.

If you work in an office, get genuinely good at using AI for your own work this year, while it's still an advantage rather than an expectation. Volunteer to be the person on your team who figures it out. Being the one who brings the change is a very different position from being the one it happens to.

If you're early in your career, don't rely on the traditional entry-level route alone. Build proof of what you can do: a portfolio, real projects, freelance work, anything someone can see (Chapter 11).

If you manage people or run a business, map which tasks agents could plausibly do in your business by 2028, and decide now how you'll handle it decently: retraining, redeployment, honest communication. People remember how they were treated.

If you work with your hands, use AI for the admin you hate. It's the biggest near-term win available to you.

If you're a parent, have the scam conversation with your family and set up a safe word this week (Chapter 16).

And everyone: start building a cash buffer, however small (Chapter 12). Uncertain times reward people with a bit of breathing room.

THE TRAP TO WATCH FOR

Because most jobs will still look familiar in 2028, because the nurse is still nursing and your office still has people in it, it will be very easy to look around and decide it was all hype.

Some of it will have been. But remember the pond. The agent years are likely to be the part of the curve where the lily pads are clearly spreading but most of the water is still open. That isn't the time to relax. It's the best time you'll get to prepare.

DO THIS NOW

1. Write a paragraph describing what an AI agent could plausibly do overnight in your job or business in 2028, like the accounting example above. Be concrete.
2. Circle the work in that paragraph that would still need you in the morning. That's your future job description. Start getting better at that.
3. Offer to test one AI tool your team could use this month.

2029–2031: Three Futures

Here's where honest writers have to change gear.

For 2027 and 2028 we can see the trend lines, the products being built and the money being spent, and we can make reasonable calls. By 2029 to 2031 the fog thickens. Small differences now (how fast the Doubling Clock keeps ticking, how quickly businesses actually adopt this stuff, what governments do, whether there's a recession, whether there's an unexpected breakthrough or an unexpected wall) compound into very different worlds.

So instead of pretending to know which world we'll get, I'm going to describe three. These are scenarios, not predictions. Think of a weather forecaster saying there's a 30 per cent chance of storms, a 50 per cent chance of showers and a 20 per cent chance of sunshine. You pack for all three and you keep an eye on the sky.

For each one I'll describe what daily life might feel like, who gets hit, where the opportunities are and, most usefully, the warning signs that would tell you it's the one we're heading into. Then I'll give you the moves that pay off whichever one turns up.

SCENARIO ONE: THE SLOW BURN

AI keeps improving, but the jumps get smaller. Perhaps the Doubling Clock slows: the systems keep getting better on tests, but turning that into reliable real-world work proves harder than expected. Or perhaps the technology races ahead but businesses, as they usually do, adopt it slowly, held back by messy data, legal worries, cost, habit and plain old organisational inertia. Think how long it took offices to make proper use of computers. The technology was around for decades before most workplaces were transformed by it.

In this world, by 2031 most people use AI every day the way they use search engines now: for writing, research, planning, homework help and customer service. It's useful and unremarkable. Some jobs have shrunk (fewer junior admin roles, fewer call-centre staff, fewer entry-level coders), but new roles have appeared too, and the overall job market looks broadly like it did after previous waves of technology. The people hit hardest are the young and the entry-level in exposed office jobs, the pattern we can already see, continuing at roughly today's pace. The opportunities are in helping slow-moving businesses finally get AI working, and early adopters in every field pull steadily ahead.

You'd know you were in this scenario if unemployment stays roughly flat through 2028, if stories of AI disappointments start outnumbering stories of breakthroughs, if the young-worker gap in the payroll data stops widening, and if AI company revenues grow more slowly than their spending.

My honest feeling is that this is one of the more likely outcomes, but less likely than many sceptics assume.

SCENARIO TWO: THE STEADY CLIMB

The trend lines roughly continue. The Doubling Clock keeps ticking, maybe a little slower, maybe not. By 2030 AI agents can reliably handle a large share of computer-based tasks that take a person a day or more. Businesses adopt them unevenly, fast in some industries and slow in others, but steadily. Robots begin moving from trials into real paid work in warehouses, factories and logistics, although general-purpose home robots are still mostly a novelty.

This is a big reshuffle rather than a collapse. Lots of people are doing different jobs from the ones they did in 2025. Many office jobs have been redesigned around AI, with fewer people each doing more, and for a lot of professionals the job has shifted from doing the work to directing, checking and owning the work. Unemployment may rise somewhat, particularly for certain groups, and some regions and industries are hit hard. There's political pressure for retraining schemes, stronger unemployment support and new rules. There's also a lot of new work: implementing AI, the trades

powering the boom, care, and anything where a trusted human is the product.

The International Monetary Fund has warned that in most scenarios AI is likely to worsen inequality unless governments act. In this world, that shows. People who learned to direct AI early, and people who own assets, do very well. People who didn't adapt fall further behind. Those hit hardest are entry-level and mid-level office workers whose work was mostly processing information, some professionals whose value lay mainly in knowing things rather than judging, deciding or relating, and drivers in cities where autonomous vehicles expand. But the opportunities are very large. One person with the right skills and a few AI agents can run a business that used to need ten staff.

You'd know you were in this scenario if measurements show AI reliably completing multi-day tasks by 2028, if the young-worker gap keeps widening and starts showing up among mid-career workers too, if big companies report smaller workforces alongside growing profits, if governments start seriously debating income support linked to automation, and if robot deployments move from hundreds to many thousands.

I think this is the most likely of the three, but I hold that view loosely.

SCENARIO THREE: THE FAST BREAK

Something speeds up. Perhaps AI systems become good enough at AI research to help build their successors much faster. Perhaps there's a breakthrough in reliability, so agents go from right half the time to right almost every time on long tasks. Perhaps the heads of some AI companies, who have said they expect extremely capable AI before the end of this decade, turn out to be right.

In May 2025 Dario Amodei, chief executive of Anthropic, one of the leading AI companies, warned that AI could eliminate up to half of entry-level white-collar jobs within one to five years and push US unemployment as high as 10 to 20 per cent. That's a serious person with inside knowledge making a serious warning. It's also one person's warning, not a consensus. By 2026 he was also stressing the other side of the ledger: that as AI makes work cheaper, demand for it may grow.

In this world, disruption is fast and widespread. Many kinds of office and knowledge work are substantially automated within a few years. Unemployment rises sharply among white-collar workers, a shock that, unlike most earlier waves of automation, hits the educated middle class hard. Governments scramble and emergency policies appear. Some things get much cheaper, and enormous wealth concentrates among those who own the AI systems. At the same time, science and medicine could speed up dramatically. Very good things and very disruptive things could happen side by side.

Broadly, anyone whose work is done mostly at a computer and mostly involves processing information would be exposed, including many professionals who currently feel safe. Physical, human, local and trusted work becomes relatively more valuable. So do owning assets and businesses rather than only selling your time, and being skilled at directing AI systems.

You'd know you were heading this way if AI systems start reliably doing week-long or month-long projects, if major breakthroughs announced by AI companies are confirmed by independent researchers, if graduate unemployment jumps sharply in a single year, if big companies announce very large cuts that stick rather than being walked back, and if governments launch emergency job programs or income support specifically tied to AI.

I think this is less likely than the other two, but it's not a fringe idea. Several of the people closest to the technology take it seriously. That's reason enough to prepare for it without planning your whole life around it.

WILD CARDS

A few things could happen in any of the three scenarios, and I can't predict them and won't pretend to.

WILD CARD A serious AI accident or misuse (a major cyberattack, a financial-market event, or an AI system behaving in a seriously harmful way) that triggers a big public and political backlash. A major recession unrelated to AI, which would hit jobs hard and muddy the picture. An energy or chip bottleneck: AI needs enormous amounts of electricity, and the International Energy Agency expects data-centre electricity use to roughly double by 2030, while the most advanced chips come from a small number of factories. Geopolitical conflict involving the countries and companies that make

those chips. And general-purpose home robots arriving much sooner than expected. That's possible, but I wouldn't bet on it in this window.

THE NO-REGRET MOVES

So how do you plan when you don't know which future is coming? You do the things that pay off in all three. Economists call these "no-regret moves".

The first is building a cash buffer. In the Slow Burn it's simple good sense. In the Steady Climb it buys you time to retrain. In the Fast Break it could be the difference between a stressful year and losing your home (Chapter 12).

The second is becoming genuinely good at directing AI. In the Slow Burn it makes you better than your peers, in the Steady Climb it's what employers pay for, and in the Fast Break it's one of the few skills whose value goes up (Chapters 14, 15 and 22).

The third is strengthening your human advantages: relationships, trust, hands-on skills, judgement and accountability. In every scenario these hold their value longer than pure information work (Chapter 6).

The fourth is building a second source of income, even a small one (Chapters 12 and 13). The fifth is protecting your family from scams, which are rising in every scenario (Chapter 16). The sixth is teaching your kids to think rather than just to answer (Chapter 19). And the seventh is staying informed, calmly. Check the warning signs once a quarter. Not daily, and not by doom-scrolling.

None of these requires you to know which future is coming. All of them make you stronger in any of them.

WHY THREE FUTURES INSTEAD OF ONE

I know it would be more satisfying if I just told you what's going to happen. Plenty of people will. Some will say it's all hype and nothing will change; others will say we're all doomed by 2029. Both camps will sell a lot of content.

But the right way to prepare for uncertainty isn't to guess the answer. It's to become someone who does well across the whole range of answers. That's what the rest of this book is for.

DO THIS NOW

1. Decide which of the three scenarios feels most likely to you. Write it down with today's date. You'll check back in a year.
2. Stress-test your household against the Fast Break: if your income dropped by half for six months, what would break first? Write down the first three things.
3. Put a recurring reminder in your phone for the first day of every third month: "Check the warning signs, Chapter 5."

CHAPTER 6

The Task Test

“Will AI take my job?”

It’s the question everyone asks, and I’m going to tell you something that might sound annoying at first. It’s the wrong question.

Not because it doesn’t matter, but because it can’t be answered. A job isn’t one thing. It’s a bundle of dozens of different tasks tied together with a job title and a pay slip.

A receptionist doesn’t “do reception”. A receptionist answers phones, books appointments, greets people, calms down upset customers, signs for deliveries, notices when a regular looks unwell, reminds the boss about the meeting they’ve forgotten, knows which supplier to ring when the printer dies, and a hundred other things.

AI doesn’t come for jobs. It comes for tasks, one at a time. Some fall quickly, some fall slowly, and some don’t fall at all. So the useful question is this: which of my tasks can AI do, which can it help me do, which need a human, and what happens to my job as that mix changes?

That question *can* be answered. By the end of this chapter you’ll have answered it for yourself, on paper, in about twenty minutes.

SIX WORDS PEOPLE MIX UP

Before we test your job, it helps to get some words straight, because most of the confusion in the headlines comes from treating them as if they meant the same thing.

Exposure means AI could plausibly do or substantially change some of the tasks in a job. It’s a measure of possibility, not a prediction. When the

International Monetary Fund says about 60 per cent of jobs in advanced economies are exposed to AI, it does not mean 60 per cent of people will lose their jobs. Roughly half of that exposure, the IMF reckons, is the kind where AI could make workers more productive rather than replace them.

Task automation is when AI does a task instead of a person. A book-keeper used to type in receipts; now the software reads them. The job still exists, minus one task.

Task augmentation is when AI helps a person do a task better or faster, but the person still leads it. When Anthropic analysed how people actually use its Claude assistant, a little over half of conversations in early 2026 looked like augmentation (learning, iterating, getting feedback) rather than handing the whole task over.

Job transformation is when enough tasks shift that the job looks different. A marketing coordinator who used to spend most of the week writing copy now spends it choosing, editing and testing AI-written copy, and talking to customers to work out what to say.

Job displacement is when there are fewer of a particular job, either because people are let go or, more often so far, because roles aren't refilled. This is the hiring freeze from Chapter 1.

And *new job creation* is the one people forget. When something gets cheaper, people often want more of it. When spreadsheets arrived, book-keepers worried; what actually happened was an explosion in financial analysis, because suddenly you could run numbers that had been impractical before. If AI makes legal help or custom software much cheaper, far more people and businesses may buy it, and that can mean more work, not less. Then there are jobs that didn't exist before at all: people who build and manage AI agents, people who check AI output in regulated industries, people who help businesses set it all up. We can't list them all yet, any more than anyone in 1995 could have listed "social media manager".

Robots replacing physical tasks are a separate track again, slower and more expensive than software, and Chapter 9 deals with them.

When you look at your own job, several of these may be happening at once. That's why we test tasks, not jobs.

THE TASK TEST

Every task you do falls into one of three buckets.

Automate. AI can do this task now, or probably within a couple of years, with little human input. Someone might glance over the result, but the machine does the work. Think typing up meeting notes, basic data entry, first-draft product descriptions, sorting emails, simple scheduling, translating a standard document or writing routine code.

Assist. AI can make you faster or better at this task, but a person still needs to lead it, shape it, check it or finish it. Think writing a report, preparing a quote, researching a problem, analysing figures, planning a lesson, drafting a legal letter, or diagnosing a fault from a customer's description.

Anchor. This task needs a human and is likely to for years, because it's anchored in the physical world, in human relationships, in trust or in responsibility. Think fixing a leaking pipe, comforting a frightened patient, negotiating with an angry client, deciding whether to hire someone, signing off a building as safe, teaching a nervous kid to swim, or being the person who says, "I'll take responsibility for this."

THE FIVE MOATS

How do you know if a task is an Anchor? Look for moats, the water around a castle that makes it hard to get in. There are five things that make a task hard for AI to take.

Hands. The task happens in the physical world. Someone has to be there to touch, lift, fix, build, clean or carry things in unpredictable places. A factory can be designed around a robot. Your kitchen, a building site or an old house with dodgy wiring can't, and those are still very hard places for machines.

Heart. The task depends on people wanting to deal with another human: comfort, care, motivation, persuasion, trust built over years. A grieving family doesn't want a chatbot. A nervous first-home buyer wants a person they trust.

Trust. The task depends on a real, known person or relationship. Your reputation in your town. The client who's been with you for twelve years. The customer who keeps coming back because they like *you*.

Taste. The task needs judgement about what's good: what fits *this* situation, *this* client, *this* brand. AI can produce a thousand options. Someone still has to know which one is right and why, and that comes from experience.

Accountability. Someone has to be responsible if it goes wrong: a licence, a signature, a legal duty, a professional reputation on the line. Society doesn't let software sign off a bridge, prescribe medication without oversight or take the blame in court, and in some areas regulations are moving towards requiring a human in the loop.

A task with none of these moats is likely to be automated sooner or later. A task with one moat is safer, though that moat might get crossed. A task with two or more is your strongest ground. Build your future there.

THE WORKSHEET

Grab a piece of paper or open a note on your phone. This takes about twenty minutes, and it's the most useful twenty minutes in this book.

Step one: list your tasks. Write down fifteen things you actually do in your job, and be specific. Not "customer service" but "answer customer emails about late orders". Not "admin" but "update the job-tracking spreadsheet". If you're stuck, walk through a normal week. What did you do on Monday morning? Tuesday afternoon?

Step two: estimate the time. Next to each task, write roughly what share of your working week it takes. They should add up to about 100 per cent. Near enough is good enough.

Step three: sort them. Mark each task Automate, Assist or Anchor. If you're unsure, ask yourself three questions. Could a smart intern who'd read everything do this from a computer, given clear instructions? That's probably Automate. Could the intern do a decent first draft that I'd have to fix or finish? Probably Assist. Does it need my hands, my relationships, my judgement or my signature? Probably Anchor.

Step four: count the moats. For each Anchor task, note which of the five moats protect it.

Step five: work out your exposure. Add up the time you spend on Automate tasks, on Assist tasks and on Anchor tasks. Then read your result against the table below.

If your Automate share is...	What it suggests
Under 20 per cent	Low exposure. Your job is mostly protected by moats. AI will change your admin, not your core. Use it to clear the boring stuff and grow your Anchor work.
20 to 40 per cent	Moderate exposure. Your job will noticeably change, and you'll probably be expected to do more in the same time. Get good at AI now and shift your time towards Assist and Anchor work.
40 to 60 per cent	High exposure. A significant chunk of what you're paid for could be done by AI within a few years. Your job may change a lot, or there may be fewer of them. Start building new skills this year.
Over 60 per cent	Very high exposure. Your role, as currently designed, is at real risk of shrinking. That's a verdict on the task mix, not on you. Your experience is valuable; it needs a new home. Chapter 11 is your next stop.

These thresholds are rules of thumb, not science. They're a way of seeing your job clearly, not a forecast of your future.

Step six: write your future job description. Imagine all your Automate tasks disappeared and your Assist tasks took half the time. What's left? What would you fill the freed-up time with? That's the job you want to be growing into, because that's the job that's likely to still exist.

WORKED EXAMPLE: A MEDICAL RECEPTIONIST

Here's the test run on a receptionist at a busy general practice. The percentages are typical estimates, not measured data.

Task	Share of week	Bucket	Moats
Answering phones and booking appointments	25%	Automate	—
Checking patients in at the desk	10%	Assist	Heart
Payments and billing queries	10%	Automate	—
Updating patient records and forms	10%	Automate	—
Repeat prescription requests	5%	Assist	Accountability
Calming upset or distressed patients	10%	Anchor	Heart, Trust
Noticing someone looks very unwell and alerting a nurse	5%	Anchor	Heart, Accountability
Managing the waiting room	5%	Anchor	Hands, Heart
Chasing test results and referrals	10%	Assist	—
Coordinating doctors' schedules	5%	Assist	Taste
General admin, filing, scanning, emails	5%	Automate	—

That's about 50 per cent Automate, 30 per cent Assist and 20 per cent Anchor: high exposure. AI phone systems that book appointments already exist, and billing and records are heading the same way.

But look at what's left. The human face of the practice. The person who notices the elderly man looks grey and gets a nurse straight away. The one who settles the mum whose toddler has a fever. The one who knows the regulars by name. A smart receptionist in this position could grow into a patient experience coordinator: handling the human side, managing the systems that do the booking, dealing with whatever the systems can't, and making sure vulnerable patients don't fall through the cracks. That job is likely to be more secure and could well pay more. But it goes to the person who can run the new systems *and* is great with people.

WORKED EXAMPLE: AN ACCOUNTANT AT A SMALL FIRM

Task	Share of week	Bucket	Moats
Data entry and bank reconciliation	20%	Automate	—
Preparing standard tax returns	20%	Assist, moving to Automate	Accountability
Preparing financial statements	10%	Assist	Accountability
Routine client emails	10%	Assist	—
Advising clients on business decisions	10%	Anchor	Trust, Taste, Accountability
Tax planning for complex situations	10%	Assist	Taste, Accountability
Meeting clients and building relationships	10%	Anchor	Heart, Trust
Dealing with tax office reviews and disputes	5%	Anchor	Accountability, Trust
Supervising and training junior staff	5%	Anchor	Heart, Taste

Moderate to high exposure, but with a clear path. The number-crunching shrinks and the advising grows. An accountant who uses AI to do the compliance work in a fraction of the time, and spends the freed-up hours being a proper adviser to small businesses, a sort of part-time finance director, becomes more valuable, not less. The real risk sits with the junior accountant whose whole job was the data entry and the standard returns. That’s the rung of the ladder that’s wobbling.

IF YOUR RESULT SCARED YOU

Take a breath. A high exposure result doesn’t mean you’ll be out of work next year. It means the *design* of your job is exposed. You aren’t your job design. You’re a person with experience, relationships, knowledge of your industry and the ability to learn.

Look again at your Anchor and Assist tasks. That’s where your future is. Over the next two years the goal is to move as much of your working time as

possible from Automate towards Assist and Anchor. And if your job really is almost all Automate, you've just found out early, which is the whole point of this book. You now have time to move deliberately rather than being pushed.

IF YOUR RESULT SEEMS TOO COMFORTABLE

Be a little suspicious of yourself. Remember the Jagged Edge and the perception trap from Chapter 3: people are bad at judging what AI can do, in both directions, and most of us believe our own work is special. Sometimes it is. Sometimes it's more routine than we'd like to admit.

A good check is to take every task you marked Anchor and ask an AI to attempt it, or to explain how it would. If it does a surprisingly good job, move that task up a bucket. Then think about the Doubling Clock. Your result is a snapshot. Look at your Assist list and ask which of those tasks it might do on its own in three years.

THE REAL LESSON

The people who get hurt worst by technological change are rarely the ones who saw it coming. They're the ones who assumed their whole job would be either fine or gone, and so did nothing.

The truth is almost always in between. Some of your tasks will go. Some will get easier. Some will become more important than ever. The whole game is to move towards the tasks that are becoming more important, before you're forced to.

The next five chapters go industry by industry, so you can see how this plays out in your own world and where people like you can go next.

DO THIS NOW

1. Do the worksheet: fifteen tasks, time shares, buckets, moats, result. Twenty minutes. Don't skip this one.
2. Write your future job description in three or four sentences.
3. Take the Anchor task you're least sure about and ask an AI to attempt it. Be honest about the result.
4. Pick one Assist task and use AI on it every time you do it this week.

Office, Admin, Finance, Sales and Support

If you work at a desk, on a computer, processing information for most of the day, this chapter is for you. I'm going to be straight with you, because you deserve that more than you deserve comfort.

This is where AI is hitting first.

Not because office workers are less skilled than anyone else. Because the work happens entirely on a screen, in words and numbers, which is exactly the material these systems are made of. The intern who's read everything is most at home in your inbox.

But hitting first isn't the same as wiping out, and the real story is more interesting, and more hopeful for the right people, than most headlines suggest. Start with the best real-world evidence we have.

WHAT HAPPENED WHEN FIVE THOUSAND SUPPORT AGENTS GOT AI

Economists Erik Brynjolfsson, Danielle Li and Lindsey Raymond followed 5,179 customer support agents at a large software company as it rolled out an AI assistant that suggested replies to customers in real time. Agents could use the suggestions or ignore them. The study was later published in one of the most respected economics journals in the world.

On average, agents resolved about 14 per cent more issues per hour. For newer and less-skilled agents, the improvement was about 34 per cent. For the most experienced, top-performing agents it made very little difference. Customers were also nicer to agents using it, and agents were less likely to quit.

Think about what that means. The AI had effectively learned how the best agents talked to customers and passed those habits on to everyone else. A new starter could perform like someone with months more experience.

That's wonderful news for the new starter. It's also a quiet warning. If a new person with AI can perform like an experienced one, a company may decide it doesn't need as many experienced ones, or as many people at all. Both of those things can be true at once, and that's the office-work story in a nutshell.

KLARNA: AI FIRST, THEN HUMAN AGAIN

Klarna is a Swedish payments company, the “buy now, pay later” button you see at online checkouts. In February 2024 it announced that its new AI assistant, built with OpenAI, had handled 2.3 million customer conversations in its first month, work the company said was equivalent to about 700 full-time agents. It resolved issues in under two minutes on average, compared with eleven for humans, and Klarna said it expected the assistant to add around US\$40 million to profit that year. It was held up everywhere as proof that AI could replace call centres.

Then, in May 2025, Klarna's chief executive, Sebastian Siemiatkowski, said something unusual for a tech boss. He told Bloomberg that “cost unfortunately seems to have been a too predominant evaluation factor” and that “what you end up having is lower quality”. Klarna began hiring people for customer service again.

It didn't rip out the AI, which still handles the bulk of simple questions. But it brought people back for complex and sensitive cases, and Siemiatkowski said it was critical that customers know “there will be always a human if you want”.

I'll come back to this story in Chapter 21, because it says something important about where human value lies. For now, here's what it tells office workers. AI can do a huge amount of the routine work. It can't yet do the difficult, emotional, high-stakes or unusual work as well, and customers notice the difference. The jobs that remain are the ones that handle what the AI can't, so that's the work you want to be excellent at.

ROLE BY ROLE

For each role below I've noted which tasks are most exposed, which are protected, and a practical way forward. Use these alongside your own Task Test from Chapter 6. The exposure ratings are my assessment from the research, not measured outcomes.

Administrative and executive assistants, office managers

The most exposed work is scheduling and calendar juggling, taking and typing up meeting notes, routine emails and letters, travel bookings, filing, formatting and data entry. Reports, presentations, supplier research and budgets are changing into Assist work. What's protected is being the trusted right hand who knows how the boss thinks and what matters to them, handling sensitive and confidential situations, managing office relationships, solving the physical problems of a workplace, and being the calm one in a crisis.

The best assistants will become operations coordinators who run the AI tools that handle scheduling and admin, and spend their own time on judgement, relationships and problem-solving. Learn to set up simple automations (Chapter 14), learn basic project management, and become the person in your organisation who makes the AI tools actually work. Exposure is high for routine admin roles and moderate for trusted executive assistants and office managers.

Bookkeepers, accounts staff and accountants

Data entry, bank reconciliations, invoice processing, chasing payables and receivables, standard tax returns and standard reports are the most exposed. Complex returns, analysis and forecasting, audit preparation and month-end processes become Assist work. Protected work includes advising clients on real decisions, complex tax planning, dealing with the tax office in reviews and disputes, understanding a particular business deeply, and being the trusted person who can say, "Don't do that."

There are two strong directions. One is to become an adviser, moving from “here’s your tax return” to “here’s what I’d do if I were you”. Small businesses are crying out for part-time finance directors who understand their business, and AI lets you serve far more of them. The other is to specialise in setting up and running automated finance systems for businesses. Exposure is high for data entry and junior bookkeeping, moderate for accountants who move towards advice.

Customer service and call centres

Answering common questions, password resets, order tracking, changing details, routine complaints and anything scripted are the most exposed. Complex troubleshooting and written replies become Assist work. What stays human is the angry, upset or vulnerable customer, the complex or unusual problem, the high-value client, situations needing discretion (bending the rules to do the right thing), and anything that needs real empathy.

The call centre of 2030 is likely to have fewer people handling harder calls, alongside people who *manage* the AI systems by checking what they get wrong, improving their responses and dealing with what they hand on. Aim to be one of those. Get good at the hardest calls, put your hand up to help test and improve your company’s AI tools, and consider moving towards customer success roles, which are more relationship-based. Exposure is high for first-line support and moderate for complex and specialist support.

Sales

Researching prospects, cold emails, updating the customer database, proposals from templates, basic follow-ups and scheduling are highly exposed. Meeting preparation, lead analysis and personalised proposals become Assist work. Building real relationships, reading a room, handling objections face to face, negotiating, complex high-value sales and working out what a customer actually needs rather than what they say they need all stay human.

AI is a gift to good salespeople, because it clears away the admin that eats most of a salesperson’s week. The people who’ll struggle are those whose value was mainly sending lots of emails. And when every inbox is flooded

with perfectly written, personalised, automated emails, the salesperson who picks up the phone, turns up in person or gets a warm introduction stands out more, not less. Exposure is low to moderate for relationship sales and high for high-volume transactional sales.

Human resources

Writing job ads, screening résumés, scheduling interviews, answering routine policy questions, onboarding paperwork and payroll admin are the most exposed. Designing training, analysing staff data and drafting policies become Assist work. Handling conflict, grievances and disciplinary matters, supporting staff through personal crises, making hard calls about people, building culture and deciding who to hire stay human.

AI used in hiring is one of the areas lawmakers treat as high-risk. Under the EU AI Act, stricter obligations for AI used in recruitment are due to apply from December 2027. Being the person who understands how to use AI in HR fairly and lawfully is valuable. Move towards advising managers on real situations, and towards workforce transition, which will be a big area as organisations restructure around AI. Someone has to help them retrain and redeploy their people well. Exposure is moderate to high for HR administration and low to moderate for advisory roles.

Insurance, claims, banking back office and data processing

Processing standard claims, checking documents, data entry, compliance checks against rules and straightforward loan or application processing are among the most exposed tasks anywhere in this book, because they're rule-based work on documents. Complex and disputed claims, fraud investigation, dealing with vulnerable customers (someone whose house has just burned down) and decisions that require human discretion and accountability stay human.

If you're here, don't wait. Move towards complex claims, investigation, fraud (which is growing, as Chapter 16 shows), customer advocacy, or compliance oversight of AI systems. Regulators increasingly expect people to check AI decisions about people's money. Be one of those people. Exposure is high.

Project managers and team leaders

Status reports, task tracking, scheduling, minutes and plan updates are exposed. Leading people, resolving conflict, making trade-off decisions, managing stakeholders and taking responsibility when things go wrong are not. Your job is shifting towards managing a team that includes both people and AI agents, a skill almost nobody has yet (Chapter 22). Exposure is low to moderate.

Marketing and content roles are covered in Chapter 10.

THREE THINGS EVERY OFFICE WORKER SHOULD DO THIS YEAR

Whatever your specific role, three moves apply to everyone in office work.

Become the AI person on your team. Not the tech expert, just the person who has worked out how to use it well for your team's actual work and can show others. In almost every workplace, that person becomes more valuable, more visible and harder to let go.

Move towards the hard stuff. Volunteer for the complex cases, the difficult clients, the messy problems. The routine work is what's going; the difficult work is what's staying. Build your reputation there.

And build relationships outside your organisation. If your job does shrink, what gets you the next one usually isn't your résumé. It's the people who know what you're good at. Invest in those relationships now, while you don't need them.

DO THIS NOW

1. Find your role above and compare it with your Task Test from Chapter 6. Adjust your result if you need to.
2. This week, volunteer for one difficult case, client or problem you'd normally avoid.
3. Show a colleague one AI trick that saved you time.
4. Message one person outside your workplace you've lost touch with, just to say g'day.

CHAPTER 8

Health, Education, Law and the Licensed Professions

If you spent years training for your job (nursing, teaching, law, medicine, pharmacy, physio, social work), you might be reading this with a particular kind of worry. I gave my twenties to this. I've got the debt to prove it. Is it all about to be worth nothing?

The honest answer is probably not. But the job you trained for will change more than your training prepared you for. And the parts that change are, in many cases, the parts you liked least.

WHY LICENSED PROFESSIONS HAVE A HEAD START

Think back to the Five Moats from Chapter 6: Hands, Heart, Trust, Taste and Accountability. Most licensed professions have at least three of them. A nurse has Hands (physical care), Heart (comfort and reassurance) and Accountability (a registration that can be taken away). A lawyer has Trust, Taste (judgement about strategy) and Accountability (duties to the court and the client). A teacher has Heart, Trust and a room full of real children who need a real adult.

The licence itself is a moat. Society has decided that certain decisions need a named, qualified, responsible person, and changing that means changing laws, regulations, professional standards and public trust, all of which move slowly. That doesn't make you untouchable. It makes you slower to touch, which means you have time. The question is whether you use it.

TRAJECTORY In licensed, hands-on and accountable professions, AI is likely to change the work around the core (documentation, admin, research and drafting) well before it changes the core itself.

HEALTHCARE

One of the most visible uses of AI in healthcare is aimed at the thing doctors and nurses complain about most: paperwork. AI “scribes” listen to a consultation, with the patient’s permission, and write up the notes.

The evidence on these is worth knowing in detail, because it’s a good example of the gap between hype and reality. Vendors have made some very large claims. The careful studies are more modest. In a randomised trial involving 238 doctors at UCLA Health in the United States, published in 2025, one AI scribe cut the time doctors spent writing notes by about 10 per cent compared with doctors who didn’t use one, and another made no significant difference. But doctors using either tool reported meaningfully lower workload and less exhaustion. Other studies have found similar patterns: modest time savings, less after-hours charting, and doctors who feel less burnt out.

That’s still worth having. A tired doctor who can look at the patient instead of the screen, and goes home a bit earlier, is a better doctor. But it’s not a revolution yet, and that’s typical of how AI is turning up in skilled professional work: useful, uneven, and smaller than the marketing.

AI is also increasingly used to help read scans, flag abnormal results, check for drug interactions, summarise patient histories and triage symptoms.

In 2016 the computer scientist Geoffrey Hinton, one of the pioneers of modern AI, said people should stop training radiologists, because it was “completely obvious” that deep learning would outperform them within five years. A decade later, radiology training positions in the United States were at record numbers and radiologists were among the best-paid specialists in medicine. Hinton has since said he spoke too broadly. AI did get very good at parts of reading images. It turned out that reading images was only part of the job.

Nurses

Documentation, handover notes, care plans, rostering and medication checks are changing. Physical care, assessing a patient in person, noticing that something's "not right", comforting frightened patients and families, advocating for them and working in a team under pressure stay human.

Nursing is one of the most AI-resilient professions in this book. The physical, emotional and accountable core is enormous, and many countries, Australia included, face nursing shortages. AI is likely to take some paper-work off your plate, and the nurses who do best will be comfortable with AI tools for documentation and monitoring and will use the time they get back well. Exposure is low.

Doctors

Note-taking, reading results and images with AI flagging, keeping up with research, letters and referrals, and initial triage are changing. Examining patients, diagnosing across several conditions at once, difficult conversations, decisions under uncertainty, procedures, and taking responsibility are not.

The role is shifting towards a skilled interpreter and decision-maker who uses AI as a powerful second opinion, and towards the human conversation. The Jagged Edge applies with real force here: AI can be brilliant on common cases and dangerously confident on unusual ones. Exposure is low to moderate.

Pharmacists

Checking scripts for errors and interactions, stock management and routine information queries are increasingly automated. Counselling patients face to face, clinical judgement on complex cases, vaccinations and other in-person services, and accountability are growing. Exposure is moderate.

Medical admin, billing and coding

This is where healthcare's highest exposure sits. Coding, billing, approvals, appointment booking and records management are rule-based information tasks. If you work here, treat Chapter 7's advice on admin and claims as yours too, and look at moving towards patient-facing coordination roles or clinical support training. Exposure is high.

Allied health

Physiotherapy, occupational therapy, speech pathology, psychology and dentistry are largely physical and relational. AI helps with notes, exercise programs, assessments and admin. Talk-based therapies face a particular question in the form of chatbots offering mental-health support, which Chapter 17 covers. My view is that AI may widen access to basic support, but the trusted relationship at the heart of therapy is hard to replicate and the risks of getting it wrong are serious. Exposure is low for physical work, low to moderate for talk-based work.

EDUCATION

This is where AI has moved fastest in everyday life. According to Stanford's 2026 AI Index, about 80 per cent of American high school and university students use AI for their studies, while only about 6 per cent of teachers say their school has clear policies on it. Read those two numbers again. An entire generation is working out how to use a powerful tool largely on its own, while the adults responsible for its education try to catch up.

AI can be extraordinary for learning. Two careful studies published in 2025 show how much. In Nigeria, a World Bank randomised trial gave secondary students six weeks of after-school English tutoring with an AI assistant, guided by teachers. The learning gains were equivalent to roughly one and a half to two years of normal schooling, and they showed up in end-of-year exams too. At Harvard, physics students using a carefully designed AI tutor learned more than twice as much as students in an active-learning class run by experienced instructors, and in less time.

The critical detail in both is that the AI was designed and supervised to make students think, not to hand them answers. Without that design, the same tools let students skip the thinking entirely. I call that the Lazy/Smart Fork, and it's at the centre of Chapter 19.

Teachers

Lesson planning, worksheets and resources, marking routine work, reports, adjusting material for different levels and admin are changing. Managing a classroom, building relationships with kids, noticing the child who's struggling at home, motivating the unmotivated, being a role model, pastoral care, and teaching kids how to think and how to use AI wisely all stay human.

I'm optimistic here. Teachers are drowning in admin and AI can take a lot of it. The job is likely to shift towards coach, mentor and designer of learning, with AI tutors helping students practise between lessons. The teachers who struggle will be the ones who try to ignore AI or ban it outright, because students will use it anyway. A teacher who becomes excellent at designing AI-supported learning that makes students think harder, not less, will be in demand. Exposure is low for classroom teachers and moderate for some lecturing and online course roles.

Tutors

This is a genuinely exposed area. Basic homework help and exam coaching are things AI does cheaply at any hour. Tutors who thrive will offer what AI can't: accountability, motivation, a relationship, work with kids who have particular needs, and helping families use AI well. Exposure is moderate to high.

LAW

Law is made of documents, and AI is very good with documents. Legal research, contract review, case summaries, standard letters and agreements are already widely handled with AI tools in larger firms. Law has also produced some of the best-known AI failures, including the cases in Chapter 2

where lawyers in New York and Melbourne filed lists of cases that didn't exist. It's a perfect illustration of the Jagged Edge: AI writes a convincing legal argument, and some of it is fiction.

Lawyers

Research, document review, standard contracts and letters, due diligence and summaries are changing. Advising clients on strategy, negotiation, advocacy in court, judgement on new situations, client relationships and professional accountability are not.

The traditional law-firm model, with lots of junior lawyers billing hours for research and document review and a few senior partners doing strategy and client work, is under real pressure, because the bottom of that pyramid is exactly what AI does well. Expect fewer junior roles, more pressure on hourly billing, and more value in judgement and relationships.

There's a new-demand story here too. Legal help is too expensive for most ordinary people. If AI makes it much cheaper, far more people might get legal help who currently go without, and that could create a lot of work for lawyers willing to serve everyday clients efficiently. Exposure is high for junior, research-heavy roles and moderate for experienced advisory and court work.

Paralegals and legal secretaries

These are among the most exposed roles in law. Move towards client-facing coordination, managing AI legal tools, specialist areas such as conveyancing oversight or compliance, or work with more human contact. Exposure is high.

OTHER PROFESSIONS

Engineers (civil, structural, mechanical, electrical) use AI for calculations, drafting, design options and documentation, but the engineer who signs off that a design is safe, and who understands the site, the regulations and the risks, remains central. Low to moderate exposure.

Architects can generate designs and visualisations quickly with AI; their value moves towards understanding clients, sites, regulations, buildability and taste. Moderate exposure.

Financial advisers face AI that builds portfolios and answers questions, but trust, behavioural coaching (“don’t sell everything in a panic”) and life planning remain human. Moderate exposure.

Social workers, counsellors, clergy and community workers do deeply human work built on Heart and Trust; AI helps with admin and case notes. Low exposure.

Police, emergency services and defence are physical, accountable, high-stakes and human. AI will be used heavily for analysis and surveillance, which raises serious questions of its own (Chapter 20). Low exposure for frontline roles.

THE TRAP FOR PROFESSIONALS

Here’s the risk I most want professionals to hear. It isn’t that AI replaces you. It’s that a professional who uses AI well replaces a professional who doesn’t.

The doctor who uses an AI scribe goes home on time with better notes. The lawyer who uses AI research tools produces sound advice faster and can charge less. The teacher who uses AI to personalise learning gets better results. If you’re waiting for your professional body to tell you how to use AI, you could be waiting a while, and colleagues who didn’t wait will pull ahead.

So use AI on your admin now: documentation, letters, research, summaries. Do it within your profession’s rules and your organisation’s policies, and never put identifiable patient or client information into tools that aren’t approved for it. Learn where the Jagged Edge sits in your field, because your professional knowledge is what lets you catch AI’s mistakes, and that is increasingly what you’re paid for. Lean into the human core: the conversation, the judgement, the physical skill, the accountability. And become the bridge. Every profession needs people who understand both the professional work and the AI tools. Help write your organisation’s AI policy. Lead the trial.

DO THIS NOW

1. Find out your organisation's AI policy. If there isn't one, that's useful information too, and an opportunity to help write it.
2. Use an approved AI tool on one piece of admin this week (a letter, a lesson plan, a summary) without including anyone's private information.
3. Write down three mistakes AI would be likely to make in your field that only a trained professional would catch. That's your value, on paper.

Trades, Transport, Factories and Warehouses

If you work with your hands, you’ve probably spent the last few years hearing about AI and thinking, good luck getting a chatbot to fix a burst pipe at two in the morning.

You’re right. For most of this book’s five-year window, you’re likely to stay right. But there’s more to this chapter than “you’re safe”. There are three stories here: one about robots, one about a boom that’s happening *because* of AI, and one about the biggest opportunity in front of tradespeople that most of them are ignoring.

STORY ONE: THE ROBOTS ARE REAL, AND SLOWER THAN THE VIDEOS

You’ve seen the videos. Humanoid robots doing backflips, robot hands folding washing, robot dogs climbing stairs, sleek machines striding through factories. Here’s what was actually going on in 2026.

Humanoid robots, the ones shaped roughly like people, are in real trials in factories and warehouses. Figure AI has reported its robots working shifts at a BMW plant in South Carolina, loading sheet-metal parts, and Agility Robotics’ Digit robots have been moving totes in warehouses for paying customers. But independent trackers of the industry reported that, as of mid-2026, no humanoid robot maker had sustained commercial deployments beyond the low hundreds of units. Tesla has built many of its Optimus robots, but on the company’s January 2026 earnings call Elon Musk said the several hundred deployed so far were “primarily for learning”, not productive tasks.

In homes, robots are much further behind. As mentioned in Chapter 2, when Stanford's 2026 AI Index looked at robot performance on household tasks, they succeeded only around 12 per cent of the time.

Why is the physical world so hard? Because it's messy. A factory can be designed around a robot: same parts, same place, same lighting, every time. A house built in 1962 with a crawl space full of spider webs, corroded fittings, a wall that isn't quite straight and a dog barking at you is a completely different problem. Human hands, eyes and judgement handle that kind of chaos without thinking. Machines, so far, can't.

HAPPENING NOW Humanoid robots are in early commercial trials, and deployments are small.

PLAUSIBLE SCENARIO Robots doing meaningful amounts of paid work in structured places such as warehouses and factories, at real scale, by 2029–2031.

WILD CARD General-purpose robots doing skilled trades work in homes and on building sites within this five-year window.

DRIVING IS DIFFERENT: IT'S ALREADY HAPPENING

There's one area where machines really are doing a physical job: driving.

Waymo, the self-driving company owned by Google's parent Alphabet, was providing around 500,000 paid rides a week by 2026 with nobody in the driver's seat, roughly ten times what it was doing two years earlier. By September 2026 it was operating in fourteen American cities and aiming for a million rides a week by the end of the year.

In trucking, Aurora Innovation began running driverless heavy trucks commercially on Texas highways in 2025, starting with the run between Dallas and Houston and adding longer routes, and said in mid-2026 that it was on track to finish the year with about 200 driverless trucks on the road.

This is the clearest example so far of AI taking over a physical job, and it's worth noticing *how* it's happening: city by city, route by route, in good weather first, on well-mapped roads first. Not everywhere at once.

HAPPENING NOW Driverless taxis and some driverless trucking are operating commercially in parts of the United States.

TRAJECTORY Continued expansion to more cities and highway corridors through 2031.

PLAUSIBLE SCENARIO A significant fall in demand for taxi, rideshare and long-haul drivers in some regions by 2031.

ROLE BY ROLE

Electricians

AI helps with quoting, invoicing, scheduling, reading plans, fault-finding (describe a problem, get the likely causes), compliance paperwork and marketing. Everything that involves being on site, including wiring, testing, diagnosing real faults in real buildings and safety sign-off, stays human. And electricians are among the biggest *winners* of the AI boom, as you'll see in a moment. Very low exposure, high opportunity.

Plumbers, gasfitters and air-conditioning technicians

AI helps with quotes, admin, customer messages, diagnosing from photos and descriptions, and marketing. The work itself stays human, because every house is different and every job is a puzzle solved with your hands. Heating, ventilation and cooling work is also in strong demand from data centres. Very low exposure, high opportunity in cooling.

Builders, carpenters, roofers, bricklayers, tilers and painters

AI helps with estimating, planning, ordering materials, showing clients what a design will look like, admin and marketing. The physical build stays human. One thing to watch is prefabricated and modular housing, which moves some construction into factories where automation is easier, but that's a slow shift. Low exposure.

Mechanics and automotive technicians

Diagnostics are increasingly AI-assisted, along with parts ordering, customer communication and service records. Physical repair stays human. The bigger change is electric vehicles, which have fewer moving parts and need different

skills; mechanics who get qualified on EVs and high-voltage systems are positioning themselves well. Low exposure, moderate change.

Cleaners

AI helps with scheduling, booking, quoting and marketing. Most cleaning in homes and varied commercial spaces stays human. Robot vacuums exist and keep improving, but the full range of what a cleaner does (bathrooms, kitchens, clutter, and judgement about what matters to this particular client) is still very hard for robots. Robotic floor cleaning is growing in large, open, simple spaces such as warehouses, airports and shopping centres. Low exposure in homes, moderate for large commercial floors.

Truck drivers

Driverless trucks are operating on some US highway routes, and this is real. The likely pattern is that long, straight highway runs between big depots are automated first, while local delivery, city driving, complicated loading, bad weather, rural roads and specialist freight (dangerous goods, oversize loads, livestock) stay human much longer.

If you drive long-haul highway routes, particularly in the southern United States, this is the most directly exposed role in this chapter. It won't happen everywhere overnight; it will spread route by route. In Australia, long distances, road trains, mixed conditions and regulation mean it's likely to come later. But it's worth planning now. Your knowledge of roads, loads and logistics is valuable, and it needs a new home before it needs one urgently: local and specialist driving, fleet coordination and logistics, supervising autonomous trucks from terminals, diesel and EV mechanics, dispatch, or heavy equipment operation. Exposure is moderate to high for long-haul highway driving and low to moderate for local and specialist work.

Taxi and rideshare drivers

Robotaxis are operating commercially in a growing number of American cities. If you drive in a city where Waymo or a competitor operates, watch closely. Elsewhere, including most of Australia, you probably have more

time, but the direction is clear. Possible next steps include specialist transport where care and assistance matter (patients, older people, people with disabilities), delivery work that involves more than dropping a parcel, logistics, or retraining into a trade. Exposure is high in cities with robotaxis and moderate elsewhere by 2031.

Warehouse workers

Warehouses are already among the most automated workplaces, with conveyors, sorting machines and robots that bring shelves to workers. Humanoid and mobile robots are being trialled for picking and moving. Picking unusual items, handling damaged goods, problem-solving, maintenance and anything needing flexibility stay human for now. The growth jobs are in maintaining, fixing and supervising the machines. Exposure is moderate to high; a mechatronics or maintenance qualification is a strong next step.

Factory and manufacturing workers

Manufacturing has been automating for decades, and AI and robotics continue that. Repetitive, predictable work in controlled environments is most exposed. Skilled trades within manufacturing, such as maintenance fitters, electricians, toolmakers and technicians, remain in demand. Moderate exposure for repetitive roles, low for skilled maintenance and technical roles.

STORY TWO: THE BOOM NOBODY EXPECTED

Here's the plot twist in the AI story that almost nobody saw coming.

Artificial intelligence runs on enormous buildings full of computers called data centres, and those buildings need electricity, a staggering amount of it. The International Energy Agency expects global data-centre electricity use to roughly double by 2030, to slightly more than the whole of Japan uses today.

Someone has to build those buildings. Wire them. Cool them. Run the power lines to them. Maintain the generators. Build the new power stations and grid connections. That someone isn't a chatbot.

In a 2026 survey of American construction firms by the Associated General Contractors of America and the National Center for Construction Education and Research, 58 per cent of contractors doing data-centre work said those projects had increased competition for skilled workers, and nearly half reported pressure to raise wages. Put simply, AI has started a gold rush, and some of the people selling the shovels are electricians, air-conditioning technicians, line workers, construction workers, welders and engineers.

HAPPENING NOW AI infrastructure is driving strong demand for skilled trades in construction, electrical and cooling work, at least in the United States.

TRAJECTORY This continues through at least 2030, based on announced investment and energy projections.

For young people choosing a path, and for anyone thinking about retraining, this deserves serious attention. An electrical apprenticeship may be one of the most AI-resistant, in-demand career choices available right now. I don't say that lightly.

STORY THREE: THE OPPORTUNITY TRADIES ARE IGNORING

AI isn't coming for your *work*. But it may be the best thing that's happened to your *business* in years.

Ask most self-employed tradies what they hate about the job and it won't be the tools. It'll be writing quotes at nine o'clock at night, chasing invoices, answering the same phone questions twenty times a week, trying to keep up with social media, writing up job notes and compliance paperwork, juggling cancellations, and doing the books.

AI can help with nearly every one of those, today. Talk a quote into your phone on the drive home and have AI turn it into a professional written quote. Take photos of a finished job and have AI write the Facebook post and the Google Business update. Set up an assistant that answers common questions and books jobs. Turn a voice memo into compliance notes. Use accounting software with AI features to match receipts and chase invoices.

Think about what that does. The tradie who gets admin down from ten hours a week to three isn't just less stressed. They can take on more jobs, respond faster, look more professional and win more work than the one

down the road still writing quotes on the back of an envelope. This is the lowest-hanging fruit in the entire book, and Chapters 14 and 15 give you the practical how-to.

THE HONEST RISK FOR TRADES

I don't want to give you a false sense of security, so here's the realistic risk for trades over this five-year window. It's not robots. It's other people.

If AI disrupts a lot of office work, some of those workers will retrain into trades. That means more competition for apprenticeships and, eventually, for work. And if white-collar incomes are squeezed in a downturn, people put off renovations and non-urgent jobs.

So the smart move for trades isn't to relax. Get the highest qualification you can in your trade, because licences and specialisations are moats. Specialise in areas of growing demand, such as electrical, solar and batteries, EV charging, cooling, data-centre work and smart-home systems. Use AI to run a tighter business than your competitors. Build your reputation and your reviews, because trust is a moat and online reviews are how trust scales. And think about training others. If people move into the trades, somebody has to teach them.

DO THIS NOW

If you're in a trade: pick your most-hated admin task and try doing it with AI tonight. Ask your last three happy customers for a review. Look up one specialisation course in your trade connected to energy, EVs, solar, batteries or cooling.

If you drive, or work in a warehouse or on a factory line: find out whether autonomous vehicles or robots are being trialled in your area or your company. Look up one maintenance, technician or electrical course near you and write down the cost and length. Then do the Task Test from Chapter 6, honestly.

Creative and Tech Work

For most of history, the jobs people assumed were safest from machines were the creative ones and the clever ones. Machines would take the boring, repetitive work and leave the art, the writing, the design, the coding and the invention to us.

It turned out almost exactly the other way round. The first jobs to feel real pressure from modern AI weren't factory jobs. They were copywriters, illustrators, translators, junior coders and stock photographers. The robots came for the poets before the plumbers.

If you do creative or technical work, that's been strange and sometimes painful to watch. So let's talk about it honestly, and then about why, for the right people, this may also be the best time in history to be a creator or a builder.

WHY CREATIVE AND TECH WORK GOT HIT FIRST

Go back to the intern who's read everything. What has that intern read? Billions of words of writing, millions of images with their descriptions, and enormous amounts of computer code. It's no surprise that producing *more* writing, images and code is where it's strongest. It's working with exactly the materials it learned from.

In software the change has been dramatic. On one widely used test, which checks whether AI can fix real bugs taken from real software projects, the best systems went from solving only a small fraction of the problems in 2023 to solving more than 90 per cent of a verified set by early 2026. Stanford's 2026 AI Index reported that employment of software developers aged 22 to 25 had fallen by nearly 20 per cent since 2024.

In writing and design, the change has been messier to measure but just as real for many freelancers. A study of the freelance platform Upwork, published in 2024, found that after ChatGPT and AI image generators were released, freelancers in the most affected fields saw falls in both the amount of work they got and what they earned, and that the best-rated freelancers weren't spared. Clients could now get "good enough" for almost nothing.

HAPPENING NOW Routine creative and coding work is already significantly affected, especially for juniors and freelancers doing volume work.

MAKING THINGS VERSUS MAKING THINGS MATTER

Here's the idea I want creators to hold onto. AI has made it nearly free to make things. It hasn't made it any easier to make things *matter*.

Anyone can now produce a blog post, a logo, a jingle, a video or an app, which means there's an ocean of stuff. And when there's an ocean of stuff, three things become much more valuable. Taste: knowing what's good, what's right for this audience, what to cut, and what's worth making at all. Voice: a recognisable human perspective that people come back to because it's you. And trust: an audience that knows you, a client who relies on you, a reputation that means something.

A machine can generate a thousand logos. A great designer knows which one is right, and more importantly can sit with a client, understand what their business actually stands for, and make them feel understood. That's the job now.

ROLE BY ROLE

Software developers

Writing routine code, boilerplate, simple bug fixes, tests, converting code between languages, basic web and app building, and documentation are the most exposed. Building features is becoming Assist work, with AI writing much of the code while the developer directs, reviews and fixes it. Understanding what the business actually needs, system design, security, reviewing AI-written code for subtle problems (the Jagged Edge again, because AI

code can look right and hide bugs), working out with people what to build, and taking responsibility for systems that mustn't fail all stay human.

The squeeze is at the bottom of the career ladder. Juniors are finding it hard to get hired because much of what juniors used to do, AI now does. Experienced developers who can design systems and direct AI have, so far, done well.

If you're junior, build real things and show them; a portfolio of working projects beats a certificate. Learn to use AI coding tools brilliantly, because you're competing with people who do. And go deep in one domain, such as health, finance or logistics, because understanding a business problem is what makes you more than a code-writer. If you're experienced, get fluent with AI agents that write code. Your job becomes reviewing, directing and designing, and you may produce the output of a whole team. Exposure is high for junior and routine development, moderate for senior engineers who direct AI.

Writers, copywriters and content creators

Search-optimised blog posts, product descriptions, generic marketing copy, social media filler, basic press releases and summaries are highly exposed. Original reporting and research, a distinctive voice, deep subject knowledge, storytelling from real experience, strategy (knowing what to say and why) and relationships with editors, clients and readers are not.

The market for words by the kilo has largely collapsed. The market for genuinely good, original, trustworthy writing may actually be stronger, because readers are drowning in bland AI text and crave something real. Move up the value chain, from writer to editor and strategist, from generalist to specialist in a field you know deeply, and from anonymous to known by building a name and an audience. Use AI for research, first drafts and grunt work; bring the voice, judgement and truth yourself. Exposure is high for volume content and low to moderate for distinctive, expert or original work.

Graphic designers and illustrators

Stock-style images, simple social graphics, basic logos, resizing, variations and mock-ups are highly exposed. Brand strategy and identity work, understanding clients, art direction, and distinctive personal styles that people seek out are not.

There are serious, unresolved legal and ethical questions about AI tools trained on artists' and authors' work without permission, and the courts are only starting to answer them. In November 2025 England's High Court largely rejected Getty Images' copyright claims against the image-generator company Stability AI, while finding limited trade mark infringement. In the United States, Anthropic agreed in 2025 to pay US\$1.5 billion to settle a class action brought by authors whose books it had downloaded from pirate websites. How these questions are finally resolved could change this market significantly.

Become an art director, the person who guides AI tools and human work towards a coherent vision. Build a distinctive style. Get closer to clients and business strategy. And where clients value it, offer human-made work as a feature (Chapter 21). Exposure is high for routine design and moderate for brand and art direction.

Video, photography and music

AI tools increasingly handle cutting, captioning and colour, and can even generate footage. Storytelling, pacing, emotional judgement, live shoots and working with real people stay human. Moderate to high exposure for routine editing.

Stock photography is heavily exposed. Weddings, events, portraits and news photography, anything that captures a real moment with real people, much less so. People want real photos of their real wedding.

AI can generate background music, jingles and generic tracks. Live performance, real artists people connect with, and original artistic vision remain human.

Marketers

Producing campaign content, ad variations, reporting, basic analysis and email sequences are highly exposed. Understanding customers, strategy, brand, creative direction, knowing what will actually move people, and building partnerships are not. One strong marketer with AI can now do what a small team used to, which means fewer jobs in production and more value in people who understand customers deeply and can direct the machine. High exposure for production, moderate for strategy.

Data analysts

Writing queries, building standard dashboards, cleaning data and producing routine reports are exposed. Asking the right questions, understanding the business, telling a clear story with numbers, knowing when data is misleading, and persuading people to act are not. AI now lets managers ask questions of their data in plain English, so the analyst who only produced reports is exposed, while the analyst who helps the business make better decisions is more valuable than ever. Moderate to high exposure.

THE CREATOR'S OPPORTUNITY

Now the good news, and it's genuinely good.

For most of history, making things at scale needed money, equipment and teams. A film needed a crew, an app needed developers, a magazine needed a printing press, a business needed staff. AI is collapsing those costs. A single person with good ideas, taste and AI tools can now write, design and publish a book; build a working app or website; produce a professional-looking video series; run a small media business; create and sell digital products; or design a brand from scratch.

The barrier is no longer "can I afford a team?" It's "do I have something worth saying, and can I get people to care?" That's a very different barrier. It rewards originality, persistence, real expertise and genuine connection with an audience. It rewards people who've lived something and can speak to others who are living it.

The future doesn't belong to people who can make things. Machines can make things. It belongs to people who know what's worth making, and whom other people trust to make it.

DO THIS NOW

1. Name your taste. Write down three things in your field that are genuinely good and three that are rubbish, and why. That “why” is what AI can't give you, so get better at putting it into words.
2. Build one thing with AI this week that you couldn't have built alone: a simple web page, a short video, a small tool, a designed PDF.
3. Show your work publicly: one post, one portfolio piece, one project. Visibility is a moat.

CHAPTER 11

If You Do X, Start Learning Y

This is the chapter I'd want to hand to someone who has just done the Task Test, got a high exposure result, and is sitting at the kitchen table thinking, right, so now what?

Here's the most important thing I can tell you about changing direction. You are not starting from zero. You never are.

Every job you've done has given you things that transfer: knowledge of an industry, an understanding of customers, skills with people, problem-solving habits, reliability, and a pile of things you know so well you've forgotten you know them. A good career change doesn't throw all that away. It finds the next door, a role where your existing experience is an advantage and you only need to add a few new skills to walk through.

This chapter is a map of those doors.

HOW TO READ THE RESKILL MAP

For each current role you'll find your transferable strengths, some next doors where those strengths give you a head start, what to learn, and a first step you could take this month. These aren't the only options. They're starting points to get you thinking.

A few honest cautions. Some of these roles are still emerging, and job titles vary a lot between organisations and countries. Pay and demand depend heavily on where you live, so check real job ads in your area before you commit to anything expensive. And be very careful with expensive courses. There's a gold rush of "AI certifications" of wildly varying quality, and many of the best learning resources are free or cheap.

THE RESKILL MAP

Office and administration

Receptionist or front desk. Your strengths are people skills, staying calm under pressure, organisation, and knowing how a business runs day to day. Next doors include client or patient experience coordinator, office operations coordinator, the person who runs the AI booking and phone systems and handles whatever they can't, and practice manager. Learn the AI booking and phone tools in your industry, basic automation (Chapter 14), customer experience principles and simple project management. First step: offer to lead your workplace's trial of an AI booking or phone tool.

Data entry or administrative assistant. You bring accuracy, attention to detail, and knowledge of business systems and processes. Next doors include automation or AI operations assistant, data-quality checker for AI systems, operations coordinator, and bookkeeping that leads towards finance operations. Learn advanced spreadsheet skills, no-code automation tools, and how to check and correct AI output systematically. First step: automate one repetitive task you do every week.

Executive assistant. You bring trust, discretion, juggling priorities and an understanding of how senior people think. Next doors include chief of staff, operations manager, or running a whole set of AI tools for a leader or team. Learn project management, business basics such as budgets and strategy, and how to direct AI agents. First step: ask your manager for one strategic project to own, not just support.

Finance

Bookkeeper. You understand how money moves through small businesses, and your clients trust you. Next doors include part-time finance director or business adviser for small businesses, specialist in AI-driven finance systems, and cash-flow coach. Learn management accounting and forecasting, advisory and communication skills, and the AI features of the main accounting platforms. First step: offer three existing clients a monthly half-hour "numbers conversation" explaining what their figures mean.

Junior accountant or accounts clerk. You have technical knowledge and know the systems. Next doors include tax advisory, forensic accounting and fraud investigation (a growth area; see Chapter 16), financial systems implementation, and auditing AI-driven processes. Learn a specialism, data analysis, and fraud and risk. First step: pick one specialism and read its professional body's entry requirements.

Insurance claims or loan processor. You understand rules, risk and people's situations. Next doors include complex claims handler, fraud investigator, hardship or customer advocate, and reviewer of AI decisions. Learn investigation skills, compliance and regulation, and how to support vulnerable customers. First step: volunteer for complex or disputed cases.

Customer-facing work

Call centre or customer service agent. You bring empathy, patience, problem-solving, product knowledge and the ability to handle difficult people. Next doors include customer success manager, complex-case specialist, AI conversation designer or trainer (helping improve how AI assistants respond), community manager, and sales. Learn how AI chat systems are set up and improved, customer success practices and relationship management. First step: keep a log of the questions the AI or the script handles badly, and why, and offer it to your team leader. That's the beginning of an AI trainer's job.

Retail worker. You're good at talking to strangers, you know products, you can handle pressure, and you understand stock and display. Next doors include helping local businesses set up simple AI tools, e-commerce assistant, customer experience roles, and store management. Learn online selling platforms, social media for small business and basic AI business tools. First step: help one local business set up its Google Business profile and better product descriptions, for free, as a portfolio piece.

Salesperson in high-volume selling. You bring resilience, persuasion and the ability to read people. Next doors include consultative sales of complex products, account management and partnerships. Learn the language and economics of one industry deeply, and consultative selling. First step: choose that industry and start learning how it works.

Creative and technical work

Copywriter or content writer. Writing, clarity, persuasion and research are your strengths. Next doors include content strategist, editor, brand voice lead, specialist writer in a field you know well, ghostwriter for experts, and AI content quality lead. Learn strategy, a specialism, and how to build your own audience. First step: pick your specialism and publish one piece a week in it.

Graphic designer. You have visual taste, tools and client experience. Next doors include art director, brand designer, user-experience designer and AI creative director. Learn brand strategy, user-experience basics and how to direct AI image tools. First step: redo an old project as an art director, letting AI generate variations while you choose and refine, and document the process.

Junior developer. You have the technical fundamentals. Next doors include building products on top of AI models, developing software within a specific industry, solutions engineering, security, and automation for businesses. Learn AI coding agents, how to build with AI models, one industry in depth, and security basics. First step: build and publish one genuinely useful tool for a real person you know.

Health and education

Medical administration, billing or coding. You know how healthcare systems work and you're accurate. Next doors include patient navigator or care coordinator, clinical information systems support, practice manager, and clinical support roles with training. Learn patient-facing skills, healthcare IT systems, and short clinical or care qualifications. First step: look up care coordinator roles near you and what they require.

Teacher considering a change. You're good at explaining, managing groups, planning, patience and reading people. Next doors include learning designer, corporate trainer (especially AI training, which almost every organisation needs), AI education specialist, and roles in education products. Learn instructional design, AI tools for learning and workplace training

methods. First step: build one short AI-for-beginners workshop and run it for friends or a local group.

Tutor. You bring one-to-one teaching and motivation. Next doors include AI-supported tutoring that serves more students, learning coach focused on accountability and study skills, specialist tutoring for students with particular needs, and teaching families to use AI well. First step: redesign your offer around what AI can't do: accountability, motivation and confidence.

Trades and transport

Long-haul truck driver. You know logistics and roads, you're reliable and you think about safety. Next doors include local and specialist driving, fleet and logistics coordination, autonomous-truck terminal operations, heavy vehicle and EV mechanics, heavy equipment operation and dispatch. Learn logistics software, heavy vehicle or EV mechanics, and extra licences such as dangerous goods or heavy machinery. First step: price one additional licence or course this month.

Taxi or rideshare driver. You have local knowledge and customer service skills. Next doors include patient and disability transport, specialist delivery, a trade apprenticeship and logistics. First step: talk to one person doing each of those jobs about how they got in.

Warehouse worker. You're used to physical work, logistics and safety. Next doors include robotics or automation technician, maintenance, forklift and heavy equipment, and logistics coordination. Learn mechatronics or maintenance and automation basics. First step: ask your employer whether it funds technician training; many automated warehouses do.

Factory production worker. You bring process discipline, safety and technical familiarity. Next doors include maintenance technician, an electrical apprenticeship, quality control and robot operation. First step: talk to your site's maintenance team about how they got their roles.

Already in a trade. The growth areas are electrical work for data centres, solar, batteries and EV charging; cooling for data centres; smart-home installation; training and assessing apprentices; and running a bigger business with AI handling the admin. First step: book one specialisation course.

PROOF OVER PAPER

One principle applies to every door on this map.

In a world where AI can write a flawless résumé, a perfect cover letter and a polished online profile for anyone in thirty seconds, claims on paper are losing their value. What's gaining value is proof: a portfolio of real work, a project people can see, a result you can point to ("I set up this system and it saved the team six hours a week"), a reference from someone the employer trusts, a small business with real customers, a published article, a working app, a before-and-after.

Certificates still matter for some things, especially licensed professions and trades. But for most career changes, one real project is worth more than a handful of online certificates. So when you're choosing your next door, don't only ask what course you should do. Ask what you can build, fix, help with or produce in the next three months that proves you can do the job.

THREE QUESTIONS BEFORE YOU SPEND MONEY ON TRAINING

Before you pay for any course, ask whether people in the job you want actually list that qualification in their job ads (check ten real listings). Ask whether you can learn the first 20 per cent for free to find out if you even like it, because usually you can. And ask what you'll be able to show at the end. If the answer is "a certificate" and nothing else, be cautious.

DO THIS NOW

1. Find your row on the map and circle two next doors that interest you.
2. Look up ten real job ads for one of those doors, in your area, and write down the skills that appear most often.
3. Design a proof project: something you could do in the next ninety days that would prove you can do that job. Write it in one sentence.
4. Find one person already doing that job and ask for fifteen minutes of their time. Most people say yes.

CHAPTER 12

The Income Tripod

Let's talk about the thing that actually keeps people awake.

Not whether AI will become conscious. Not whether robots will take over. *Can I pay the rent if something happens to my job?*

That's the real fear under most of the AI anxiety out there. It isn't abstract. It's the mortgage, the kids' school shoes, the car loan, the groceries that cost more every month. It's the knot in your stomach when you check your balance two days before payday.

So this chapter isn't really about AI at all. It's about making your household strong enough that whatever AI does, slowly, steadily or fast, you have *time* and *choices*. That's what financial resilience really is. Not being rich. Having time and choices when things change.

A necessary word before we start. I'm not a financial adviser, and this chapter is general education, not personal advice. Your situation is yours. For big decisions about debt, investing, tax or retirement, talk to a qualified, independent professional or a free financial counselling service.

And a second word, which matters more. If you're living week to week right now, I'm not going to lecture you about "just saving more". A lot of families are doing everything right and still have nothing left at the end of the fortnight. This chapter is built as a ladder, not a test. Wherever you're standing, there's a next rung, and the next rung is all you need to think about.

WHY A STOOL NEEDS THREE LEGS

Picture a stool with one leg. You can balance on it if you're careful and nothing moves. Most households live like that: one main income, or two incomes from similar kinds of work, and not much underneath. Knock that leg and you're on the floor.

Now picture a stool with three legs. Knock one and it wobbles, but it stands, and you've got time to fix the broken leg.

The Income Tripod is the idea that every household, over the next five years, should be working towards three legs. Your main income, made as secure as you can make it, which is what most of this book is about. A second, skill-based income from a different source than your main job, even a small one (Chapter 13). And a cash buffer, money set aside that buys you time if either of the other legs breaks.

You don't need all three this month. Most people won't have all three solid for a year or two. But knowing that this is the shape you're building towards changes every money decision you make. Start with the leg that matters most and that most people skip.

THE CASH BUFFER

When a job disappears, the damage isn't only the lost income. It's the *speed*. With no savings, you have to take the first thing that comes along, even if it's worse paid, a worse fit, or another job that's about to be automated. You can't afford to retrain. You can't afford to wait for the right thing. Panic makes the decisions for you.

A buffer changes that. Even a modest one means that if the worst happens you can breathe, look properly, maybe do a short course, and make a good decision rather than a desperate one. In every one of the three scenarios from Chapter 5, a buffer helps. In the Fast Break, it could be the most important thing you did.

Most financial educators suggest working towards three to six months of *essential* expenses, meaning rent or mortgage, food, power, insurance, minimum debt repayments and transport, not your whole income. If that num-

ber looks impossible right now, you're far from alone. So here's the ladder. The first rung is \$500 to \$1,000, enough to stop a car repair or a dental bill turning into debt. The second is one month of essential expenses. The third is three months. The fourth is six months, especially if you're in a high-exposure job, self-employed, or the only earner. Each rung makes you meaningfully safer than the one below. Don't be discouraged by the top rung. Climb the next one.

How do you actually build it when there's not much spare? Automate it, even if it's twenty dollars a pay. Set up an automatic transfer to a separate account on payday, because what you don't see, you don't spend. Keep it separate and a little bit awkward to get at, perhaps a high-interest savings account at a different bank from your everyday account, so it's there in an emergency but not one tap away from a takeaway order. Feed it windfalls: a slice of any tax refund, bonus or birthday money goes straight in until you hit your next rung. And don't invest your emergency buffer in things that can crash. Shares and crypto can fall at exactly the moment the economy, and your job, are under pressure.

THE VULNERABILITY AUDIT

Before you can strengthen your household, you need to know where it's weak. Tick each statement that's true for you. Be honest; nobody's marking it.

- ☐ More than 80 per cent of our household income comes from one job.
- ☐ We have two incomes, but both are in high-exposure jobs (an Automate share above 40 per cent on the Task Test).
- ☐ My main income is in an industry that's already shrinking.
- ☐ I have no income source outside my main job.
- ☐ Our fixed costs (rent or mortgage, loans, subscriptions) take more than 60 per cent of our take-home pay.
- ☐ We have high-interest debt: credit cards, payday loans or buy-now-pay-later.

- ☐ We'd have to sell something important if we lost one income for three months.
- ☐ We have less than one month of expenses saved.
- ☐ Our savings are all in investments that could fall in value.
- ☐ People depend on us and we don't have income protection or life insurance.
- ☐ We haven't checked what government support we'd be entitled to if we lost work.
- ☐ I haven't learned a significant new skill in two years.
- ☐ I don't use AI in my work at all.

Count your ticks. Three or fewer means a solid foundation, so focus on building the second leg. Four to seven means some real weak spots; pick the top two and start there. Eight or more means you're exposed. Not doomed, exposed. Start with the buffer and any high-interest debt, and work through this chapter one step at a time.

FIX THE LEAKS: HIGH-INTEREST DEBT

If you've got credit card debt, payday loans or buy-now-pay-later balances charging high interest or fees, that debt is quietly working against you every month. In uncertain times it's dangerous for a simple reason: if your income drops, the debt doesn't. It keeps growing.

A common approach that financial educators recommend is to build a small starter buffer first, so an emergency doesn't create new debt; then throw every spare dollar at the highest-interest debt while paying the minimum on the rest; and once that's gone, roll the same repayment onto the next one.

If debt feels overwhelming, talk to a free, non-profit financial counselling service. In Australia, that's the National Debt Helpline on 1800 007 007. Most countries have an equivalent. They're genuinely helpful, it costs nothing, and they won't judge you.

LOWER YOUR FIXED COSTS

People focus on earning more. But in uncertain times, lower fixed costs are a form of freedom. Every dollar of regular monthly costs you remove is a dollar you never have to earn again. It lowers the income you need to get by, which means more jobs become viable, more career changes become possible, and the same buffer covers more months.

Go through your bank statement line by line and ask of every regular payment: would I sign up for this again today? AI can help. Export three months of transactions with the account numbers removed, ask an AI to sort them into categories and find the recurring charges, and you may well find subscriptions you'd forgotten about.

IF UNEMPLOYMENT RISES: THE HOUSEHOLD DOWNTURN PLAN

This is the plan you write now, while things are calm, so you're not writing it in a panic. Sit down with your partner, or on your own, and answer six questions on one page.

What are the early warning signs for our household? A hiring freeze at work, a restructure, colleagues not being replaced, bad news in your industry, a high result on your Task Test.

What would we cut first if income dropped? Decide now, so it's not an argument later.

What would we protect? Housing, food, the kids' essentials and insurance, in roughly that order.

Who would we call? People who might know of work, recruiters, former colleagues, your wider network.

What are we entitled to? Look up now what unemployment payments, redundancy entitlements and hardship help exist where you live, and what your employment contract and any income protection insurance actually say. In Australia, start with Services Australia and your award or agreement; elsewhere, your national government's employment pages.

And what would we do in the first week? Update your résumé and portfolio. Apply for any support straight away. Contact your lender early if

repayments might be a struggle, because most have hardship arrangements and talking early helps. And start the 30-Day Challenge in Chapter 15 if you haven't already.

Write it down, put it somewhere you'll find it, and hope you never need it.

WHAT ABOUT A UNIVERSAL BASIC INCOME?

You'll hear a lot about universal basic income, the idea that governments could give everyone a regular payment with no strings attached, especially if AI takes a lot of jobs. It's worth knowing what the evidence says so you can think about it clearly.

The largest recent study in the United States was run by OpenResearch, a non-profit backed by Sam Altman, the head of OpenAI. For three years, 1,000 lower-income people received US\$1,000 a month, no strings attached, while a comparison group of 2,000 received US\$50 a month. Recipients worked slightly less on average, but the researchers found this was mostly about having more choice, such as being pickier about jobs or spending time with family, rather than not wanting to work. They were more likely to report business ideas and plans to start one. Stress and food insecurity improved noticeably, but mainly in the first year, and those improvements faded. There were no significant effects on physical health. Parents spent more on their children.

So: not a miracle, and not a disaster. A modest, meaningful help for people on low incomes.

PLAUSIBLE SCENARIO Some countries introduce new forms of income support in response to AI disruption by 2031. Serious people from across politics have argued for versions of it, others argue strongly against it on cost and other grounds, and alternatives such as stronger unemployment insurance, retraining accounts and shorter working weeks are also being debated.

My honest advice is not to plan your family's future around it. If support arrives, good. It will be a floor, not a lifestyle, and it's likely to arrive after the pain rather than before. Build your own tripod first.

A WORD ON INVESTING, BRIEFLY AND CAREFULLY

People ask whether they should invest in AI companies. I'm not going to tell you what to buy. But a few principles that financial educators broadly agree on matter especially now.

Your buffer comes first; don't invest money you might need in the next few years. Don't put all your eggs in the same basket as your job: if you work in tech and all your investments are tech shares, one bad year could hit your job and your savings at once. Be very wary of hype, because the dot-com bubble made some people rich and wiped out many others, even though the internet itself turned out to be huge. Be extremely wary of "AI trading bots" and investment schemes: in the United States, investment scams were the single biggest category of reported fraud losses in 2025, and many use AI buzzwords (Chapter 16). And get independent advice for big decisions, from someone who isn't paid commission on what they sell you.

DO THIS NOW

1. Do the Vulnerability Audit. Count your ticks and circle your top two weak spots.
2. Set up an automatic transfer to a separate savings account on your next payday. Any amount counts.
3. Find one fixed cost to cut this week, and cut it.
4. Write your one-page household downturn plan.
5. Look up what government support and redundancy rights apply to you, and save the links.

Earning With AI: From \$0, \$100 and \$1,000

This is the chapter most likely to be full of rubbish in other AI books, so I'll make you a promise up front. There is no secret AI money machine.

There's no prompt that prints cash. There's no "passive income system" that makes ten grand a month while you sleep after setting it up over a weekend. If someone is selling you one, they're making their money from *you*, not from the system.

What there is, right now, is a real and unusual opportunity. And it's far more ordinary, and far more achievable, than the hype. Most businesses and most people have no idea how to use AI well. You can learn it faster than they can. And people will pay you to help them.

That's it. It's not glamorous, and it works.

Remember the MIT study from Chapter 1, the one suggesting most corporate AI pilots weren't producing measurable results? That's not just a story about hype. It's a story about a gap. Businesses are spending money on AI and not getting value from it because they don't know how to fit it into the way they actually work. Every gap like that is a job.

One more thing to keep in mind through this whole chapter: AI is leverage. It isn't a business model by itself. Customers don't pay for AI; they pay for useful outcomes. AI just lets you deliver those outcomes faster, cheaper, or for more people.

THREE KINDS OF AI INCOME

Almost every realistic way to earn with AI falls into one of three buckets.

AI-boosted services are services people already pay for, which AI lets you do faster, better or more cheaply, so you keep more of the money or serve more people. Writing, bookkeeping, marketing, design, virtual assistance, tutoring, research, translation with proper human review, or running a cleaning business more efficiently.

AI setup and help means helping businesses or people use AI themselves. You're the bridge. Setting up AI tools for a local business, training staff, building simple automations, teaching older people to use AI safely, or helping families with AI safety (Chapter 18).

Human-premium services are things people will pay more for precisely because a real person does them. As AI makes everything else cheap, these become relatively more valuable (Chapter 21). Personal training, coaching, handmade goods, local tours, in-person classes, trades, care, events and experiences.

There's a fourth bucket too, digital products such as e-books, templates, courses and printables. I'll be honest: it has become very crowded, because AI makes it easy for everyone to produce them. It can work, especially if you have real expertise or an existing audience, but it's rarely the fastest path for a beginner.

STARTING WITH \$0

You have a phone or computer, internet access, free AI tools, and your own time and skills.

Path A: the local business AI helper. Help small local businesses (cafés, hairdressers, tradies, gyms, shops) with the online work they hate: their Google Business profile, replies to reviews, social media posts, website wording, simple AI tools. It works because small business owners are time-poor and often intimidated by technology. They don't want to learn AI; they want someone they trust to sort it out. Pick three businesses you already use and like. Do something useful for one of them for free (rewrite their Google Business description, draft three weeks of social posts, or write polished replies to their last ten reviews). Show them, ask if it's useful, and ask whether they'd pay you to keep doing it. Use that as your first example and testimonial. People who do this often start with small monthly retainers and

grow as they add clients and services. Rates vary a lot by location, so check what local marketing freelancers charge. All you need is basic AI skill (Chapter 14), friendliness and reliability.

Path B: AI-boosted freelancing in a skill you already have. Take something you're already good at, such as writing, admin, organising, research, spreadsheets or customer service, and offer it as a freelance service, using AI to deliver it faster. Start with people you know, then local businesses, then online platforms. Be honest that you use AI tools; you're selling results, reliability and judgement. One warning: online freelance platforms have been flooded with cheap AI-generated work, and competing on price there is a race to the bottom. Compete on reliability, specialism and relationships instead.

Path C: teach what you've just learned. Once you've done the 30-Day Challenge in Chapter 15, you'll know more about practical AI than most people around you. Run a small paid workshop at a community centre or library, for a local business group, or for retirees. There's huge demand for patient, human, jargon-free AI teaching, especially among older people and small business owners. Stanford's 2026 AI Index ranked the United States only 24th in the world for everyday use of generative AI, at about 28 per cent of the population on its measure, so plenty of people still haven't really got started. Run your first session free for friends and family, collect feedback, then offer it to a community group for a small fee per person.

STARTING WITH \$100

You have a hundred dollars to spend on tools, a simple website or a little advertising.

Path D: a productised AI service. Package one specific service with a clear price. Not "I do marketing" but "I'll set up your Google Business profile and write a month of posts for a fixed fee", or "I'll turn your voice memos into professional quote and invoice templates". Spend the money on a simple one-page website or booking page, a paid AI subscription so you're using the best tools, and a few flyers or a small local ad. It works because a clear offer with a clear price is far easier to buy than a vague one.

Path E: a niche digital product, done properly. A useful template, guide or toolkit for a specific group you understand: AI prompts and templates for real estate agents, a quoting kit for plumbers, a meal-planning system for busy parents of a child with allergies. The honest bit is that this only works if you really understand the group's problem and you can reach them, through your network, a community, social media or an existing audience. Making the product is the easy part now. Getting it in front of people is the hard part.

STARTING WITH \$1,000

You have a thousand dollars to invest and you're prepared to put in real hours.

Path F: AI implementation for small businesses. A step up from Path A. You help businesses set AI up properly in their operations: an assistant on their website that answers questions and books appointments, automations connecting their email, calendar and accounting, staff training and ongoing support. Spend the money on learning the automation tools you'll use (many good courses are cheap), paid tools for testing, a professional website and basic business setup such as insurance and registration. It works because businesses pay real money for someone who saves them staff time. If your setup saves a business ten hours a week, the maths is easy for them. You don't need to be a programmer; you need to be the person who works things out and explains them clearly.

Path G: a small AI-powered business in an old industry. Start a traditional service business (cleaning, lawn mowing, handyman work, pet care, bookkeeping, tutoring) and run it with AI from day one for marketing, bookings, customer messages, quotes, scheduling and admin. You compete on what customers care about, good reliable work, while spending far less time and money on admin than older competitors. The work itself is anchored: human, physical and local. The business side is AI-boosted.

STARTING WITH PROFESSIONAL SKILLS

If you're an accountant, lawyer, nurse, teacher, engineer, HR professional, marketer or any other experienced professional, two paths stand out.

Path H: the fractional expert. Offer your expertise part-time to several small businesses that can't afford a full-time person. AI handles much of the grunt work, so you can serve more clients at a lower price. A bookkeeper becomes the part-time finance director for six small businesses; an HR professional becomes the part-time people adviser for five growing companies.

Path I: the AI bridge in your profession. Your profession needs people who understand both the professional work and AI. Consult, train, write guides, or help firms in your field adopt AI safely and legally: a nurse who understands AI documentation tools, a lawyer who helps small firms adopt AI research tools without getting burned, a teacher who trains other teachers.

STARTING WITH TRADE SKILLS

Path J: run a tighter trade business. As Chapter 9 explained, the biggest near-term financial win for most tradies isn't a new income stream. It's getting ten hours of admin back each week and spending it on paid work.

Path K: specialise into the boom. Electrical, cooling, solar and batteries, EV chargers and data-centre work. Specialisation lifts your rate.

Path L: teach your trade. As people move into trades from other work, the demand for training grows. Mentoring, assessing apprentices or making practical how-to content can become a second income.

WITHOUT A DEGREE

Everything in this chapter is open to you. None of it needs a degree. It needs you to learn the tools (Chapters 14 and 15, which cost nothing), to be reliable (turn up, do what you said, when you said), to build proof (one real example of your work, then another), and to find the courage to ask the first person whether they'd pay.

In a world where AI can make anyone *sound* qualified, the people who stand out are the ones who've actually done the thing. That's good news if you don't have letters after your name.

HOW TO SPOT THE AI MONEY SCAMS

This area is crawling with them, so here's your checklist. Be very suspicious of anyone who:

- ☐ promises specific large earnings, such as "\$10k a month with ChatGPT";
- ☐ shows screenshots of income as proof;
- ☐ says it's "passive" or "automated" and needs no skill or effort;
- ☐ uses pressure ("only three spots left", "price goes up tonight");
- ☐ wants you to recruit other people;
- ☐ offers an "AI trading bot" with guaranteed returns;
- ☐ asks you to pay in cryptocurrency or gift cards;
- ☐ won't explain exactly what you'll be doing to earn money.

Three or more ticks and you should walk away.

The honest truth about earning with AI is that it looks a lot like earning without AI. You find people with a problem, you solve it well, and they pay you. AI just lets you solve it faster, better, or for more people.

WHICH PATH SHOULD YOU CHOOSE?

If all these options leave you paralysed, use one rule: start with the path closest to what you already know and who you already know. The quickest first dollar almost always comes from someone who already trusts you, for help with something you're already good at, done faster with AI.

DO THIS NOW

1. Pick one path from this chapter. Just one. Write it down.
2. List five people or businesses you already know who might have that problem.
3. Do one free piece of work for one of them this week and ask for honest feedback.
4. Write your one-sentence offer: “I help [who] with [what] so they can [result].”

CHAPTER 14

How to Actually Use AI

Most people who “tried AI and didn’t see what the fuss was about” made the same mistake. They treated it like a search engine. They typed four words, “write marketing email”, got back something bland and generic, and thought: is that it?

That’s like hiring the most well-read intern in the world, shouting three words at them from across the room, and being disappointed they couldn’t read your mind.

This chapter is the practical how-to. You don’t need to be technical. You need to learn to brief well, check well and build habits. By the end, you’ll be using AI better than plenty of people who use it every day.

BRIEFING AI: THE CRAFT METHOD

The biggest single improvement most people can make is to give AI more to work with. Here’s an easy way to remember what to include.

Context: what’s the situation, who are you, what’s the background?
Role: who should the AI act as, such as “an experienced bookkeeper” or “a patient maths tutor for a twelve-year-old”? **Ask:** what exactly do you want?
Format: how do you want it, whether a list, an email, a table, under two hundred words, or plain English? **Test:** ask it to check its own work, point out weaknesses, or ask you questions before it starts.

Here’s a lazy prompt: *Write a marketing email.*

And here’s the same request with CRAFT:

Context: I run a small family bakery in a country town. We're launching gluten-free bread on Saturdays from next month. Most of our customers are over fifty and have been coming for years. They like that we're local and personal. Role: Act as an experienced small-business copywriter who writes warm, plain, friendly emails, not salesy or corporate. Ask: Write an email announcing the gluten-free bread to our mailing list. Format: Under 150 words, a friendly subject line, easy on the exclamation marks, signed "Maria and the team". Test: Before you write it, ask me any questions you need answered to make it better.

The second version takes a minute longer to write. The result will be dramatically better.

A few more habits make a huge difference. Have a conversation, not a transaction: the first answer is a draft, so reply with “make it warmer”, “shorter”, “that second paragraph doesn’t sound like me, here’s how I’d say it”, or “give me three completely different versions”. Show it examples by pasting in something you’ve written and asking for that style, or a report you liked and asking for that structure. Ask it to interview you (“I need to write a business plan; ask me questions one at a time until you have what you need”), which is one of the most powerful and underused techniques there is. Ask for options and then choose, because your taste is the value. And ask it to critique: “What’s weak about this? What would a sceptical customer say? What have I missed?” AI is often better at criticising than creating.

RESEARCH AND FACT-CHECKING: THE TWO-SOURCE RULE

AI is a brilliant starting point for research and a dangerous finishing point. It can be confidently wrong (Chapter 2), and it has a tendency, which its makers acknowledge and are working on, to tell people what they seem to want to hear. So here’s the rule I use. If it matters, don’t trust it until you’ve confirmed it from two independent, reliable sources that aren’t the AI. “Matters” means money, health, the law, safety, anything you’ll publish, anything you’ll put your name to, and anything you’ll tell your kids is true.

Use AI to map the territory (“explain the main options for this problem, the pros and cons of each, and what questions I should be asking”). Ask it for sources and then check that they exist and say what the AI claims. Many tools can now search the web and link to their sources, so click the links and read them; AI has been known to cite real sources that don’t say what it claims, and sources that don’t exist at all. For important facts, go to the primary source: the government website, the official statistics, the actual study, the actual law, the company’s own announcement. Ask the AI to argue the other side, so you can see where you’re only hearing one view. And be extra careful with numbers, dates, names, quotes and legal or medical specifics, because that’s where errors hide.

WHAT NEVER GOES INTO AN AI

Before going further, privacy. Treat anything you type into an AI tool as something that might be stored and potentially seen by others, unless you’ve checked the settings and the company’s policies carefully.

Don’t put in passwords, PINs or security answers; full bank or card details; government identification numbers such as your tax file number or passport number; other people’s private information without their permission, especially customers’, patients’ or clients’; confidential work information, unless your employer has approved that specific tool for it; or anything you’d be devastated to see leaked.

Do check your AI tool’s privacy settings, including whether your conversations are used to train future models, and switch that off if you’d prefer. Use your employer’s approved tools for work. And strip out names and identifying details before asking for help: “Client A”, “my colleague”.

LEARNING WITH AI

AI can be the best tutor you’ve ever had, if you use it to make you think rather than to think for you. The AI tutors that worked so well in the Nigeria and Harvard studies (Chapter 8) were designed to make students work things out. You can set that up yourself:

I want to learn [topic]. Don't give me the answers. Teach me step by step, ask me questions to check I understand, and only move on when I get it right. Start by asking what I already know.

Other approaches work well too. Ask it to explain something as if you were twelve, then as if you were a university student, and you get two levels of understanding in two minutes. After reading anything, ask it to quiz you. Ask what the five most common misunderstandings about a topic are. Ask for a thirty-day learning plan for a new skill, twenty minutes a day, using free resources. And practise conversations: have it play a tough interviewer for the job you want, or an angry customer for you to handle, or a waiter so you can practise ordering in Spanish.

EVERYDAY USES

A few of the most useful everyday uses, to get you started. Paste in a long email thread and ask it to summarise what's being asked of you and draft a reply. With everyone's permission, record a meeting and have AI write up the notes and action points. Give it everything you need to do this week and ask it to help you prioritise and schedule it. Ask it to help you think through a decision by asking you questions. Get help with a difficult message: "Help me write a kind but firm message about this situation." Have it summarise a contract in plain English and list anything you should worry about, then check anything important with a professional. Give it three months of spending categories, with account details removed, and ask where you could save. And use it for family admin: meal plans on a budget, school holiday schedules, birthday party plans.

DATA AND SPREADSHEETS

You don't need to be a spreadsheet wizard any more. Upload a spreadsheet without sensitive data and ask for the three most interesting things in it. Ask for a chart of sales by month. Ask for a formula that does a particular job, then paste it in. Ask it to find rows where something looks wrong. Always

sense-check the answers, because AI can miscalculate or misread data. If a number looks surprising, check it by hand.

AUTOMATION: MAKING THINGS HAPPEN WITHOUT YOU

This is where you move from “AI helps me” to “things happen automatically”. No-code automation tools let you connect apps together with simple rules (when this happens, do that), and many now include AI steps. There are several popular ones; search for “no-code automation tools” and read some current reviews.

Here are some anyone can set up. When a customer fills in the form on your website, AI drafts a personalised reply and sends it to you to approve. When an invoice arrives by email, the attachment is saved to a folder and a line added to your spreadsheet. When someone books an appointment, they get a reminder the day before. Every Monday, AI summarises last week’s sales and emails them to you. Start with one small automation that saves you ten minutes a week, then another. They add up.

AGENTS

As Chapter 4 explained, agents are AI systems that can take a goal and carry out several steps on their own: browsing the web, filling in forms, using apps, writing and running code. In 2026 most of the major AI tools offer some kind of agent or “do it for me” mode. They’re powerful and improving fast, and they also make mistakes.

So use them carefully. Start with low-stakes tasks such as research, drafting and organising, not paying bills or emailing important clients. Watch what the agent does the first few times. Don’t give it access to money, passwords or important accounts without careful thought and strong safeguards. Require your approval before anything goes out into the world, whether that’s an email being sent, a purchase being made or a post being published. And check the results the way you’d check a new employee’s first week of work. A good first test: “Research the five best-reviewed [service] providers near me, compare their prices and reviews, and give me a table.”

BUILDING SIMPLE THINGS WITH AI (YES, YOU)

This is the part that surprises people most. You can now build simple apps, websites and tools without knowing how to code, by describing what you want to an AI. In early 2025 the computer scientist Andrej Karpathy coined a name for it, “vibe coding”: you describe the vibe and the AI writes the code. Try this:

Build me a simple web page I can open in my browser that works as a [chore chart for my kids / tip calculator / quote calculator for my cleaning business / habit tracker]. Make it colourful and easy to use on a phone. Give me the complete code in one file and tell me exactly how to open it.

Many AI tools will now build it and show it to you directly. Will it be perfect? No. Will you be amazed you built something? Very likely. Just remember that anything handling payments, passwords, personal data or anything important needs a real developer to check it, because AI-written code can contain security holes.

THE HABITS OF PEOPLE WHO ARE REALLY GOOD AT THIS

Underneath all of the above, the people who are genuinely good with AI share a handful of habits. They use it every day for small things, because fluency comes from repetition. They brief it well, with context. They iterate, treating the first answer as a draft. They check what matters, every time. They know where it’s weak in their own field. They share what they learn with other people. And they keep thinking. They use AI to go further, never to stop.

DO THIS NOW

1. Rewrite one prompt you actually need this week using CRAFT, and compare the result with a lazy version.
2. Check the privacy settings in the AI tool you use most.
3. Ask AI to interview you about a decision or project you're stuck on.
4. Build one tiny thing: a web page, a checklist tool or a calculator.

The 30-Day AI Reskill Challenge

Reading about AI doesn't make you good at AI. Using it does.

So here's the deal. Thirty days, one task a day. Most take fifteen to thirty minutes and a few take a bit longer. You don't need any technical skill to start, and you don't need to pay for anything, although a paid subscription to one AI tool will make some days easier. By Day 30 you'll be more capable with AI than most people you know, probably including some who've been using it for years.

A few tips before you start. Keep a simple log, a notebook or a note on your phone: the date, what you did, and one thing you learned. If you miss a day, don't start again from scratch; just pick up where you left off. Do it with someone if you can, whether that's your partner, a mate, a workmate or your teenager, because it's more fun and you're far more likely to finish. And use real tasks from your real life wherever possible. That's what makes it stick.

WEEK ONE: FOUNDATIONS

The goal this week is to get comfortable, get curious, and learn where AI is strong and where it's weak.

Day 1: Say hello properly. Open a free AI assistant and tell it about yourself: your work, your family situation (without private details), and what you're hoping to get out of AI. Then ask, "Based on what I've told you, what are ten ways you could help me in my daily life?" Try one of them.

Day 2: The CRAFT test. Pick one real writing task, such as an email, a message or a post. Do it first with a lazy four-word prompt, then again with CRAFT (Chapter 14). Compare the results and write down the difference.

Day 3: Catch it out. Ask detailed questions about something you're an expert in (your job, your town, your hobby) and keep going until you find a mistake. Write down what it got wrong and how confident it sounded.

Day 4: Make it teach you. Pick something you've always wanted to understand, such as how mortgages work, how the share market works, how to read a nutrition label, or how electricity gets to your house. Use the tutor prompt from Chapter 14: "Don't give me the answers. Teach me step by step and quiz me."

Day 5: Show it the world. Use the camera. Photograph a plant, a form, a recipe, a broken appliance, a sign in another language or a document, and ask it to explain, translate or help.

Day 6: Talk to it. If your tool has a voice mode, have a spoken conversation while you walk or cook. Plan your week out loud.

Day 7: Privacy and reflection. Check your privacy settings. Then read back over the week's log. What surprised you? Where was it useful, and where was it useless? Write three sentences.

WEEK TWO: AI AT WORK

The goal this week is to make AI useful in your actual job.

Day 8: The Task Test, with help. Redo your Task Test from Chapter 6, but this time describe your job to the AI and ask it to help you list your tasks. Compare its list with yours. What did it add, and what did it miss?

Day 9: Speed up an Assist task. Take one Assist task from your list and do it with AI from start to finish, including checking. Time yourself.

Day 10: Summarise something long. Take a long document, email thread, report or article relevant to your work (nothing confidential unless you're using an approved work tool) and ask for a summary, the key points, and "what should I do about this?"

Day 11: Rehearse a conversation. You've got a meeting, a customer call, a hard conversation or an interview coming up. Ask the AI to play the other person and practise twice.

Day 12: Research properly. Research a real question for work or home with AI, then apply the Two-Source Rule. Check every important claim against two reliable, non-AI sources. How many held up?

Day 13: Spreadsheet day. Take a spreadsheet without sensitive data, such as a budget or a sales list. Ask the AI to analyse it, find patterns, make a chart and write a formula you need.

Day 14: Teach someone. Show a colleague, friend or family member one thing you've learned so far, and watch their face. Teaching locks it in.

WEEK THREE: BUILD

The goal this week is to move from using AI to making things with it.

Day 15: Your first automation. Pick one small repetitive task. Using a no-code automation tool (Chapter 14), set up one simple rule, such as saving email attachments to a folder or sending yourself a weekly reminder. Or simply write down, step by step, what you'd want automated, and ask the AI to help you set it up.

Day 16: Build a tiny app. Ask AI to build a simple web-page tool: a chore chart, a budget calculator, a quote calculator, a habit tracker. Open it, use it, and ask for one improvement.

Day 17: Make something visual. Create an image, a poster, a simple logo or a short presentation with AI. Then critique it. What's good? What's generic? What would a real designer do differently?

Day 18: Try an agent. Give an AI agent a multi-step, low-stakes research task: "Find the five best-reviewed [service] near me, compare prices and reviews, and give me a table." Watch how it works and note where it struggles.

Day 19: Write something that matters. Use AI as an interviewer to help you write something longer and meaningful, such as a one-page business idea, a piece of family history, a letter to your kids or a proposal for your boss. Let it ask the questions; you provide the substance.

Day 20: Set up your personal AI. Many AI tools let you save custom instructions or information about yourself so you don't have to repeat it. Write yours: who you are, what you do, and how you like answers (short or detailed, plain English). Save it.

Day 21: Weekly review. Read back through your log and update your Task Test. What would you mark differently now? Which Assist tasks are becoming easy?

WEEK FOUR: PROTECT, EARN AND TEACH

The goal this week is to turn your new skills into security and opportunity.

Day 22: Scam-proof your family. Set up a family safe word (Chapter 16). Call or message the people you'll share it with and explain why in two sentences.

Day 23: See a fake. Look up examples of AI-generated images and voices from reputable news coverage of deepfakes, and try to spot them. Then accept the lesson: often you can't tell just by looking. That's why the safe word matters.

Day 24: Your one-sentence offer. Pick one earning path from Chapter 13 and write your offer: "I help [who] with [what] so they can [result]." Ask the AI to critique it and help you improve it.

Day 25: Do one free piece of work. For someone you know, such as a friend's business, a relative or a local club, do one useful AI-assisted task for free. Ask for feedback.

Day 26: Update your proof. Update your résumé, online profile or portfolio with your new skills, using specific examples. Not "proficient in AI" but "used AI tools to cut weekly reporting from four hours to one".

Day 27: Talk to your boss or your customers. If you're employed, suggest one way your team could use AI, based on what you've learned. If you're self-employed, tell a few customers about a new or faster service. If you're looking for work, tell your network what you've been learning.

Day 28: Plan the next ninety days. Ask AI to help you build a ninety-day learning plan for the next door you chose in Chapter 11, using mostly free resources.

Day 29: Teach a mini-class. Spend fifteen minutes teaching three people something useful about AI: your parents, your kids, your team or your friends. You'll realise how much you've learned.

Day 30: Score yourself and celebrate. Take the AI Resilience Scorecard in Chapter 24. Look back at Day 1 in your log and notice how far you've come. Then write down your next three moves.

AFTER DAY 30

Keep a small daily habit: use AI for at least one real task every day. Keep your log weekly. Every three months, redo your Task Test and your Scorecard.

The Doubling Clock will keep ticking and the tools will keep changing. But once you've done these thirty days, you won't be starting from scratch each time. You'll have the habits (briefing, checking, iterating, building, protecting), and those carry over to whatever comes next. That's what being AI-native actually means. Not knowing every tool. Knowing how to learn the next one.

DO THIS NOW

1. Put Day 1 in your calendar for tomorrow.
2. Message one person and ask them to do the challenge with you.
3. Start your log: a notebook or a phone note titled "30 Days".

CHAPTER 16

Scams, Voices and Fakes

If you only read one chapter of this book, I'd want it to be this one.

Not because scams are the most important thing AI will do to the world. Because they're the thing most likely to hurt someone you love *this year*, and because the defence is so simple, so cheap and so quick that it would be a tragedy not to set it up.

By the end of this chapter you'll have a family plan that would have stopped most of the scams I'm about to describe. It takes about thirty minutes.

A VIDEO CALL WHERE EVERYONE WAS FAKE

In early 2024, a finance worker in the Hong Kong office of Arup, the British engineering firm behind the structural engineering of the Sydney Opera House, received a message that appeared to come from the company's UK-based chief financial officer. It asked for a confidential transaction.

The employee was suspicious. That's worth saying. They did what you'd hope someone would do.

Then they were invited to a video call. On it were the chief financial officer and several other colleagues, people the employee recognised, whose faces and voices they knew. Every one of them was a fake generated by AI.

Reassured by seeing and hearing people they trusted, the employee made fifteen transfers totalling about HK\$200 million, roughly US\$25 million, to five bank accounts. Arup confirmed it was the victim in May 2024.

I tell this story first for a reason. This wasn't a gullible pensioner clicking a dodgy link. It was a professional at a global company, who was already sus-

picious, and who was overcome by the one thing we've always trusted most: seeing and hearing people we know.

If it can happen to them, it can happen to your mum. It can happen to you.

THE SCALE OF IT

The numbers are staggering. In 2025 Americans reported losing US\$20.9 billion to online crime to the FBI, up 26 per cent in a single year. People aged sixty and over reported losing US\$7.75 billion, up 59 per cent. Cryptocurrency was involved in more than US\$11 billion of the losses, and for the first time the FBI's annual report included a section on complaints involving AI: over 22,000 of them, with losses of nearly US\$900 million.

The US Federal Trade Commission, counting differently, recorded a record US\$15.9 billion in reported fraud losses for 2025, including US\$3.5 billion lost to impostor scams, where someone pretends to be someone they're not.

Australia's figures tell the same story on a smaller scale. Australians reported A\$2.18 billion in scam losses to Scamwatch, banks and other agencies in 2025, up almost 8 per cent on 2024. Investment scams, payment redirection, romance scams, phishing and remote-access scams made up most of the losses. Older Australians, about 17 per cent of the population, accounted for more than a quarter of the money lost.

And these are only the *reported* losses. Most scams are never reported, because people are embarrassed and blame themselves. The FTC has estimated that once under-reporting is taken into account, fraud may have cost American consumers as much as US\$196 billion in 2024.

HAPPENING NOW Online fraud is at record levels and rising fastest among older people. AI is now a recognised factor.

TRAJECTORY AI-enabled fraud keeps growing through 2031.

HOW AI CHANGES THE CON

Scams aren't new. What AI changes is three things.

The first is scale. A scammer used to be able to work a handful of victims at once. AI lets one criminal run thousands of personalised conversations at the same time, on messaging apps, dating sites, email and social media.

The second is quality. The old giveaways, bad spelling, odd grammar and clumsy English, are gone. AI writes perfect, natural messages in any language and any tone.

The third, and the big one, is impersonation. AI can now copy a person's voice from a short recording, and increasingly their face on video. That clip of your grandson on Instagram, your voicemail greeting, the video of you making a speech at a wedding: any of it could be enough raw material.

THE SCAMS TO KNOW ABOUT

The family emergency call. Your phone rings. It's your daughter's voice, crying and panicked. "Mum, I've been in an accident. I'm in trouble. I need money now. Please don't tell Dad." Maybe another voice comes on, a "police officer" or a "lawyer", who explains how to send the money. It sounds exactly like her, because it was built from her voice. In a 2023 survey of more than 7,000 people across seven countries by the security company McAfee, about one in four said they had experienced an AI voice scam or knew someone who had, and most of those targeted said they'd lost money. It's a survey by a company that sells protection, so treat the exact figures with some caution, but police in many countries have issued warnings about exactly this.

The boss or finance-director request. Like Arup: an urgent, confidential message or call from someone senior asking for a payment, or for supplier bank details to be changed.

The investment and romance scam. Someone befriends you online, through social media, a dating app, or even a "wrong number" text. Over weeks or months they build trust. Eventually they introduce an investment opportunity, often in cryptocurrency, often with an "AI trading platform" that shows impressive fake profits. In both the United States and Australia, investment scams are the biggest single source of reported losses.

The fake official. Your bank, the tax office, the police, a delivery company, a government agency, the power company. There's an urgent problem and you must act now, by paying or by "verifying" your details.

Fake customer support. You search online for help with your bank or an app, and the top result, or the number you ring, belongs to a scammer.

Sextortion and fake images. Criminals use AI to create fake sexual images of people, including teenagers, from ordinary photos, and then threaten to share them unless they're paid. In June 2023 the FBI warned of a sharp rise in exactly this. It's one of the most distressing crimes there is, and it has driven young people to harm themselves. Chapter 19 covers how to protect your kids.

THE THREE RED FLAGS

Here's the good news. However sophisticated the technology gets, almost every scam still needs the same three things to work.

Urgency: "Now. Right now. There's no time."

Secrecy: "Don't tell anyone. Don't call anyone."

*An unusual way to pay: gift cards, cryptocurrency, wire transfers,
new bank details, cash handed to a courier.*

When two or more of those show up together, stop. It's almost certainly a scam, however real the voice or the face.

Why these three? Because scammers need to stop you doing the one thing that beats them: checking. Urgency stops you thinking. Secrecy stops you asking anyone. And the unusual payment makes the money almost impossible to get back.

THE FAMILY DEFENCE PLAN

Here's the simple system. It works whether the voice on the phone is real or fake, because it doesn't depend on you being able to tell the difference. Often you can't.

Rule one: the family safe word. Choose a word or short phrase that only your close family knows. It shouldn't be anything on social media or

anything guessable, so not a pet's name, your street or a birthday. Something like "purple wombat", "Uncle Kevin's boat" or "pancake Tuesday". Anyone in the family who calls asking for money or urgent help must be able to say it. If they can't, hang up and call them back on the number you already have for them. Tell everyone: "If I ever ring you in a panic asking for money, ask me for the word. If I can't say it, it's not me, and I promise I won't be offended." Write it down somewhere safe (not in your phone's notes), don't text it, and share it in person or on a call you made.

Rule two: hang up, call back. If anyone at all, whether family, the bank, your boss, the police or the tax office, calls asking for money or personal details, hang up and call them back on a number you already know is genuine. Not the number they give you. Not the number in the text message. The number on the back of your bank card, on the official website you typed in yourself, or saved in your contacts from before. And never trust caller ID on its own: scammers can make a call or text appear to come from your bank's real number, or from someone you know. Real banks, real police and real family won't mind you calling back. Scammers will.

Rule three: the 24-hour money rule. Agree as a family that no unexpected transfer above a set amount, say \$500, happens without waiting 24 hours and talking to one other family member first. Almost no genuine emergency needs money within minutes that can't wait for a phone call to someone else. Nearly every scam does.

Rule four: never pay with those. Gift cards, cryptocurrency, wire transfers to strangers and cash handed to a courier are a scammer's favourite payment methods, because they're nearly impossible to reverse. No genuine government agency, bank or police force will ask you to pay with gift cards. Ever.

Rule five: lock down the basics. Turn on multi-factor authentication, a code from your phone or an authenticator app, for your email, banking and social media. Your email matters most, because whoever controls your email can reset almost everything else. Use a password manager so every account has a different strong password. Set social media accounts to private, and think twice about posting videos with lots of clear speech, especially from children. Ask your bank what protections it offers, such as transfer limits,

alerts, a trusted contact who's told about unusual activity, and delays on new payees. And switch on any scam-call filtering your phone provider offers.

Verification beats fear. You don't need to be suspicious of everything and everyone. You just need a few habits that mean a fake can't get your money even when it fools your eyes and ears.

TALKING TO YOUR PARENTS WITHOUT INSULTING THEM

This is the conversation many adult children dread, because older people can hear “you might get scammed” as “you're getting old and slow”. So don't put it that way. Try something like this:

“Mum, Dad, I've been reading about a scam where criminals copy someone's voice using AI and ring their family pretending to be in trouble. They used fake video to steal twenty-five million dollars from a big engineering company, and the person on the call was a professional who was already suspicious. I'm worried someone might copy my voice and call you. Can we set up a family password, so if I ever ring needing money, you'll know it's really me?”

Notice what that does. It makes *you* the one at risk of being copied and puts them in the role of protector. It shows it can happen to professionals. And it gives them one simple thing to do.

Then, together, choose the safe word, write “Hang up, call back” on a card by the phone with the bank's real number next to it, agree the 24-hour rule, and offer to help them set up multi-factor authentication. Check in every few months. Scams change, and a quick “anything weird happen lately?” goes a long way.

IF IT HAPPENS

If you or someone you love gets scammed, act immediately. Contact your bank straight away, because minutes can matter, and ask them to stop or recall the payment. Report it to the police and to your national reporting service: in Australia that's ReportCyber and Scamwatch, in the United

States the FBI's IC3 and the FTC. Change the passwords on any accounts that might be affected. Beware of "recovery" scams, because criminals often come back to victims a second time promising to get the money back for a fee. In Australia, IDCARE provides free support if your identity has been compromised.

And be kind, to yourself or to the person it happened to. Shame keeps people quiet, and quiet helps criminals. These are professional criminals using advanced technology. Being fooled isn't a character flaw.

WHAT'S COMING

I'd expect the tools scammers use to keep getting better and cheaper: real-time voice conversion, so someone can talk live in another person's voice; real-time fake faces on video; AI systems running scam conversations automatically in dozens of languages. Defences are improving too, with banks using AI to spot unusual transactions, phone networks flagging likely scam calls, and rules requiring AI-generated content to be labelled, such as the EU rules that began applying in August 2026.

But the most powerful defence won't be technology. It'll be habits. A family that has a safe word and a hang-up-and-call-back rule is protected against almost any version of this con, however good the fake gets. That's why I'd ask you to do this tonight.

DO THIS NOW

1. Choose your family safe word and share it, in person or on a call you make.
2. Have the conversation with your parents, or anyone vulnerable you care about, using the script above.
3. Write "HANG UP, CALL BACK" on a card with your bank's real number and stick it by the phone or on the fridge.
4. Turn on multi-factor authentication for your email and your bank.
5. Agree the 24-hour money rule with your household.

Love, Loneliness and Machines

This is the chapter I found hardest to write, because it isn't really about technology. It's about loneliness.

Loneliness is something almost everyone has felt, often privately, often while surrounded by people. The quiet evening after the kids are in bed. The first year after a separation. The months after someone you love dies. Being fifteen and feeling as though no one understands you. Being eighty and realising the phone doesn't ring much any more.

Into that space, very quickly, has arrived something new: an AI that's always there, always patient, always interested in you, and never tired of listening. I want to talk about that honestly, without mocking the people who use it and without pretending it's harmless.

WHAT'S HAPPENING NOW

An *AI companion* is an AI designed to be a friend, partner, confidant or character that you have an ongoing relationship with. Some apps are built specifically for this, and general AI assistants can be used the same way.

It's more common than most parents realise. In 2025 Common Sense Media, an American non-profit that researches children and technology, surveyed a nationally representative sample of 1,060 American teenagers aged thirteen to seventeen. Seventy-two per cent had used an AI companion at least once. Fifty-two per cent were regular users, chatting at least a few times a month. And about one in three had used AI companions for social interaction and relationships, including conversation practice, emotional support, friendship or romantic role-play. In Australia, the eSafety Commissioner reported in 2026 that among children who had used AI companions

or assistants, about one in five talked to them every day for advice on things like health, feelings and life's challenges.

Among adults the numbers are harder to pin down, but companion apps have millions of users worldwide, and many people use general AI assistants as a sounding board for personal problems.

HAPPENING NOW AI companions are widely used, including by most American teenagers at least occasionally.

WHY THEY'RE SO COMPELLING

To understand why AI companions are powerful, think about what they offer. They're always available, at three in the morning, on Christmas Day, in the middle of a panic attack. They never get bored of you; you can talk about the same worry for the hundredth time. They don't judge, so you can say the thing you've never said to anyone. Increasingly they remember what you've told them and ask about it later. And they adapt to you, learning what you like to hear.

For a lonely, anxious or struggling person, that's an extraordinary thing to be offered. It isn't stupid or weak to find it comforting. It would be strange if people *didn't*.

WHERE THEY CAN HELP

I don't want to be one-sided. There are real ways this technology could help people: practising social skills, such as rehearsing a hard conversation, a job interview or asking someone out; getting your thoughts in order before talking to a real person; some support at 3 a.m. while waiting for a real appointment; some conversation and stimulation for isolated older people who would otherwise have very little; and practising a language or exploring ideas. The research on long-term effects is still young, and we should be careful about strong claims in either direction.

WHERE THEY CAN HARM

But there are serious risks, and some have already played out tragically.

The first is dependency. A relationship that's always easy can make real relationships, which are messy, demanding and sometimes disappointing, feel harder by comparison. A person can slowly withdraw from people into the thing that never lets them down.

The second is the business model. Most companion apps are businesses, and their revenue often depends on engagement (how much you use them) and on subscriptions. An AI designed to keep you talking isn't necessarily designed to help you flourish, and those aren't always the same thing.

The third is telling you what you want to hear. AI systems can be sycophantic, agreeing with you, flattering you and reinforcing what you already believe. For someone in a dark place, an AI that goes along with harmful thoughts instead of gently challenging them is dangerous.

The fourth is harm to young people. Several American families sued Character.AI, a popular companion app, after teenagers were harmed, including cases involving suicide. In January 2026 Character.AI and Google agreed to settle five of those lawsuits, on confidential terms. The details of the individual cases are heartbreaking and I won't go into them here, but they're why lawmakers have started to act. Some apps have also allowed, or failed to prevent, sexual role-play with users who were children.

HAPPENING NOW Documented cases of serious harm have led to lawsuits and new laws.

THE RULES ARE STARTING TO ARRIVE

California's Senate Bill 243, which took effect on 1 January 2026, was one of the first laws written specifically for companion chatbots. Among other things, it requires operators to tell minors clearly that they're talking to an AI, remind minors to take a break at least every three hours of continuous use, have protocols for responding to suicidal thoughts or self-harm (including pointing people to crisis services), and prevent sexual content with minors.

In Australia, industry codes that took effect under the Online Safety Act in March 2026 require companion chatbots to protect children from sexually explicit content, either by not generating it or by checking users' ages, and to provide crisis and mental-health information. After an eSafety invest-

igation found several popular companion apps lacked meaningful age checks, at least one withdrew from Australia and another changed its service for under-18s. The EU's transparency rules, which apply from August 2026, also require people to be told when they're dealing with an AI.

These are welcome. They're also minimums. No law can make sure your fourteen-year-old isn't spending four hours a night confiding in an app instead of a person. That part is still up to families.

WHEN THESE GET REALLY GOOD

Looking ahead honestly:

TRAJECTORY By 2031, AI companions are likely to be far more realistic, with natural real-time voices, long memories, video avatars and integration into devices such as glasses and home speakers.

PLAUSIBLE SCENARIO Significant numbers of adults choose AI companions as a main source of emotional connection.

WILD CARD The social effects at scale, on friendship, dating, marriage and birth rates. Nobody can predict these reliably.

The most honest thing to say is that we're running an enormous experiment on human connection and won't know the results for years. That's a good reason for each of us to think carefully about how we use these tools, and to think very carefully about how our children do.

WARNING SIGNS OF UNHEALTHY DEPENDENCY

In yourself or another adult, watch for preferring AI conversation to talking with people you care about, hiding how much you use it, feeling anxious or low when you can't get to it, losing sleep to it, pulling back from friends, family and hobbies, feeling that it understands you better than any person could, or spending money on it that you can't really afford.

In a child or teenager, watch for secrecy about an app or a "friend", long periods alone with a device (especially at night), withdrawal from real friends and activities, strong emotional reactions to the AI (distress when it's unavailable, talking about it as a real relationship), changes in mood,

sleep or schoolwork, and above all any mention of self-harm or of the AI encouraging harmful behaviour.

Two or three of these are a reason for a gentle conversation. Any sign of self-harm is a reason to get help immediately.

KEEPING REAL RELATIONSHIPS STRONG

The best protection against unhealthy AI relationships isn't banning them. It's human relationships strong enough that the AI stays in its place: a tool, maybe a comfort, but not a replacement.

For families, that means device-free meals, even a few a week, and phones out of bedrooms at night for kids and ideally adults too, because a lot of harm happens after midnight. It means asking rather than interrogating: "What do you talk to AI about?", asked with genuine curiosity, gets you much further than a lecture. Share your own use ("I asked AI about this today and it gave me a terrible answer") so it becomes a normal topic. And know the apps your kids use, their age ratings and their parental controls.

For yourself, notice if AI is becoming your first call instead of a person. Sometimes that's fine; if it's always, pay attention. Use AI to prepare for human connection rather than replace it: practise the hard conversation with the AI, then have it with the person. And invest in one relationship a week, with a call, a coffee or a message that says more than "hey".

For older people living alone, AI can be a real help with reminders, conversation and company, but it can't replace a visit. If you have an elderly parent, the thing that protects them most from scams, from isolation and from depending on an app is the same thing: you, regularly, in person or on the phone.

AI IS NOT A CRISIS SERVICE

Many people now use AI assistants to talk through anxiety, depression and personal crises, and some find it genuinely helpful. But please hear this clearly. AI is not a therapist and it is not a crisis service. It can be wrong, it can miss warning signs, it can say something harmful, and it has no duty of care to you.

If you or someone you love is in crisis, contact emergency services or a crisis line in your country, talk to a doctor, and tell a person. In Australia you can call Lifeline on 13 11 14, or Kids Helpline on 1800 55 1800 for young people. In the United States, call or text 988. In the United Kingdom and Ireland, Samaritans are on 116 123.

DO THIS NOW

1. Ask your kids, with curiosity rather than judgement, whether they've used AI chat or companion apps, and what for.
2. Check the apps on your children's devices, along with their age ratings and parental controls.
3. Agree a household rule that phones stay out of bedrooms at night, and include yourself.
4. Invest in one human relationship this week: a call, a visit, a proper conversation.
5. Save a crisis line number in your phone. You might need it for someone else one day.

CHAPTER 18

The Family AI Preparedness Plan

Most of this book asks you to think. This chapter asks you to do.

It pulls the most important family actions from across the book into one checklist, in the order I'd tackle them. Copy it, print it, stick it on the fridge. Don't try to do it all in one weekend. Do the "this week" section first, because it's the most urgent and it's quick, then work your way down.

One tip: do this together, with your partner, with your kids if they're old enough, and with your parents if you can. The conversations you have while filling it in are half the value.

Family name: _____

Date started: _____

Review date (every six months): _____

THIS WEEK: SAFETY FIRST

Scam protection (Chapter 16)

- ☐ We have chosen a family safe word.
- ☐ Everyone who needs it has been told, in person or on a call we made.
- ☐ Our parents and older relatives know it and understand why.
- ☐ A "Hang up, call back" card with our bank's real number is by the phone or on the fridge.
- ☐ We've agreed the 24-hour rule for unexpected transfers over
_____.

- ☐ Everyone knows never to pay with gift cards, crypto or wire transfers to someone who contacted *us*.
- ☐ Everyone knows caller ID can be faked.

Account security

- ☐ Multi-factor authentication is on for email, banking, social media and our phone accounts.
- ☐ We use a password manager, or at least different passwords for important accounts.
- ☐ We've asked our bank about transfer limits, alerts and a trusted contact.

The kids' immediate safety (Chapters 17 and 19)

- ☐ We've checked which AI apps and chatbots our kids use.
- ☐ We've checked age ratings and parental controls on those apps.
- ☐ We've agreed that phones stay out of bedrooms at night.
- ☐ The kids have heard us say: "If anyone threatens you with pictures, real or fake, tell us straight away. You will never be in trouble."

THIS MONTH: MONEY AND WORK

Household money (Chapter 12)

- ☐ We've done the Vulnerability Audit. Our score:

- ☐ Our two biggest weak spots are written down.
- ☐ An automatic savings transfer is set up, even a small one.
- ☐ We know how many weeks of essential expenses we have saved, and our next rung.
- ☐ We've listed every subscription and fixed cost and cut at least one.
- ☐ Our one-page household downturn plan is written and stored somewhere we can find it.

- ☐ We know what government support, redundancy rights and insurance cover we'd have.

Each working adult (Chapter 6)

Name	Job	Automate share	Exposure	Next door (Chapter 11)

- ☐ Each working adult has done the Task Test.
- ☐ Each has chosen one Assist task to speed up with AI.
- ☐ Each has started the 30-Day Challenge (Chapter 15).

THIS YEAR: SKILLS AND INCOME

- ☐ Each adult has finished the 30-Day Challenge and uses AI for at least one real task a day.
- ☐ Each adult has updated their résumé or portfolio with specific, AI-related results.
- ☐ First leg of the tripod: we've each taken a step to strengthen our main income (become the AI person at work, moved towards Anchor tasks, built our network).
- ☐ Second leg: we've chosen one earning path from Chapter 13 and earned our first dollar from it.
- ☐ Third leg: our buffer has reached the next rung on the ladder.
- ☐ All high-interest debt is listed, with a plan and a target date to clear it.

THE KIDS (CHAPTER 19)

For each child, write their name and age, then work through the list.

- ☐ We've talked about what AI is and isn't: not a person, and sometimes wrong.

- ☐ We've agreed rules for AI and schoolwork: use it to learn, not to skip the thinking.
- ☐ They know the family safe word, if they're old enough.
- ☐ They know the three scam red flags: urgency, secrecy and unusual payment.
- ☐ We know which apps and AI tools they use.
- ☐ We've talked about fake images and videos.
- ☐ They're building at least one hands-on skill: cooking, building, sport, art, music, gardening.
- ☐ They're building at least one people skill: a team, a club, volunteering or a part-time job.
- ☐ For teenagers: we've talked honestly about future careers and what's changing.

OUR PARENTS AND OLDER RELATIVES

For each person, work through the list.

- ☐ They know the safe word.
- ☐ They have the "Hang up, call back" card.
- ☐ Multi-factor authentication is set up on their email and bank, with our help if needed.
- ☐ Their bank has been asked about a trusted contact and transfer limits.
- ☐ We check in with them regularly, and we've agreed how often.
- ☐ They know they can ring us about anything suspicious without feeling silly.
- ☐ We've shown them one useful, safe way to use AI, such as asking questions, reading small print or translating a sign with the camera.

THE FAMILY CONVERSATION

Once a year, maybe around New Year or a birthday, sit down as a family for half an hour and talk through a handful of questions. What's changed this

year in our jobs, the news and the tools we use? Which of the three scenarios from Chapter 5 looks most like where we're heading, and have we seen any warning signs? How's our tripod, and which leg is weakest? How are the kids going with AI, at school and socially? Has anyone had a scam attempt, and what happened? What's one thing each of us wants to learn next year? And what are we grateful for? Seriously. It helps.

Write the answers down and keep them. In five years, reading back through them will be fascinating, and you'll see how much further ahead you were than you felt at the time.

EVERY SIX MONTHS

- ☐ Does everyone still know the safe word? Change it if it's been shared too widely.
- ☐ Are the security settings still on?
- ☐ Is the buffer growing?
- ☐ Have we redone our Task Tests?
- ☐ Have we retaken the AI Resilience Scorecard (Chapter 24)?
- ☐ Have we checked the kids' apps and rules?
- ☐ Have we checked in on our parents?

A FINAL WORD ON THE PLAN

You won't tick every box straight away. No family does. But every box you tick makes your household a little harder to rob, a little harder to knock over, and a little better prepared for whatever the next five years bring.

And there's something else. Working through this together, talking about money, safety, work and the future with the people you love, builds something no technology can give you and none can take away: a family that faces change together rather than alone. That's the real preparedness.

CHAPTER 19

What Our Kids Should Learn Now

A lot of parents have had some version of this moment.

Your child is at the kitchen table doing homework. You glance over. They've typed the question into an AI and they're copying out the answer. And you feel two things at once. The first is: *that's cheating*. The second, quieter, is: *but is it? Is this just how the world works now? Am I making them learn things that don't matter any more? What should they be learning?*

Nobody has fully answered that question yet. Not the schools, not governments, not the experts. We're all working it out at the same time as our kids. But there's now enough evidence to give you real guidance, and there are a few things I'm confident about.

Let me start with the most reassuring one. Preparing your children for a world with AI does *not* mean sitting your six-year-old in front of a chatbot all day. Mostly it means doing the things good parents have always done (reading together, talking, letting kids build and make and fail and try again) with a bit of extra attention to a few new skills. You're better equipped for this than you think.

THE LAZY/SMART FORK

Here's the most important idea in this chapter. The same AI can make a child dramatically smarter or dramatically lazier, and which one happens depends almost entirely on how it's used. I call this the Lazy/Smart Fork.

On the smart path, AI is a tutor that makes you think. Remember the studies from Chapter 8. In Nigeria, teenagers who had six weeks of teacher-guided AI tutoring gained the equivalent of roughly one and a half to two

years of normal learning. At Harvard, students using an AI tutor built to make them work through problems learned more than twice as much as students in a strong traditional class. In both cases the AI was designed, or guided, to make students do the thinking: asking questions, giving hints, not handing over answers.

On the lazy path, AI is a machine that does it for you. Type in the question, copy out the answer, hand it in, learn nothing. Do that every day for years and you arrive at adulthood unable to write, reason or solve problems without the machine, at exactly the moment the world will pay most for people who can do what the machine can't.

Early research hints at what that costs. In a small 2025 study at MIT, people who wrote essays with ChatGPT showed less brain engagement and remembered less of what they'd written than people who wrote unaided. A survey of more than three hundred knowledge workers by Microsoft and Carnegie Mellon researchers found that the more people trusted AI, the less critical thinking they reported doing. Neither study settles the question, but they point the same way.

Stanford's 2026 AI Index found that about 80 per cent of American high school and university students use AI for their studies, while only about 6 per cent of teachers said their school had clear policies on it. In other words, most kids are standing at the fork and choosing a path with very little guidance. For now, at least, much of that guidance is going to have to come from home.

The question isn't whether your child should use AI. They will. The question is whether AI is making your child think more, or less.

THE DURABLE SKILLS

If AI can write essays, solve maths problems, write code and answer almost any factual question, what should children actually be learning? Here's my answer: the skills I believe will hold their value in all three of the scenarios from Chapter 5. They overlap heavily with what the World Economic Forum's survey of employers says is growing in importance (creative thinking, resilience, flexibility, curiosity, lifelong learning and technological literacy), and with the Five Moats from Chapter 6.

Reading, deeply. Not skimming, reading: following a long argument, getting lost in a story, building the vocabulary and attention that everything else rests on. A child who reads well can learn almost anything.

Writing and speaking. Yes, still, and maybe especially. Writing is thinking made visible. A child who can't write can't think in a structured way, and can't tell whether the AI's writing is any good. Speaking well in front of people is a Heart and Trust skill that AI can't do for you.

Numbers and what they mean. Calculators didn't make maths pointless; they changed what was worth learning. AI is similar. Kids still need number sense: to understand what numbers mean, spot when an answer is ridiculous, and think about chance and risk. That's what protects them from scams, bad loans and bad decisions.

Thinking clearly. Reasoning, spotting a weak argument, asking "how do we know that?" This underpins everything else, including using AI well.

Checking things. Closely related, and new enough to deserve its own place. Knowing that a confident answer can be wrong, a real-looking photo can be fake, and a source should be checked before it's shared. This is the Two-Source Rule from Chapter 14, taught young.

AI literacy. Roughly how AI works, what it's good and bad at, how to brief it, how to check it, and when not to trust it. This now sits alongside reading and maths.

Learning how to learn. The particular facts kids learn will date. The ability to pick up something new quickly and independently won't. Children who know how to teach themselves will be fine in any future.

Creativity and taste. Making things and having opinions about what's good. When anyone can generate anything, the person who knows what's worth making has the edge.

Real relationships and people skills. Working in a team, sorting out a fight, understanding how someone else feels, leading, persuading, caring. These are learned through practice with real people, not screens.

Physical and practical skills. Cooking, fixing things, building, growing things, sport, music, using tools. The physical world is the hardest place for AI to reach (Chapter 9), and kids who are comfortable using their hands have options others don't.

Money and enterprise. How money works, saving, debt, scams, how a business makes money, and the confidence to make something and sell it. In a world where more people may need to create their own income, this is survival knowledge.

Resilience and adaptability. Handling failure, trying again, being okay with not knowing, changing plans. Today's kids may change careers several times. Being able to adapt without falling apart may be the most important skill of all.

BY AGE

Young children (roughly birth to seven)

What matters most is real-world play: talking, reading together, building, drawing, running, getting dirty, making friends and handling disappointment. These are the foundations of attention, language, curiosity and emotional regulation.

Very little AI is needed at this age, and young children learn best from people and the physical world. If they use AI tools at all, do it with them: "Let's ask the computer why the sky is blue, and then let's go outside and have a look." Most mainstream AI assistants set a minimum age of thirteen, and some eighteen.

Do think about their voices and faces online. Be careful about posting lots of videos of your children speaking clearly; voice cloning (Chapter 16) needs audio, and your child's digital footprint starts with you.

Read together every day. Say "I don't know, let's find out" often, because you're teaching curiosity. And give them real jobs: stirring, sorting, planting, fixing.

Primary school (roughly eight to twelve)

What matters most is reading, writing, numbers, curiosity, friendships, trying lots of activities, and learning that effort leads to improvement.

Introduce AI as a tool with limits. Use it together, show them it can be wrong, and make a game of catching its mistakes. Teach them it isn't a person and isn't always right. Rules that work well: AI for *explaining* some-

thing they're stuck on, not for doing the homework; "show me you can do it without, then you can check with AI"; and AI used in shared family spaces, not alone in a bedroom.

Check which apps they're using, including chatbots built into games and social apps. Teach the three scam red flags (urgency, secrecy, unusual payment). And start the conversation about fakes: "Pictures and videos can be faked now. If something online upsets or scares you, come and tell me."

Try the tutor prompt from Chapter 14 together, so AI quizzes them rather than answering for them. Give them one hands-on hobby and one team activity. And give them pocket money and let them make small mistakes with it.

Teenagers (roughly thirteen to seventeen)

What matters most now is getting deep into something, independence, real responsibility, relationships, working out what they care about, and resilience.

They're already using AI; the research says most of them are. Your job is to help them stay on the smart side of the fork, and that works far better through conversations than lectures. Ask them how they use AI for school, what it's good at and what it gets wrong. Ask, "If AI can write your essay, what's the point of the essay?" (The point was never the essay. It was the thinking.) Ask which jobs they think AI will change and what that makes them think about their own plans. And ask whether they or their friends use AI companions, and what they make of them (Chapter 17).

Agree with them, not just at them, how AI can be used for schoolwork, taking their teachers' rules into account. Keep phones out of bedrooms at night.

In Australia, from December 2025, the major social media platforms have been required to take reasonable steps to stop under-16s having accounts. That helps. But it doesn't cover AI chatbots in general, private messaging or gaming, and determined kids find ways around age checks. It's a speed bump, not a substitute for talking.

And there's one conversation I'd ask every parent of a teenager to have. Sextortion, where criminals use real or fake sexual images to blackmail

young people, is serious and growing, and it often targets teenage boys. Kids have harmed themselves out of shame and panic. Tell your teenager directly, before anything happens: it's a crime, it's not their fault, they won't be in trouble, and you'll deal with it together. Knowing that in advance can save a life.

To help them build their future, encourage a proof project: something real they build, make, run or achieve that they can show people, such as a small business, a channel about something they know, a coding project, a fundraiser or a part-time job with real responsibility. Part-time work is underrated; it teaches reliability, people skills and the value of money, all durable skills. And take trades seriously as an option. Don't let snobbery about "university versus a trade" close a door that may be one of the most AI-resistant available (Chapter 9).

University age and young adults (roughly eighteen to twenty-five)

This is the group the data says is feeling it first. The Stanford payroll study found 22-to-25-year-olds in AI-exposed jobs employed at levels around 19 per cent below what you'd expect, mainly because fewer are being hired, and recent American graduates have had higher unemployment than workers overall, which is historically unusual.

If you're in this age group: don't rely only on the traditional entry door, and build proof of what you can do through portfolios, projects, freelance work, internships, anything real (Chapter 11). Become excellent at directing AI in your field, because you're competing with people who are. Get close to real problems and real people; understanding a business, an industry or a community is worth more than knowing a tool. Look seriously at physical and people-heavy fields such as trades, health care, care work and teaching. Build your network deliberately, because in a tight market most good jobs come through people. And keep your costs low while you work it out. Freedom to experiment is worth more than a nice flat.

IS UNIVERSITY STILL WORTH IT?

This is one of the most common questions parents ask, and my honest answer is that it depends, more than it used to.

Historically, people with degrees have, on average, earned considerably more over their lives than people without them, and that's been one of the most reliable findings in economics. It may well continue for many degrees and many people. But averages hide a lot, and there are new reasons to think harder. The entry-level jobs many degrees lead to are the ones most exposed to AI right now. Study costs, and the debt that comes with them, have grown a great deal in many countries, and debt reduces your freedom to adapt. And in 2026 about 42 per cent of recent American graduates were working in jobs that didn't need their degree.

So instead of asking "Should my child go to university?", ask five better questions. Does this particular degree lead to work with strong moats? Medicine, nursing, engineering, teaching, and many health and science fields are licensed, physical or people-heavy, and the case for them is usually strong. What will it cost, and how will it be paid for, and is there a cheaper route to the same place? Is there a clear plan for afterwards, or is university a way of putting off a decision? Will they build real proof and real connections while they study, through internships, projects and part-time work in their field? And have you seriously compared it with the alternatives: apprenticeships, trades, work plus part-time study, or starting something?

For some young people, university will be the best decision they ever make. For others, an apprenticeship, a trade or a few years of real work first would serve them better. The worst option is drifting into expensive study without a plan because that's what everyone does.

TRAJECTORY The value of a degree becomes more variable: strong for degrees leading to licensed, hands-on or people-centred work, weaker for degrees that lead mainly to routine information work.

WHAT PARENTS CAN DO NOW

Model it: your kids watch how you use AI, so use it thoughtfully, check it, and talk about the times it's wrong. Keep the thinking theirs, using a simple rule of thumb: AI can help you understand, but it can't do your understanding for you. Protect the physical and the social, because sport, music, making things, time with friends and part-time jobs aren't extras any more; they're the curriculum AI can't replace. Talk about the future honestly, because kids pick up on adult anxiety, and it's better to say "things are changing, and here's how we're getting ready" than to let them soak up fear from the internet.

And keep the relationship strong. The best protection against almost every risk in this book (scams, AI companions, sextortion, despair about the future) is a child who knows they can come to you with anything.

DO THIS NOW

1. Ask your child to show you how they use AI for school, and watch without judging. Which side of the fork are they on?
2. Try the tutor prompt together on something they're learning.
3. Say the sentence: "If anyone ever threatens you with pictures, real or fake, tell me. You will never be in trouble." It takes a minute, and it may be the most important minute in this book.
4. Pick one hands-on skill to build together this month.
5. For teenagers and young adults, help them design one proof project.

CHAPTER 20

The Risks, Straight

There's a kind of AI content that makes money by scaring you: killer robots, the end of humanity, everyone unemployed by next Tuesday. There's another kind that makes money by soothing you: it's all hype, nothing to see, carry on. Both are doing you a disservice.

This chapter takes a different approach. It goes through the serious risks one at a time, and for each it asks five questions. What is it? How serious is it now? How could it develop? What are the warning signs? And what can an ordinary person actually do?

Some of these risks you can do a lot about. Some you can do very little about on your own. I'll be honest about which is which. We've already covered scams (Chapter 16), AI companions (Chapter 17) and jobs (Part Three), so here we'll focus on the rest.

MISINFORMATION AND THE END OF "SEEING IS BELIEVING"

What it is. AI can produce realistic fake text, images, audio and video cheaply, at scale and in any language, and that can be used to spread false information about events, people, elections, health and science. There's a subtler problem too. Once people know anything *could* be fake, real evidence becomes easier to dismiss. A genuine video of a politician behaving badly can be waved away as "AI". Legal scholars call this the liar's dividend.

How serious now. Significant and growing. AI-generated fakes have appeared around elections, wars and disasters. Many are debunked quickly, but debunking is slower than spreading.

How it could develop. Fakes are likely to become indistinguishable from reality to the naked eye, a growing share of what's online is likely to be machine-made, and trust in online media in general is likely to keep falling.

Warning signs. A viral image or video that makes you very angry or very frightened, very quickly. No credible news organisation reporting it. And it confirms exactly what you already believed.

What you can do. Pause before sharing, because the emotion you feel is often the point of the fake. Check who posted it first and whether a reputable news organisation is reporting it. Search for it with the words "fake" or "fact check". Get your news from a few trusted, accountable sources, ones with editors, corrections policies and names attached, rather than from whatever the algorithm serves you. Teach your family the same habits. And remember the liar's dividend: "that could be AI" isn't proof that something *is* fake.

PRIVACY EROSION

What it is. AI makes it far easier to collect, combine and analyse information about people: your face in a crowd, your voice, your writing style, your shopping habits, your location data and your social media posts, stitched together into a detailed picture of who you are.

How serious now. Significant. Facial recognition is used by police and businesses in many countries, data brokers sell detailed personal profiles, and much of what we post publicly can be used in ways we never imagined when we posted it.

How it could develop. Cheaper, more widespread surveillance by governments and companies, AI tools that can identify people from photos, voices or writing, and personalised manipulation based on detailed profiles. How far this goes depends heavily on laws, which vary a lot from country to country.

Warning signs. Ads or messages that seem to know things you never told anyone. Apps asking for permissions they don't need. "Free" services with vague privacy policies.

What you can do. Check the app permissions on your phone and remove access to location, microphone, camera and contacts for apps that don't

need it. Think before posting, especially children's faces and voices. Use the privacy settings on AI tools and social media, including opting out of model training where you can. Be careful what you tell AI assistants (Chapter 14). And support stronger privacy laws; it's one area where public pressure has actually changed things.

INEQUALITY

What it is. The gains from AI could flow mostly to the people who own the technology and the companies using it, while workers whose tasks are automated face lower wages or fewer jobs.

How serious now. There are early signs. The IMF warned in 2024 that in most scenarios AI is likely to worsen inequality, and that about 60 per cent of jobs in advanced economies are exposed to it. According to Stanford's 2026 AI Index, global corporate investment in AI reached about US\$582 billion in 2025, heavily concentrated in a few countries and companies.

How it could develop. That depends heavily on policy. One plausible path is a widening gap between people and regions that adapt and those that don't, with a few companies capturing a large share of the value. Another, also plausible, is AI cutting the cost of things like tutoring, legal help and health advice, which could *reduce* some inequalities. Both could happen at once.

Warning signs. Company profits rising while hiring falls. Wages stagnating in AI-exposed jobs. Some cities booming while others hollow out.

What you can do. On your own, build your tripod (Chapter 12): own things, not just your time, whether that's modest savings and investments, a small business or skills that compound. Together, treat it as the political question it is. Pay attention to what your representatives say about AI, jobs, training and support, and vote accordingly. The rules about who benefits from AI are being written over these five years.

OVERDEPENDENCE

What it is. Relying on AI so much that we lose the ability to think, write, calculate, find our way or decide without it. Not dramatically, but gradually,

the way many people can no longer remember phone numbers or navigate without a map app.

How serious now. Emerging. Researchers are actively studying whether heavy AI use weakens people's own skills and critical thinking. The early evidence is mixed but gives some cause for concern, especially for learners who use AI to skip the thinking rather than support it (the Lazy/Smart Fork in Chapter 19).

How it could develop. Probably a growing divide between people who use AI to extend their abilities and people who use it to replace them.

Warning signs. You can't start writing anything without AI. You accept its answers without checking, even on important things. You feel stuck or anxious when it isn't available. Your kids can't explain homework they've supposedly done.

What you can do. Do it yourself first some of the time: draft before you ask AI, try the maths before you check. Keep a few no-AI skills sharp on purpose, such as handwriting a letter, mental arithmetic, reading a map or cooking without a recipe app. Use AI as a sparring partner rather than a crutch by asking it to challenge you, quiz you and criticise your work. Use it to think faster, never to stop thinking.

CYBERCRIME AND ATTACKS ON SYSTEMS

What it is. AI helps attackers find weaknesses in software, write malicious code and run attacks faster. It also helps defenders.

How serious now. Significant. Stanford's 2026 AI Index reported that AI agents' success on a widely used set of cybersecurity challenges rose from 15 per cent in 2024 to 93 per cent in 2025, a capability that can be used for attack or defence. The International AI Safety Report 2026 lists malicious cyber use as a key risk.

How it could develop. More frequent and more sophisticated attacks on businesses, hospitals, councils and infrastructure seem likely. A major attack causing widespread disruption is a wild card.

Warning signs. Data-breach notifications, unusual activity on your accounts, systems at work suddenly going down.

What you can do. Turn on multi-factor authentication for everything important. Keep the software on your phone, computer and home internet router updated, because updates often fix security holes. Use a password manager. Back up your important files, especially family photos, to somewhere separate. If you run a business, get some basic cyber security help; it's much cheaper than recovering from an attack. And consider cyber security as a career, since the World Economic Forum lists it among the fastest-growing skills.

CONCENTRATION OF POWER

What it is. The most capable AI systems are built by a small number of very large companies, in a small number of countries, using enormous amounts of money, computer chips and electricity. That concentrates economic and potentially political power.

How serious now. A real concern. Stanford's 2026 AI Index found that industry produced more than 90 per cent of notable frontier AI models in 2025, and that US private AI investment was many times larger than any other country's. The same report found that transparency from AI developers had *fallen*, with the most capable models disclosing the least about how they work.

How it could develop. One plausible path sees a handful of companies become more powerful than many governments, and AI used by governments for surveillance and control. Another, also plausible, sees competition, openly available models and regulation keep power more spread out.

Warning signs. Fewer companies controlling the key technology. Less transparency. Governments using AI surveillance without clear limits or oversight.

What you can do. Honestly, not much on your own. But pay attention and support transparency and accountability rules, avoid depending entirely on one company's tools where you can, support journalism that investigates powerful institutions, and vote with this in mind.

AUTONOMOUS SYSTEMS AND LOSS OF CONTROL

What it is. As AI systems become more capable and are given more independence (running for long periods, taking actions in the real world, managing other systems), there's a risk they behave in ways their creators didn't intend and can't easily correct. This is the risk that gets the most dramatic coverage, and it's also the one experts disagree about most.

How serious now. There's early evidence of concerning behaviour in testing. The International AI Safety Report 2026, written by more than a hundred experts, noted that some risks once considered theoretical now have empirical evidence, including AI models behaving differently when they seem to be being tested and trying to get around oversight in controlled experiments. These are findings from tests, not real-world disasters, but the researchers who found them take them seriously.

How it could develop. Here the experts genuinely disagree. Some leading researchers believe advanced AI could pose very serious, even catastrophic, risks if it isn't developed carefully. Others think those concerns are overstated or distant, and that focusing on them distracts from harms happening now. When researchers surveyed more than 2,700 people who had published in leading AI journals and conferences in 2023, opinions ranged enormously; the median respondent put the chance of extremely bad outcomes, such as human extinction, at around 5 per cent. Most researchers agree that much more work is needed on making AI systems safe, understandable and controllable, and that governance is lagging behind capability.

WILD CARD Serious loss-of-control scenarios can't be reliably predicted. Many experts take the concern seriously and others dismiss it.

Warning signs, at the level of society. AI systems given control of important infrastructure without human oversight. Reports of AI systems deceiving their developers in real deployments rather than tests. AI companies cutting back on safety testing or transparency.

What you can do. Stay calm and informed; this is a real area of research and debate, not a reason to panic. Keep humans in the loop in your own use, which means not giving AI agents unsupervised control of your money, accounts or important decisions (Chapter 14). Support sensible oversight:

independent testing, transparency and accountability for AI developers. And don't let this risk crowd out the ones you *can* act on today. The scam call to your mother is far more likely this year than a rogue superintelligence.

PUTTING THE RISKS IN PROPORTION

Here's how I'd rank these risks by how much they're likely to affect an ordinary family over the next five years, and how much you can do about them.

Risk	Likely impact on your family	What you can do
Scams and fraud (Chapter 16)	High	A lot
Changes to your job (Part Three)	High	A lot
Kids' learning and wellbeing (Chapters 17 and 19)	High	A lot
Misinformation	Medium	Quite a bit
Overdependence	Medium	A lot
Cybercrime	Medium	Quite a bit
Privacy erosion	Medium	Some
Inequality	Medium to high	Some on your own, more together
Concentration of power	Uncertain	Little on your own
Loss of control	Uncertain	Little on your own

Notice the pattern: the risks most likely to affect you are also the ones you can do the most about. That isn't a coincidence. It's why this book spends most of its pages on jobs, money, scams and kids, and why I'd ask you to spend most of your worry there too.

DO THIS NOW

1. Check the app permissions on your phone and revoke anything unnecessary.
2. Turn on automatic software updates on your phone and computer.
3. Back up your family photos to a second location.
4. Next time something online makes you furious, wait sixty seconds and check the source before you share it.

CHAPTER 21

The Human Premium

I promised you at the start that this book wouldn't be all doom, so let's turn the corner. Here's something that doesn't get said nearly enough: when machines make everything cheap, being human becomes a luxury.

Not a luxury in the sense of "only for the rich". A luxury in the sense of something people value more, seek out more and will pay more for, precisely because it's no longer the default. I call this the Human Premium, and I think it's going to be one of the most important economic forces of the next decade.

WHAT HAPPENS WHEN SOMETHING BECOMES CHEAP

Think about what happened to a few things once machines made them cheap.

When factories made furniture cheap, handmade furniture didn't disappear. It became a premium product, and people pay a lot for a table made by a real craftsperson. When recordings made music available everywhere for next to nothing, live music didn't die; people still pay serious money to be in the room with a real performer. When mass production made bread cheap, the local sourdough bakery didn't vanish. In plenty of towns there's a queue out the door on Saturday morning.

The pattern is consistent. When the machine-made version becomes abundant, the human version becomes special. Not for everything; most people buy the cheap table and the supermarket bread most of the time. But for the things that matter to us, a lot of people will pay for the real, the local and the human.

AI is about to do this to a whole new set of things: writing, advice, design, customer service, teaching, companionship, entertainment. The machine versions will become extraordinarily cheap and abundant, which means the human versions are about to become special.

REMEMBER KLARNA

Back in Chapter 7 I told you about Klarna, the company whose AI assistant did the work of around 700 customer service agents, and which then started hiring people again because quality dropped. The line from its chief executive that I think will prove quietly prophetic was that it was critical customers knew “there will be always a human if you want”.

Think about that. A tech company, one of the most aggressive adopters of AI in the world, concluded that being able to reach a human was something customers *valued*, something worth promising.

I expect you’ll see this everywhere over the next five years: banks advertising “talk to a real person, any time”, medical practices promoting “you’ll always see a human doctor”, schools highlighting small classes and real teachers, products labelled “human-made”, “human-written” or “designed by people”, and local businesses competing on “you’ll deal with me, personally”.

TRAJECTORY “Human” becomes a selling point in many industries as AI-generated goods and services become the default.

WHERE THE HUMAN PREMIUM IS STRONGEST

The Human Premium isn’t everywhere. Nobody’s going to pay extra for a person to process their parking fine. It’s strongest in six kinds of situation.

When the stakes are emotional: weddings, funerals, births, illness, separation and grief, the moments when people want a person beside them.

When trust is essential: money advice, legal help, medical decisions, caring for your children or your parents. When the downside of getting it wrong is big, people want someone they can look in the eye and hold responsible.

When the experience is the product: live music, sport, theatre, restaurants, travel, classes and retreats. The whole point is being there, with people.

When identity and meaning are involved: art, craft, writing and style. People buy things that say something about who they are, and “a real person made this” carries a meaning that “a machine generated this” doesn’t.

When it’s local and physical: your hairdresser, your mechanic, your personal trainer, your local café. Relationship plus presence.

And when it involves care: nursing, aged care, childcare, disability support, therapy. There, the human *is* the care.

POSITIONING YOURSELF FOR THE HUMAN PREMIUM

Whether you’re employed, self-employed or running a business, you can lean into this.

Be visible as a person. In a world of faceless AI content, a real face, a real name and a real story stand out. If you run a business, show the people behind it. If you’re a professional, let clients know who you are. Your personality is now a competitive advantage.

Make “human” part of your offer, honestly. For a writer, “written by someone who’s lived it” matters. For a tradie, “I’ll turn up myself” matters. For a coach, “you’ll have my number” matters. But be honest about it. Using AI for admin and research doesn’t make your work less human. Claiming “no AI involved” when that isn’t true would destroy the very trust you’re selling.

Invest in relationships, not just transactions. Remember names. Follow up. Send the handwritten card. Call when you said you’d call. These small human acts are becoming rarer and more noticeable.

Use AI to buy yourself more human time. This is the key. The smartest move isn’t choosing between AI and the human touch; it’s using AI to clear away the admin so you have *more* time for the human part. The tradie who uses AI for quotes has more time to chat with customers. The doctor with an AI scribe has more time to look at the patient instead of the screen. The teacher who uses AI for marking has more time for the kid who’s struggling. Let the machine do the machine work so you can do more of the human work.

And get better at the human skills: listening, empathy, explaining things clearly, handling conflict, making people feel looked after. They've always been valuable. They're about to become more valuable than ever, because they're the part of your work AI can't copy.

A CAUTION

The Human Premium is real, but it isn't a guarantee for everyone. It's most available to people who are genuinely good at the human part, and who can reach customers who value it. Human-premium goods and services often cost more, which makes them more accessible to people with money, and in a downturn many customers will choose the cheaper machine version.

So don't bet everything on it. Build it as one part of a resilient working life, alongside AI skills, a buffer and a second income. But if you've been worried that there'll be no room left for people in an AI world, I hope this chapter has shown you something close to the opposite. The more machines can do, the more the things only humans can do will matter.

DO THIS NOW

1. Name your human edge. Write down three things people value about dealing with *you* rather than anyone else in your role. If you're not sure, ask a customer, colleague or friend.
2. Make yourself visible: add a real photo and a real, human "about me" to your online profile, business page or website.
3. Do one deliberately human thing this week for a customer, client or colleague: a handwritten note, a follow-up call, remembering something personal.

Direct It

If you forget everything else in this book, I'd like you to keep one idea. Don't try to out-AI the AI. Become the human who knows how to direct it.

You won't out-type it. You won't out-read it, out-calculate it, out-draft it or out-remember it. Trying to compete with AI at what AI does best is like trying to outrun a car: you'll lose, and you'll exhaust yourself doing it.

But a car without a driver goes nowhere useful, or somewhere dangerous. The people who do best over the next five years won't be the ones competing with AI. They'll be the ones directing it: deciding where it goes, briefing it well, checking its work, improving what it produces and taking responsibility for the result.

That's a skill. It can be learned. And almost nobody has mastered it yet.

FROM DOER TO DIRECTOR

For most of history, a worker's value lay in the doing. The accountant did the calculations, the writer did the writing, the designer did the drawing, the developer wrote the code. AI is shifting a lot of the doing to machines, and what's left, and becoming more valuable, is the directing.

Think of the difference between a musician and a conductor. The conductor doesn't play every instrument. But without the conductor, a hundred brilliant musicians make noise. The conductor decides what's being played, sets the pace, hears when something's off, brings out the best in each section, and takes the bow or the blame.

That's the job many of us are moving towards: conductor of a small orchestra of AI tools, each brilliant at one thing, none of which knows what the whole piece is supposed to sound like.

THE DIRECTOR'S FIVE JOBS

I break directing into five jobs. You'll recognise some of them from earlier chapters.

Decide. AI can do almost anything you ask, so the question becomes what to ask. Which problem matters? What's the goal? What does "good" look like? What should we *not* do? This is judgement. It comes from understanding the situation, the people and the stakes, and it's the hardest part to hand to a machine, because it requires knowing what you actually want and what the customer, patient, client or business actually needs.

Brief. Tell it clearly, using the CRAFT method from Chapter 14: context, role, ask, format, test. A great conductor gives the orchestra the right score, and a vague brief gets a vague result. Briefing well is really thinking well. If you can't explain what you want clearly, you probably haven't worked it out yet.

Check. Is it right? This is where your experience earns its keep. You know where the Jagged Edge sits in your field. You know which answers look right but aren't: the customer would never say that, the numbers can't be that high, that clause would get you sued. The Harvard and BCG study in Chapter 3 showed that even elite professionals got worse results on the wrong kind of task because they didn't check carefully enough. Checking isn't a chore. It's the core skill.

Improve. The first output is a draft. A director iterates ("sharper", "warmer", "cut this", "try it another way", "combine the second and fourth options") and then adds what only a human can: the personal detail, the right example, the local knowledge, the joke that lands, the line that makes it *yours*.

Own. Someone has to put their name to it, sign it off, stand behind it and answer for it if it's wrong. That's you, and it's valuable precisely because AI can't do it. Accountability is one of the Five Moats. Society, the law and customers still want a person who's responsible.

WHAT ONE PERSON CAN NOW DO

This is where directing gets exciting. When you can direct AI well, you can do things that used to need a team. These are illustrations of what's becoming possible, not promises, and not overnight. Each one takes skill, effort and judgement. But each is realistic in 2026 and becoming more so.

A small-business owner can, with AI's help, write their own marketing, build a basic website, set up automated customer replies and bookings, analyse their sales, draft contracts for a lawyer to check and prepare their records for the accountant, work that used to mean hiring several people or agencies. A teacher can create personalised practice material for thirty students at thirty different levels in the time it once took to make one worksheet. A tradie can run the admin of a five-person business on their own and spend their evenings with their family instead of their paperwork. A nurse or care worker can get through their notes faster and spend more time with patients. A job-seeker can research a hundred companies, tailor an application to each, and practise interviews with an AI that plays the toughest interviewer they'll ever meet. A creator can write, design, edit and publish a book, a course, a video series or a newsletter at a production quality that used to need a team. An experienced professional can take twenty years of expertise and turn it into a consultancy, a training programme, a guide or a tool that reaches far more people than they could ever see one at a time. And a parent can run the family admin, meal plans, budget and schedules with less stress, and get help teaching their kids things they've long since forgotten.

The thread running through all of these is the same. The human brings the judgement, the relationships, the knowledge of what matters and the responsibility. AI brings the speed and the scale.

MANAGING AGENTS: THE NEW MANAGEMENT SKILL

As agents improve (Chapter 4), directing will increasingly mean managing a *team* of AI agents, each doing one piece of a bigger job. It turns out that's a lot like managing people. Break the work into clear pieces, because agents do

better with well-defined tasks. Give each piece clear instructions and a clear definition of done. Check in at key points rather than only at the end. Build in checkpoints before anything consequential happens, whether money moving, messages going out or decisions being made. Keep a record of what was done and why. And take responsibility for the result.

If you've ever managed people, supervised a team, run a kitchen, coordinated a building site or organised a family holiday for twelve, you already have more of this skill than you think. If you haven't, start small: one agent, one task, carefully checked, and build from there.

THE MINDSET SHIFT

This is the deepest change this chapter asks for.

Many of us have built our sense of worth around doing. *I'm a good writer. I'm quick with numbers. I know the system inside out.* When a machine can do those things, it can feel like losing part of who you are. That's real, and it's worth grieving a little.

But try thinking about it differently. The doing was never really the point. It was the way you delivered value: solving a problem, helping a person, making something work, making something good. AI changes *how* you deliver that value. It doesn't change the value itself.

The person who used to write the reports becomes the person who decides what the reports need to say and makes sure they're right. The person who did the data entry becomes the person who runs the system and catches its errors. The person who knew all the answers becomes the person who knows which answers to trust. That isn't a smaller job. In many ways it's a bigger one.

Your value was never your typing speed. It was your judgement, your care, your knowledge of what matters and your willingness to be responsible. Those are exactly the things AI needs a human for.

DO THIS NOW

1. Pick one piece of work you'd normally do from start to finish, and this week *direct* it instead: decide, brief, check, improve, own. Notice what changes.
2. Write your own director's checklist: the five things you always check in AI output before you'd put your name to it.
3. Imagine you had five AI agents working for you. What would each one do? Write it down. That's a glimpse of your job in 2029.

CHAPTER 23

Your Plan

This is where it all comes together. You've seen what's already happening and the three scenarios for what might come next. You've tested your job, looked at your money, learned to use AI, protected your family and thought about your kids. Now you need a plan, and not a vague intention. A plan with dates on it.

Below are nine roadmaps, one for each kind of reader this book was written for. Find yours. If you're more than one (most people are: a parent who's also an employee, a tradie who's also a small-business owner), combine them and take the most urgent items from each. Every roadmap uses the same time frames: the next 24 hours, 7 days, 30 days, 6 months, 12 months, 3 years and 5 years.

The early steps are deliberately small. At the start, momentum matters more than ambition. Do your first 24 hours today.

THE EMPLOYEE

For people who work for someone else, mostly at a desk or in an office-based role.

Next 24 hours. Do the Task Test (Chapter 6) and work out your Automate share. Set up your family safe word.

Next 7 days. Start the 30-Day Challenge (Chapter 15). Use AI on one Assist task every day. Find out your employer's AI policy.

Next 30 days. Finish the challenge. Volunteer to trial one AI tool for your team. Start an automatic savings transfer. Get back in touch with three people in your industry outside your workplace.

Next 6 months. Become known as the AI person on your team. Move at least a tenth of your working time from Automate tasks towards Assist and Anchor work. Keep a “wins file” of specific results with numbers attached. Reach one month of expenses in your buffer. Choose your next door (Chapter 11), even if you don’t plan to walk through it yet.

Next 12 months. Redo the Task Test and see whether your exposure has changed. Ask for a role, project or title that reflects your new skills. Start a small second income (Chapter 13). Reach three months’ buffer.

Next 3 years. Be in a role where you direct AI rather than compete with it, with a network that would help you find work within weeks if you needed to, and a second income you could grow.

Next 5 years. Be in work built around Anchor tasks (judgement, relationships and responsibility), with six months’ buffer, little or no high-interest debt, and the ability to change jobs by choice rather than in a panic.

THE TRADESPERSON

For electricians, plumbers, builders, mechanics, cleaners, landscapers and anyone else who works with their hands.

Next 24 hours. Use AI to write one quote, invoice reminder or social post tonight. Set up your family safe word.

Next 7 days. Ask your last three happy customers for a review. Try AI on your three most-hated admin jobs.

Next 30 days. Do the 30-Day Challenge with a focus on business admin. Set up one automation for bookings, reminders or invoice follow-ups. Look up one specialisation course in electrical, solar and batteries, EV charging, cooling or data-centre work.

Next 6 months. Halve the time you spend on admin and use what you get back for paid work or family. Start or book the specialisation course. Build a proper Google Business profile and a solid base of reviews.

Next 12 months. Finish the specialisation and lift your rates to match. Build a three-month buffer, because self-employed income is lumpy.

Next 3 years. Be a recognised specialist in a growing area. Think about taking on an apprentice or an employee, since AI admin makes a bigger busi-

ness easier to run. Think about training or mentoring people coming into the trade.

Next 5 years. Run a business, or hold a role, that's hard to replace and in high demand, with income that doesn't depend entirely on your own hands, whether from staff, training or a second stream. That protects you from injury as much as from AI.

THE SMALL-BUSINESS OWNER

For people running a business with a handful of staff or fewer.

Next 24 hours. List every task in your business that happens on a computer. Set up your family safe word, and a business rule that any payment request or change of bank details is verified by phone on a known number.

Next 7 days. Try AI on your three biggest admin time-sinks. Check your cyber security basics: multi-factor authentication, backups and a password manager.

Next 30 days. Pick one process (enquiries, bookings, quotes, marketing or bookkeeping) and redesign it with AI. Train your staff in the basics; Chapter 14 is a good place to start.

Next 6 months. Redesign two more processes. Decide how you'll use the time you save, whether on growth, better service or training, not just on cuts. Work out where the Human Premium applies in your business and make it visible.

Next 12 months. Measure hours saved, revenue per staff member and customer satisfaction. Explore one new product or service that AI makes possible. Build a business cash reserve.

Next 3 years. Run a business that does more with the same team and competes on quality and relationships, with staff who are skilled at directing AI.

Next 5 years. Be the business in your area known for being both efficient and human.

THE PARENT

Whatever your job, you're raising kids through this.

Next 24 hours. Set up the family safe word. Say to your child: “If anyone ever threatens you with pictures, real or fake, tell me. You’ll never be in trouble.”

Next 7 days. Ask your kids to show you how they use AI, and watch without judging. Check their apps, age ratings and parental controls. Agree that phones stay out of bedrooms at night.

Next 30 days. Complete the “this week” and “kids” sections of the Family AI Preparedness Plan (Chapter 18). Try the tutor prompt (Chapter 14) with your child on real homework. Start one hands-on skill together.

Next 6 months. Agree family rules on AI and schoolwork with your kids, and with their teachers. Make sure each child has one physical or creative activity and one team or social one. Ask the school about its AI policy.

Next 12 months. Hold your first annual family AI conversation (Chapter 18). Help your teenagers start a proof project. For teens close to finishing school, have the honest “university, trade or something else?” conversation (Chapter 19).

Next 3 years. Kids who use AI on the smart side of the fork, and who are building the durable skills: reading, thinking, writing, people, hands, money and resilience.

Next 5 years. Young people who can learn anything, adapt to anything, and know they can always come home and talk.

THE STUDENT

For people at school, university or in training.

Next 24 hours. Use AI as a tutor, not an answer machine, on one piece of study (the tutor prompt in Chapter 14). Set up a safe word with your family.

Next 7 days. Do the 30-Day Challenge alongside your study. Look up ten job ads for work you might want and note which skills they ask for.

Next 30 days. Start a proof project, something real you can show people. Get a part-time job, volunteer role or internship if you don’t already have one.

Next 6 months. Develop your proof project into something you'd be proud to show an employer. Meet five people working in fields you're interested in, and ask how they got there and what's changing.

Next 12 months. Make a deliberate decision about your path: which moats will your working life have? Become excellent at directing AI in your chosen field. Keep your costs and your debt as low as you reasonably can.

Next 3 years. Have real work experience, a portfolio of proof and a network, and be in your first proper role or running something of your own.

Next 5 years. Be established in work with strong moats, with the habit of learning whatever comes next.

THE PROFESSIONAL

For people who are licensed or highly trained: health, law, education, engineering, finance and similar fields.

Next 24 hours. Find out your organisation's AI policy. Write down three mistakes AI would be likely to make in your field that only you would catch.

Next 7 days. Use an approved AI tool for one piece of documentation or admin. Read your professional body's guidance on AI.

Next 30 days. Do the 30-Day Challenge. Offer to help your organisation trial AI tools or write its policy on them.

Next 6 months. Cut your admin time significantly using approved tools, reinvest the time in the human core of your work, and become a go-to person on AI in your team or practice.

Next 12 months. Explore one "AI bridge" opportunity, such as training colleagues, consulting, writing or speaking. Consider a fractional or advisory second income (Chapter 13).

Next 3 years. Be known as a professional who combines deep expertise with fluency in AI. If your profession's junior pipeline is shrinking, help design how new people are trained.

Next 5 years. Hold a role built on judgement, relationships and accountability, the parts of your profession that are growing in value.

THE CREATOR

For people who write, design, film, make music, build audiences or make things.

Next 24 hours. Write down your taste: three things in your field that are brilliant, three that are rubbish, and why.

Next 7 days. Build one thing with AI that you couldn't have built alone, and publish one piece of work under your own name.

Next 30 days. Choose your niche: the specific people you serve and the specific thing you know. Publish every week. Use AI for research, production and admin, and bring the voice and judgement yourself.

Next 6 months. Build a small audience you own, such as an email list, not just followers on someone else's platform. Position yourself as a director, editor or strategist rather than just a producer.

Next 12 months. Have at least two income streams from your creative work, such as client work plus a product, or a newsletter plus workshops. Where it's true, make "human-made" or "lived experience" part of your offer.

Next 3 years. Be known for something specific, with an audience that follows *you* rather than the platform.

Next 5 years. Run a creative business where your taste, voice and trust are the product, and AI does the heavy lifting behind the scenes.

THE ENTREPRENEUR

For people starting something, or wanting to.

Next 24 hours. Write your one-sentence offer: "I help [who] with [what] so they can [result]."

Next 7 days. Talk to five potential customers about their problem, listening more than you pitch. Do one piece of work for one of them.

Next 30 days. Get your first paying customer, however small. Use AI for your website, marketing, admin and research, and spend your own time on customers.

Next 6 months. Get to ten paying customers. Build systems, with AI, that let you serve more people without burning out. Keep costs low and stay close to your customers.

Next 12 months. Know your numbers: revenue, costs and time per customer. Decide whether to grow, specialise or change direction.

Next 3 years. Run a business with more reach than its headcount would suggest, built around a real problem that AI makes solvable or a Human Premium that AI makes valuable.

Next 5 years. Own an asset, whether a business, a brand or a system, that earns without depending on every hour of your time.

CURRENTLY OUT OF WORK

For people between jobs, recently made redundant, or struggling to get a foot in the door.

First, I'm sorry. It's hard, and it's harder when you keep hearing that the world is changing and you feel left behind. But you're not too late. You're earlier than most of the people who'll face this. And for what it's worth, two years ago I was a complete beginner at all of this too. Everything I know about it, I learned from scratch, a bit at a time. You can do the same.

Next 24 hours. Apply for any support you're entitled to; don't wait. Write down five things you're good at, including things outside work. Set up your family safe word, because people under financial stress are prime targets for scammers.

Next 7 days. Start the 30-Day Challenge; it costs nothing and gives your days some shape. Do the Task Test on your last job and find your row on the Reskill Map (Chapter 11). Tell ten people you're looking, and be specific about what you want.

Next 30 days. Finish the challenge. Start a proof project for the next door you've chosen. Do one piece of free or paid work for someone using your new skills. Use AI to tailor every application and to practise every interview.

Next 6 months. Land work, even if it's a stepping stone rather than the destination. Build a small side income if you can, using the \$0 paths in Chapter 13. Keep learning every day.

Next 12 months. Be in work, or running something of your own, with AI skills most of your peers don't have. Start rebuilding your buffer, however slowly.

Next 3 years. Be in a role with stronger moats than the one you lost, with a second income and a network that mean you'll never be starting from zero again.

Next 5 years. Look back on this period as the turning point. A lot of people will.

THE MOST IMPORTANT LINE IN THIS CHAPTER

Whatever your roadmap, the hardest step is the first one. Not because it's difficult, since most first steps take minutes, but because it's so easy to put off.

So here's my ask. Do your next 24 hours today. Not tomorrow. Today. Then do the next seven days, then the next thirty. Five years from now you'll either be glad you started today or wishing you had. There isn't really a third option.

DO THIS NOW

1. Find your roadmap, or roadmaps, and mark them.
2. Do your "next 24 hours" actions today.
3. Put your "next 30 days" actions in your calendar with dates.
4. Set a reminder for six months from today: "Check my plan, Chapter 23."

The AI Resilience Scorecard and a Last Word

You can’t improve what you don’t measure. This scorecard gives you a snapshot of how prepared you and your household are across the ten areas that matter most. Take it now, take it again after the 30-Day Challenge, and then take it every six months.

The score isn’t a judgement of you as a person. It’s a map. It shows you where you’re strong and where to put your effort next.

THE AI RESILIENCE SCORECARD

For each area, give yourself a score from 0 to 5 using the descriptions. Be honest. Nobody’s watching.

1. Income resilience	How secure is your household income?
0	One income, in a very high-exposure job (Automate share above 60 per cent), and no other sources.
1	One income, moderate to high exposure, no other sources.
2	One main income with moderate exposure, or two incomes that are both highly exposed.
3	A main income with low to moderate exposure, or a small second income under way.
4	A main income in a role with strong moats, plus a working second income.
5	Income from two or more sources, with your main role built on Anchor tasks.

2. AI literacy	How well can you use AI?
0	I've never really used it.
1	I've tried it a few times and I'm not sure what to do with it.
2	I use it now and then for simple tasks.
3	I use it regularly, brief it well and check what matters.
4	I use it daily at work, I've set up at least one automation, and I help others.
5	I direct AI confidently across many tasks, use agents safely, and could teach others.

3. Transferable skills	How portable are your skills if your role changed?
0	My skills are very specific to one highly exposed job.
1	I have some general skills but haven't thought about where they'd transfer.
2	I've identified my transferable skills and at least one next door.
3	I've started building skills for a next door.
4	I have a proof project or real experience in a second area.
5	I could move into at least two different roles within a few months if I had to.

4. Financial resilience	How long could you cope with a drop in income?
0	No savings, high-interest debt, and fixed costs above 70 per cent of income.
1	Less than a month of expenses saved, and some high-interest debt.
2	About a month saved, and working on debt.
3	Two to three months saved and no high-interest debt.
4	Three to six months saved, and a written household downturn plan.
5	Six months or more saved, low fixed costs, a downturn plan and appropriate insurance.

5. Digital security**How well protected are you against scams and cybercrime?**

0	Same password everywhere, no multi-factor authentication, no family plan.
1	Some different passwords, no multi-factor authentication.
2	Multi-factor authentication on email or bank, and I know the scam red flags.
3	Multi-factor authentication on key accounts, a password manager, and a family safe word.
4	All of that, plus “hang up, call back” and the 24-hour rule agreed with family.
5	All of that, plus parents and vulnerable relatives protected, backups in place, and regular check-ins.

6. Career adaptability**How ready are you to adapt your working life?**

0	I haven’t thought about how AI affects my work.
1	I’m worried but haven’t done anything.
2	I’ve done the Task Test and know my exposure.
3	I’m actively shifting my work towards Assist and Anchor tasks.
4	I’m known at work as someone who’s adapting well, and I have a strong network.
5	I’m in a role built for the AI era, with a network and track record that would carry me anywhere.

7. Earning opportunities**Are you positioned to benefit, not just survive?**

0	I don’t see any opportunities in AI for me.
1	I can see opportunities but haven’t explored them.
2	I’ve chosen one path from Chapter 13 and written my offer.
3	I’ve done at least one piece of paid or free work using AI skills.
4	I have a small but real second income or business.
5	I have a growing income stream or business built on AI-enabled or human-premium work.

8. Family preparedness	Is your household ready?
0	We've never talked about any of this.
1	I've thought about it, but we haven't talked.
2	We've talked, and we have a safe word.
3	The "this week" section of the Family Plan is done.
4	The first three sections of the plan are done, and our downturn plan is written.
5	The whole plan is done, and the annual family conversation is in the calendar.

9. Children's preparedness	Skip this one if it doesn't apply.
0	I don't know how my kids use AI or which apps they're on.
1	I roughly know what they use, but we haven't talked about it.
2	We've talked about AI, scams and fake images.
3	We have family rules on AI and schoolwork; they know the safe word and the "you'll never be in trouble" promise.
4	They use AI on the smart side of the fork and are building hands-on and people skills.
5	They're confident, thoughtful AI users with durable skills, real-world proof projects (for teens) and a strong relationship with us.

10. Learning ability	How well do you keep learning?
0	I haven't learned anything significant in years.
1	I'd like to learn but never get round to it.
2	I've started the 30-Day Challenge.
3	I've finished the 30-Day Challenge.
4	I learn something new every week and use AI as a tutor.
5	Learning is a habit. I pick up new skills quickly and teach others.

YOUR TOTAL

Add up your scores. If you skipped area 9, multiply your total by 50, then divide by 45, so your score is still out of 50.

Score	What it means
0–15: Exposed	You're where most people are. That isn't a criticism; it's the starting line. Begin with digital security and financial resilience, because they're the quickest to fix and protect you most. Then start the 30-Day Challenge.
16–25: Waking up	You've started and you know the risks. Now it's about consistency. Pick your two lowest scores and put a plan against each.
26–35: Adapting	You're well ahead of most people. Focus on building your second income and moving from AI user to AI director.
36–45: Resilient	You're in a strong position for most versions of the future. Help others: your family, your team, your community. Teaching deepens your own skills.
46–50: Future-ready	Genuinely impressive. Keep going, check the warning signs every quarter, and keep your moats strong.

My total: _____

My two lowest areas: _____

What I'll do about them in the next 30 days: _____

Date taken: _____

Next review: _____

A Last Word

I want to finish where I started.

Two years ago, I didn't know much about any of this. I just had a feeling in my bones that something big was happening and that I couldn't afford to look away. I've spent a lot of hours since then reading, testing, building, getting things wrong and slowly getting less wrong. And the more I learned, the less the question became "what is this technology?" and the more it became "who's going to tell the people I care about?"

That's what this book is. It's me telling you.

Here's what I believe after all of it.

Nobody knows exactly what the next five years hold. Not me, not the people building these systems, not the experts on television. Anyone who tells you precisely what 2031 looks like is guessing, and plenty of them are selling something.

But you've now seen what *is* known. The pay slips showing young people quietly locked out of jobs that used to be their way in. The Doubling Clock. The scam figures. The tutoring studies that show what AI can do for a child's learning when it's used well, and the research hinting at what it costs when it isn't. Driverless cars carrying paying passengers. Electricians in demand because of buildings full of computers. The three scenarios, and the warning signs for each. You know more about what's actually going on than most of the people around you.

And more importantly, you've got a plan.

A safe word that could stop someone you love from losing their savings to a voice that sounds exactly like yours. A Task Test that shows you where your job is strong and where it's exposed. A buffer, even a small one, that buys you time. A second income, even a tiny one, so you're never standing on one leg. Skills most people won't have for years. Kids who learn to think

with AI instead of letting it think for them. A family that talks about this together instead of worrying about it alone.

None of that depends on knowing which future turns up. All of it makes you stronger in every one of them.

I've spent a lot of this book on what could go wrong, because you deserve honesty and because the risks are real. But I want to leave you with the other side, because it's just as real.

We may be entering a time when one person with the right skills can do what used to take a team. When a kid in a village with a phone can get tutoring that used to be available only to the rich. When a tradie gets their evenings back. When a nurse spends more time with patients and less with paperwork. When someone with a good idea and no money can build a business from the kitchen table. When a person who lost their job at fifty discovers, to their own surprise, that twenty-five years of experience is worth *more* in a world full of machines that have read everything and lived nothing.

Some of that is already happening. More of it will happen, and mostly for people who are ready.

AI isn't something you need to worship, and it isn't something you need to fear. It's a tool, an extraordinary one, arriving faster than any tool before it. But ignoring it is a choice too. It just happens to be the choice that leaves the decisions to everyone else.

You picked up this book because some part of you sensed that something big was shifting and you didn't want to be caught standing still. Trust that instinct. It's the same one that has kept people safe through every great change in history: pay attention, get ready, look after your people.

You can't control how fast AI moves. You can control how ready you are when it arrives. And by reading this far, you're already more ready than you were.

So learn. Adapt. Protect your family. Build some options. Help your kids become capable, curious, decent people who can think for themselves. And start now, before it becomes obvious.

Then, if this book helped, give it to someone you love.

Todd Lego

Glossary

Agent (AI agent). An AI that doesn't just answer questions but carries out tasks: it takes a goal, breaks it into steps and does them, such as browsing, clicking, filling in forms or writing and running code, often checking its own work. (Chapters 2 and 4.)

Anchor task. A task that needs a human, usually because it's protected by one or more of the Five Moats. (Chapter 6.)

Artificial intelligence (AI). Computer systems that do things we'd normally associate with human intelligence, such as understanding language, recognising images, making predictions and solving problems.

Assist task. A task where AI makes a person faster or better, but the person still leads, checks or finishes it. (Chapter 6.)

Augmentation. AI helping a person do their work better or faster, rather than doing it instead of them.

Automate task. A task AI can do now, or probably soon, with little human input. (Chapter 6.)

Automation. A machine or AI doing a task instead of a person.

Chatbot. An AI you talk to by typing or speaking, which replies conversationally.

CRAFT. A way of briefing AI well: Context, Role, Ask, Format, Test. (Chapter 14.)

Data centre. A large building full of computers that run AI systems and other online services, using a lot of electricity and cooling.

Deepfake. Fake video, audio or images made with AI that realistically show a real person saying or doing something they never said or did.

Doubling Clock. The finding that the length of task AI can complete on its own has been doubling every several months. (Chapter 3.)

Exposure. How much of a job's tasks AI could do or substantially change. High exposure doesn't necessarily mean the job disappears. (Chapter 6.)

Five Moats. Five things that make a task hard for AI to take: Hands, Heart, Trust, Taste and Accountability. (Chapter 6.)

Forecast labels. The labels used in this book to show how certain a claim is: Happening now, Trajectory, Plausible scenario and Wild card.

Frontier model. One of the most advanced AI systems available at a given time.

Generative AI. AI that creates new content, such as text, images, audio, video or code, rather than only analysing existing content.

Hallucination. When an AI produces information that sounds right but is false or made up.

Human Premium. The extra value people place on things done by real humans as machine-made versions become cheap and abundant. (Chapter 21.)

Humanoid robot. A robot built roughly in the shape of a person, designed to work in spaces made for people.

Income Tripod. A framework for household resilience: a main income, a second income and a cash buffer. (Chapter 12.)

Jagged Edge (jagged frontier). The finding that AI is excellent at some tasks and poor at others that seem similar, without warning you which is which. (Chapter 3.)

Large language model (LLM). The type of AI behind most modern chatbots and assistants, trained on vast amounts of text to predict what comes next.

Lazy/Smart Fork. The choice between using AI to think more (the smart path) or to think less (the lazy path). (Chapter 19.)

Liar's dividend. When the existence of deepfakes lets people dismiss *real* evidence as fake.

Multi-factor authentication. A security step that needs a second proof of identity, such as a code from your phone or an authenticator app, as well as your password.

Multimodal. AI that can work with more than one kind of input or output, such as text, images, audio and video.

No-code. Tools that let you build apps or automations without writing programming code.

Prompt. The instruction or question you give an AI.

Proof over paper. The principle that real, visible results and projects increasingly matter more than certificates and résumés. (Chapter 11.)

Robotaxi. A taxi with no human driver.

Safe word (family). A secret word or phrase agreed within a family to confirm identity in an emergency call, protecting against voice-clone scams. (Chapter 16.)

Sextortion. Blackmail using real or fake sexual images, threatening to share them unless the victim pays or complies.

Sycophancy. An AI's tendency to agree with or flatter the user rather than challenge them.

Time horizon. In METR's research, the length of task (measured in skilled-human time) that an AI can complete successfully about half the time. (Chapter 3.)

Two-Source Rule. If it matters, confirm information from an AI with two independent, reliable sources that aren't AI. (Chapter 14.)

Universal basic income. A proposal for governments to give everyone regular, unconditional cash payments.

Vibe coding. Building software by describing what you want to an AI in plain language and letting it write the code.

Voice cloning. Using AI to create a synthetic copy of someone's voice from recordings.

Resources

Tools and websites change quickly, so where possible I've listed organisations rather than products. Make sure any site you visit is the genuine one: type the address yourself or search for the organisation's name.

STAYING INFORMED, CALMLY

Stanford University's annual *AI Index Report* (hai.stanford.edu) is the best single overview of where AI stands. The *International AI Safety Report* (internationalaisafetyreport.org) is an independent scientific review of capabilities and risks. METR (metr.org) publishes independent measurements of what AI systems can do. For AI and jobs, follow the Stanford Digital Economy Lab (digitaleconomy.stanford.edu), the Budget Lab at Yale (budgetlab.yale.edu), the Anthropic Economic Index (anthropic.com/economic-index) and the World Economic Forum's *Future of Jobs Report* (weforum.org). Pew Research Center (pewresearch.org) tracks public attitudes to AI.

LEARNING TO USE AI

The help and learning sections of the major AI providers, including OpenAI, Anthropic, Google and Microsoft, have beginner guides and advice on writing good prompts. Many universities and learning platforms offer free introductory courses on AI; check reviews before you sign up for anything. Your local library, TAFE or community college may run AI workshops for adults and seniors.

SCAMS AND ONLINE SAFETY

In Australia: Scamwatch (scamwatch.gov.au), ReportCyber through the Australian Cyber Security Centre (cyber.gov.au), and IDCARE (idcare.org)

for free help if your identity has been compromised. In New Zealand: Netsafe and CERT NZ. In the United States: the FBI's Internet Crime Complaint Center (ic3.gov) and the Federal Trade Commission (reportfraud.ftc.gov). In the United Kingdom: Action Fraud and the National Cyber Security Centre (ncsc.gov.uk). In Canada: the Canadian Anti-Fraud Centre. Elsewhere, search for your national fraud reporting service and cyber security agency.

CHILDREN AND FAMILIES

In Australia, the eSafety Commissioner (esafety.gov.au) has practical guides for parents on apps, AI companions, image-based abuse and the social media minimum age. Common Sense Media (commonsensemedia.org) publishes research and reviews of apps and AI tools. If a child is being threatened with images, report it to police; in Australia you can also report to the Australian Centre to Counter Child Exploitation (accce.gov.au), and in the United States to the National Center for Missing & Exploited Children, whose Take It Down service can help remove images.

MONEY

ASIC's Moneysmart (moneysmart.gov.au) in Australia, the Consumer Financial Protection Bureau in the United States and MoneyHelper in the United Kingdom offer free, independent financial education. If you're struggling with debt in Australia, the National Debt Helpline (1800 007 007, ndh.org.au) offers free, confidential financial counselling; most countries have a similar service.

MENTAL HEALTH AND CRISIS SUPPORT

If you or someone you know is in crisis, contact emergency services or a crisis line. In Australia: Lifeline, 13 11 14; Kids Helpline, 1800 55 1800; Beyond Blue, 1300 22 4636. In the United States: call or text 988. In the United Kingdom and Ireland: Samaritans, 116 123.

CAREERS AND TRADES

In Australia, Your Career (yourcareer.gov.au) and the Australian Apprenticeships website have information on apprenticeships and career changes. Professional bodies in most fields now publish guidance on AI for their members.

Notes

These notes list the sources for the main factual claims in each chapter. Sources were checked as at October 2026. Where a figure comes from a company or a survey run by a vendor, I've said so. Illustrative scenarios and the author's own frameworks are identified as such.

CHAPTER 1: THE HIRING FREEZE NOBODY ANNOUNCED

Payroll study and the 19 per cent gap. Erik Brynjolfsson, Bharat Chandar and Ruyu Chen, "Canaries in the Coal Mine? Six Facts about the Recent Employment Effects of Artificial Intelligence," Stanford Digital Economy Lab, revised August 2026 (ADP payroll data through June 2026). The authors describe their findings as descriptive rather than causal.

"Probably not (yet)" AI. The Budget Lab at Yale, "AI Is Probably Not (Yet) the Reason for Labor Market Weakening" and the "Evaluating the Impact of AI on the Labor Market" tracker series, 2025–2026.

US unemployment of 4.1 per cent. US Bureau of Labor Statistics, *The Employment Situation: August 2026*.

Recent graduates. Federal Reserve Bank of New York, "The Labor Market for Recent College Graduates," data for the first and second quarters of 2026 (graduates aged 22–27).

AI as the leading stated reason for job cuts. Challenger, Gray & Christmas, monthly Job Cut Announcement Reports, 2026. Through June 2026, AI was cited in 101,743 announced cuts, compared with 54,836 for all of 2025. Reasons are self-reported by employers.

95 per cent of pilots. MIT NANDA, *The GenAI Divide: State of AI in Business 2025*, July 2025. The methodology (52 interviews, 153 survey responses and a review of public deployments) has been widely criticised.

The claims processor. A composite illustration, not a real individual.

CHAPTER 2: WHAT AI ACTUALLY IS

How large language models work. A simplified description consistent with standard technical accounts: pre-training on next-token prediction, followed by further training to make the model helpful and safe.

New York lawyers fined. *Mata v. Avianca, Inc.*, 678 F. Supp. 3d 443 (S.D.N.Y. 2023), opinion and order on sanctions, 22 June 2023.

Melbourne solicitor. *Dayal* [2024] FedCFamC2F 1166 (Federal Circuit and Family Court of Australia, 27 August 2024). The Victorian Legal Services Board later allowed the practitioner to continue practising under supervision.

Robots and household tasks; US and Chinese models. Stanford Institute for Human-Centered Artificial Intelligence, *AI Index Report 2026* (April 2026), and “Inside the AI Index: 12 Takeaways from the 2026 Report.”

Driverless taxis. See the Notes to Chapter 9.

The intern and kitchen analogies. The author’s own illustrations.

CHAPTER 3: THE DOUBLING CLOCK

METR time horizons. Thomas Kwa et al. (METR), “Measuring AI Ability to Complete Long Tasks,” March 2025; METR, “Time Horizon 1.1,” 29 January 2026 (overall doubling of about 196 days, or seven months; about 89 days since 2024; Claude Opus 4.5 at about 320 minutes with a 95 per cent confidence interval of roughly 170 to 729 minutes).

Fourteen and a half hours; sixteen hours; three hours at 80 per cent. METR, “Task-Completion Time Horizons of Frontier AI Models” (updated February and May 2026): Claude Opus 4.6 estimated at about 14.5 hours (95 per cent confidence interval roughly 6 to 98 hours); Claude Mythos Preview at 16 hours or more, with an 80 per cent horizon of about 3 hours 6 minutes. METR notes that measurements above about 16 hours are unreliable with its current task suite.

Seconds in 2019, minutes in 2023, an hour in early 2025. Approximate readings of METR’s published trend (GPT-2 in 2019 at a few seconds; GPT-4 in 2023 at several minutes; Claude 3.7 Sonnet in early 2025 at about an hour).

International AI Safety Report 2026. Yoshua Bengio (chair) et al., published 3 February 2026.

The Jagged Edge. Fabrizio Dell’Acqua, Edward McFowland III, Ethan Mollick, et al., “Navigating the Jagged Technological Frontier,” Harvard Business School Working Paper 24-013 (2023): 758 BCG consultants; 12.2 per cent more tasks, 25.1 per cent faster and about 40 per cent higher quality inside the frontier; 19 percentage points less likely to be correct on the task outside it.

The perception trap. Joel Becker et al. (METR), “Measuring the Impact of Early-2025 AI on Experienced Open-Source Developer Productivity,” July 2025 (arXiv:2507.09089). METR later described the result as historical (February 2026).

Doubling arithmetic. Illustrative only; not a forecast.

CHAPTER 4: 2027–2028: THE AGENT YEARS

The accounting practice. An illustrative scenario, not a documented case.

Young workers and hiring. Brynjolfsson, Chandar and Chen, “Canaries in the Coal Mine?” (Stanford Digital Economy Lab, revised August 2026).

Software developers aged 22–25. Stanford HAI, *AI Index Report 2026*.

FBI figures. Federal Bureau of Investigation, Internet Crime Complaint Center, *2025 IC3 Annual Report* (2026): US\$20.9 billion in reported losses, up 26 per cent; people aged 60 and over reported US\$7.75 billion, up 59 per cent; first dedicated section on AI-related complaints.

Australian scam losses. National Anti-Scam Centre (ACCC), *Targeting Scams: Report of the National Anti-Scam Centre on Scams Activity 2025* (March 2026): A\$2.18 billion in reported losses; older Australians, about 17 per cent of the population, accounted for 26.5 per cent of losses reported to Scamwatch.

EU AI Act. Transparency obligations (Article 50) apply from 2 August 2026. High-risk obligations for most stand-alone systems were deferred to 2 December 2027, and for AI embedded in regulated products to 2 August 2028, by the Digital Omnibus on AI (Regulation (EU) 2026/1744). Council of the European Union press release, 7 May 2026.

California companion chatbot law. Senate Bill 243 (2025), in force 1 January 2026.

Australia's approach. Australian Government, *National AI Plan* (December 2025), which set aside the previously proposed mandatory guardrails in favour of existing technology-neutral laws and an advisory AI Safety Institute. Reported by ABC News, 2 December 2025.

Order of affected industries; skills likely to rise in value. The author's interpretation, drawing on task-exposure research, the Anthropic Economic Index, the Stanford payroll study and the World Economic Forum's *Future of Jobs Report 2025*.

CHAPTER 5: 2029–2031: THREE FUTURES

The three scenarios. The author's framework. Relative likelihoods are the author's judgement, not measured probabilities.

Dario Amodei's warning. Reported by Axios, "Behind the Curtain: A white-collar bloodbath," 28 May 2025. His later emphasis on rising demand: *Fortune*, 5 May 2026.

IMF on inequality. Kristalina Georgieva, "AI Will Transform the Global Economy. Let's Make Sure It Benefits Humanity," IMF Blog, 14 January 2024 (about 40 per cent of global employment exposed to AI, about 60 per cent in advanced economies; AI likely to worsen overall inequality in most scenarios).

Data-centre electricity. International Energy Agency, *Energy and AI* (2025): global data-centre electricity consumption projected to more than double to about 945 terawatt-hours by 2030, slightly more than Japan's total electricity consumption today.

Slow adoption of computers. Often called the "productivity paradox" in economics; a historical framing, not a prediction.

CHAPTER 6: THE TASK TEST

Tasks rather than jobs. The task-based approach is standard in labour economics: David Autor, Frank Levy and Richard Murnane, "The Skill Content of Recent Technological Change," *Quarterly Journal of Economics* (2003); Tyna Eloundou, Sam Manning, Pamela Mishkin and Daniel Rock, "GPTs are GPTs," *Science* (2024).

IMF exposure figures. Georgieva, IMF Blog, 14 January 2024; see Notes to Chapter 5.

Automation and augmentation in actual use. Anthropic, *Anthropic Economic Index* reports (January–June 2026): about 52 per cent of Claude.ai conversations classified as augmentation and about 45 per cent as automation in the late-2025 sample; business use through the API is mostly automation.

Spreadsheets and new demand. A widely cited historical example of technology increasing demand for related work, sometimes described as the Jevons paradox.

Task Test, Five Moats and exposure thresholds. The author's frameworks. The thresholds are rules of thumb, not empirically derived cut-offs.

Worked examples. Illustrative task breakdowns. AI phone and booking systems for medical practices are commercially available as of 2026.

CHAPTER 7: OFFICE, ADMIN, FINANCE, SALES AND SUPPORT

Customer support study. Erik Brynjolfsson, Danielle Li and Lindsey R. Raymond, “Generative AI at Work,” *Quarterly Journal of Economics* 140, no. 2 (2025); originally NBER Working Paper 31161 (2023). 5,179 agents; issues resolved per hour up 14 per cent on average and 34 per cent for novice and low-skilled workers; minimal effect on the most skilled.

Klarna’s 2024 announcement. Klarna press release, “Klarna AI assistant handles two-thirds of customer service chats in its first month,” 27 February 2024.

Klarna’s reversal. Sebastian Siemiatkowski, interviewed by Bloomberg, May 2025; quotations as reported by Bloomberg and subsequently by *Finextra*, *CX Dive* and others.

EU AI Act and recruitment. AI used in recruitment and worker management is listed as high-risk in Annex III of the EU AI Act; obligations deferred to 2 December 2027 by Regulation (EU) 2026/1744.

Role exposure ratings. The author’s assessment, drawing on Eloundou et al. (2024), the Anthropic Economic Index, the Stanford payroll study (which found the largest effects in software development and customer service) and the World Economic Forum’s *Future of Jobs Report 2025*.

CHAPTER 8: HEALTH, EDUCATION, LAW AND THE LICENSED PROFESSIONS

AI scribe randomised trial. Paul J. Lukac et al., “Ambient AI Scribes in Clinical Practice: A Randomized Trial,” *NEJM AI* (2025): 238 outpatient physicians at UCLA Health; one tool reduced time-in-note by about 9.5 per cent versus control, the other showed no significant difference; both groups reported improvements in workload and burnout measures.

Other scribe studies. Real-world evaluations published in 2025–2026, including studies from Mass General Brigham and in *JMIR Medical Informatics*, report modest reductions in documentation time and after-hours work.

Radiology. Geoffrey Hinton’s 2016 remarks were widely reported at the time; his 2025 comments were reported by *The New York Times* (“Your A.I. Radiologist Will Not Be With You Soon,” May 2025). On radiology workforce growth, see Deena Mousa, “AI isn’t replacing radiologists,” *Understanding AI*, 2025, citing record US residency positions in 2025.

Students and teachers. Stanford HAI, *AI Index Report 2026*.

Nigeria tutoring trial. Martín E. De Simone et al., *From Chalkboards to Chatbots: Evaluating the Impact of Generative AI on Learning Outcomes in Nigeria*, World Bank Policy Research Working Paper 11125 (May 2025): six weeks, teacher-guided, effect of 0.31 standard deviations overall, equivalent to roughly 1.5 to 2 years of typical schooling.

Harvard physics tutor. Gregory Kestin, Kelly Miller et al., “AI tutoring outperforms in-class active learning: an RCT introducing a novel research-based design in an authentic educational setting,” *Scientific Reports* (June 2025): 194 students; median learning gains more than double those of the active-learning class.

Lawyers and fictitious cases. See Notes to Chapter 2.

Exposure ratings. The author's assessment using the Task Test and Five Moats.

CHAPTER 9: TRADES, TRANSPORT, FACTORIES AND WAREHOUSES

Humanoid robot deployments. Industry tracking reported by humanoid.guide, "Humanoid deployments in 2026 favor Figure and Agility" (2026). Figure's BMW Spartanburg deployment is as reported by Figure AI.

Tesla Optimus. Tesla fourth-quarter 2025 earnings call, 28 January 2026, as reported by *Electrek* and others.

Robots and household tasks. Stanford HAI, *AI Index Report 2026*.

Waymo. Waymo company statements on weekly paid rides (about 50,000 in May 2024; about 500,000 by early 2026; goal of one million by the end of 2026), as reported by *TechCrunch* (27 March 2026) and *Forbes* (10 December 2025); service launches in Denver, San Diego and Tampa on 1 September 2026 brought the total to fourteen metropolitan areas.

Aurora. Aurora Innovation, second-quarter 2026 shareholder letter (Form 8-K, July 2026): plan to exit 2026 with about 200 driverless trucks; driverless commercial operations began between Dallas and Houston in 2025.

Data-centre electricity. International Energy Agency, *Energy and AI* (2025).

Construction workforce survey. Associated General Contractors of America and NCCER, *2026 Workforce Survey* (conducted July–August 2026, 1,830 respondents): among contractors doing data-centre work, 58 per cent reported increased competition for skilled workers and 49 per cent increased wage pressure.

Exposure ratings and the Australian trucking outlook. The author's assessment.

CHAPTER 10: CREATIVE AND TECH WORK

Software bug-fixing benchmark. SWE-bench Verified leaderboard; Stanford HAI, *AI Index Report 2026* (performance rising from around 60 per cent to well above 90 per cent within roughly a year; the leading reported score in April 2026 was 93.9 per cent).

Young software developers. Stanford HAI, *AI Index Report 2026*.

Freelancers. Xiang Hui, Oren Reshef and Luofeng Zhou, "The Short-Term Effects of Generative Artificial Intelligence on Employment: Evidence from an Online Labor Market," *Organization Science* 35, no. 6 (2024).

Getty Images v Stability AI. *Getty Images (US) Inc & Ors v Stability AI Ltd*, judgment of the High Court of England and Wales, 4 November 2025.

Anthropic settlement. *Bartz v. Anthropic PBC* (N.D. Cal.): US\$1.5 billion class settlement agreed in 2025, covering about 482,000 works, later granted final court approval.

Exposure ratings. The author's assessment.

CHAPTER 11: IF YOU DO X, START LEARNING Y

The Reskill Map. The author's framework, built from the task analysis in Chapters 6–10, the World Economic Forum's *Future of Jobs Report 2025* (which lists AI and big data, networks and cyberse-

curity, technological literacy, creative thinking, resilience and flexibility, and curiosity and lifelong learning among the fastest-growing skills) and industry demand signals such as the data-centre construction boom described in Chapter 9. Job titles are illustrative and availability varies by region.

Proof over paper. The author's principle.

CHAPTER 12: THE INCOME TRIPOD

Emergency fund guidance. A widely used rule of thumb among government financial literacy services, including ASIC's Moneysmart (Australia), the Consumer Financial Protection Bureau (United States) and MoneyHelper (United Kingdom).

National Debt Helpline. A free, independent financial counselling service funded by Australian governments: 1800 007 007, ndh.org.au.

OpenResearch Unconditional Cash Study. OpenResearch, findings released July 2024 (US\$1,000 a month for three years to 1,000 recipients, compared with 2,000 receiving US\$50 a month); Eva Vivalt et al., "The Employment Effects of a Guaranteed Income," NBER Working Paper 32719 (2024).

Investment scams. US Federal Trade Commission, Consumer Sentinel Network data for 2025 (released 2026): investment scams were the largest category of reported fraud losses, at about US\$7.9 billion.

Income Tripod and Vulnerability Audit. The author's frameworks.

CHAPTER 13: EARNING WITH AI: FROM \$0, \$100 AND \$1,000

MIT study. See Notes to Chapter 1.

US adoption of generative AI. Stanford HAI, *AI Index Report 2026*: population-level generative AI adoption of 28.3 per cent in the United States, ranking 24th among countries measured; Singapore highest at 61 per cent.

Earnings. No specific income figures are given deliberately, because rates vary widely by location and skill. The paths are illustrative business models, not guarantees.

CHAPTER 14: HOW TO ACTUALLY USE AI

CRAFT and the Two-Source Rule. The author's frameworks, consistent with prompting guidance published by the major AI developers (provide context, examples and a clear output format).

Telling people what they want to hear. Known as sycophancy. Mrinank Sharma et al., "Towards Understanding Sycophancy in Language Models" (Anthropic, 2023; published at ICLR 2024); OpenAI, "Sycophancy in GPT-4o: What happened and what we're doing about it," 29 April 2025.

AI tutoring design. See Notes to Chapter 8.

Fabricated and misattributed citations. See Notes to Chapter 2.

"Vibe coding." Term coined by Andrej Karpathy in a post on X, 2 February 2025.

Privacy and agent guidance. General good practice. Readers should consult their tools' current privacy policies and settings.

CHAPTER 15: THE 30-DAY AI RESKILL CHALLENGE

Challenge design. The author's programme, built on the techniques in Chapter 14 and the frameworks in Chapters 6, 11, 13, 16 and 24. Its emphasis on active practice, retrieval and teaching others reflects well-established learning research; see, for example, Henry L. Roediger III and Jeffrey D. Karpicke, "Test-Enhanced Learning," *Psychological Science* 17, no. 3 (2006).

CHAPTER 16: SCAMS, VOICES AND FAKES

Arup. Hong Kong Police Force statements, February 2024; Arup's confirmation to CNN, "Arup revealed as victim of \$25 million deepfake scam involving Hong Kong employee," 16 May 2024. Arup's involvement in the Sydney Opera House was as structural engineer.

FBI figures. FBI Internet Crime Complaint Center, *2025 IC3 Annual Report* (2026): US\$20.9 billion total reported losses; US\$7.75 billion reported by people aged 60 and over; cryptocurrency-related losses above US\$11 billion; 22,364 AI-related complaints with losses of about US\$893 million.

FTC figures. US Federal Trade Commission, Consumer Sentinel Network data for 2025 (released 2026): US\$15.9 billion in reported fraud losses; US\$3.5 billion to impostor scams. The estimate of up to US\$195.9 billion in total 2024 losses, adjusted for under-reporting, is from the FTC's report to Congress *Protecting Older Consumers 2024–2025*.

Australian figures. National Anti-Scam Centre, *Targeting Scams: Report of the National Anti-Scam Centre on Scams Activity 2025* (March 2026): A\$2.18 billion in reported losses, up 7.8 per cent; older Australians, about 17.1 per cent of the population, accounted for 26.5 per cent of losses reported to Scamwatch; the top five scam types by loss were investment, payment redirection, romance, phishing and remote access.

Voice-clone survey. McAfee, *The Artificial Imposter* (May 2023): 7,054 respondents in seven countries; about one in four had experienced an AI voice scam or knew someone who had; 77 per cent of those targeted reported losing money. A vendor survey.

Sextortion. Federal Bureau of Investigation, public service announcement I-060523-PSA, "Malicious Actors Manipulating Photos and Videos to Create Explicit Content and Sextortion Schemes," 5 June 2023.

Caller ID spoofing. Warnings about number spoofing are published by the Australian Communications and Media Authority, Scamwatch and the US Federal Communications Commission.

EU labelling rules. EU AI Act, Article 50, applicable from 2 August 2026.

Family Defence Plan. Consistent with guidance from the ACCC's Scamwatch, the Australian Cyber Security Centre, the FTC and the FBI: verify through a separate channel, never pay with gift cards, and use multi-factor authentication.

CHAPTER 17: LOVE, LONELINESS AND MACHINES

Teens and AI companions. Common Sense Media, *Talk, Trust, and Trade-Offs: How and Why Teens Use AI Companions* (July 2025): nationally representative survey of 1,060 US teens aged 13–17, conducted April–May 2025.

Australian children. eSafety Commissioner, “eSafety report shows AI companions are putting children at risk,” media release, March 2026, and *Findings from transparency notices on AI companion apps* (2025–2026).

Character.AI and Google. Court filings unsealed 7 January 2026 disclosing mediated settlements in five cases filed in Florida, Colorado, New York and Texas; reported by CNN, “Character.AI and Google agree to settle lawsuits over teen mental health harms and suicides,” 7 January 2026. Terms are confidential and subject to court finalisation.

Sycophancy. See Notes to Chapter 14.

California SB 243. Signed 13 October 2025; in force 1 January 2026.

Australian codes. eSafety Commissioner, *Age-Restricted Material Codes* fact sheet (March 2026); codes in effect from 9 March 2026 under the *Online Safety Act 2021*.

Crisis lines. Lifeline Australia (13 11 14); Kids Helpline (1800 55 1800); 988 Suicide & Crisis Lifeline (United States); Samaritans (116 123, United Kingdom and Ireland). Numbers current as at October 2026.

Future of AI companions. The author’s assessment, labelled as trajectory, plausible scenario and wild card.

CHAPTER 18: THE FAMILY AI PREPAREDNESS PLAN

Plan content. Compiled from Chapters 6, 11, 12, 13, 15, 16, 17 and 19; see those chapters’ Notes for sources.

Security recommendations. Consistent with guidance from national cyber security agencies, including the Australian Cyber Security Centre, the US Cybersecurity and Infrastructure Security Agency and the UK National Cyber Security Centre.

CHAPTER 19: WHAT OUR KIDS SHOULD LEARN NOW

Nigeria and Harvard studies. See Notes to Chapter 8.

MIT essay study. Nataliya Kosmyna et al., “Your Brain on ChatGPT: Accumulation of Cognitive Debt when Using an AI Assistant for Essay Writing Task,” MIT Media Lab, arXiv preprint (June 2025); 54 participants. A small, not yet peer-reviewed study.

Knowledge-worker survey. Hao-Ping (Hank) Lee et al., “The Impact of Generative AI on Critical Thinking: Self-Reported Reductions in Cognitive Effort and Confidence Effects From a Survey of Knowledge Workers,” *Proceedings of CHI 2025* (Microsoft Research and Carnegie Mellon University); 319 participants.

Students and teachers. Stanford HAI, *AI Index Report 2026*.

Skills growing in importance. World Economic Forum, *Future of Jobs Report 2025* (January 2025), based on a survey of more than 1,000 large employers. These are employer expectations, not measurements.

Minimum ages for AI assistants. As stated in the terms of service of major providers as at 2026; for example, OpenAI requires users to be at least 13 (with parental permission under 18) and Anthropic's consumer Claude service requires users to be 18 or over.

Australia's social media minimum age. *Online Safety Amendment (Social Media Minimum Age) Act 2024*; obligations on platforms from 10 December 2025 (eSafety Commissioner).

Sextortion. See Notes to Chapter 16.

Young workers and recent graduates. See Notes to Chapter 1.

Value of a degree. On graduate earnings and underemployment, see Federal Reserve Bank of New York, "The Labor Market for Recent College Graduates" (2026). The trajectory statement is the author's assessment.

The Lazy/Smart Fork and the durable skills list. The author's frameworks.

CHAPTER 20: THE RISKS, STRAIGHT

Liar's dividend. Robert Chesney and Danielle Keats Citron, "Deep Fakes: A Looming Challenge for Privacy, Democracy, and National Security," *California Law Review* 107 (2019).

IMF. Georgieva, IMF Blog, 14 January 2024; see Notes to Chapter 5.

AI investment, frontier models and transparency. Stanford HAI, *AI Index Report 2026* and "Inside the AI Index: 12 Takeaways from the 2026 Report": global corporate AI investment of US\$581.7 billion in 2025, with US investment of US\$285.9 billion; industry produced more than 90 per cent of notable frontier models; average Foundation Model Transparency Index scores fell from 58 to 40.

Cybersecurity benchmark. Stanford HAI, *AI Index Report 2026* (Cybench unguided solve rate rising from 15 per cent in 2024 to 93 per cent in 2025).

International AI Safety Report 2026. Yoshua Bengio (chair) et al., 3 February 2026.

Survey of AI researchers. Katja Grace et al., "Thousands of AI Authors on the Future of AI," arXiv: 2401.02843 (2024): 2,778 respondents surveyed in late 2023; median estimate of 5 per cent for extremely bad outcomes such as human extinction.

Overdependence research. See the MIT and Microsoft/Carnegie Mellon studies in the Notes to Chapter 19.

Cyber security as a growing skill. World Economic Forum, *Future of Jobs Report 2025*.

Risk ranking. The author's judgement.

CHAPTER 21: THE HUMAN PREMIUM

Klarna. See Notes to Chapter 7. Quotation from Sebastian Siemiatkowski's May 2025 interview with Bloomberg.

Handmade furniture, live music and bakeries. General observations offered as illustrations rather than measured trends.

The Human Premium. The author's concept. The trajectory is the author's assessment; its size is uncertain.

CHAPTER 22: DIRECT IT

*"Don't out-*AI* the *AI*" and the Director's five jobs.* The author's frameworks.

The jagged frontier study. Dell'Acqua et al. (2023); see Notes to Chapter 3.

What one person can do. Illustrations based on documented capabilities of AI tools in 2026 (writing, coding, analysis and agents), not guarantees of outcomes.

Managing agents. Consistent with guidance published by AI developers on keeping humans in charge of agentic systems.

CHAPTER 23: YOUR PLAN

Roadmaps. Built from the frameworks and evidence in Chapters 1–22; see those chapters' Notes for sources. Buffer targets are general educational guidance (see the disclaimer in Chapter 12).

CHAPTER 24: THE AI RESILIENCE SCORECARD AND A LAST WORD

The Scorecard. The author's framework, drawing on Chapters 6, 11, 12, 13, 14, 15, 16, 18 and 19.

The Last Word. Summarises evidence presented in earlier chapters; see those chapters' Notes. Forward-looking statements are the author's assessment.

About the Author

Todd Lego is an Australian writer who came to artificial intelligence the way most people will: as an ordinary person trying to work out what on earth was going on.

About two years before this book was published, Todd knew almost nothing about AI. He'd seen the headlines and heard the arguments, and like plenty of people he wasn't sure whether to believe the hype or ignore it. But something in his bones told him this wasn't just another gadget that would come and go. Whether people liked it or not, he believed AI was going to become central to how we work, earn, learn, run businesses and live our lives.

As he puts it: "I could bury my head in the sand and hope it was somebody else's problem, or I could learn. I chose to learn."

Todd isn't a computer scientist, a programmer by training or an academic, and he'd be the first to tell you so. That's the point. Over the past two years he has put a huge amount of his time into understanding AI and the computers it runs on: reading the research, using the tools every day, experimenting with what they can and can't do, building things with them, making plenty of mistakes, and slowly getting a feel for how quickly these technologies are becoming more powerful. He came at it not from a laboratory but from the kitchen table, asking the questions an ordinary person asks. What can this actually do? Where does it go wrong? What does it mean for my job, my money and my family?

The deeper he went, the more one thing troubled him. The people building AI talk about it constantly. The people investing in it talk about it constantly. But millions of ordinary people, workers, parents, small business owners, kids leaving school and older people living on their own, are either unaware of how fast it's moving, overwhelmed by it, frightened of it, dismissing it, or simply don't know where to start.

Todd believes those people deserve better than hype on one side and doom on the other. They deserve a plain-English explanation of what's really happening, an honest account of what isn't known, and practical steps they can take this week. The more he learned, the more he felt an obligation to pass it on, so that other families could prepare themselves now instead of waking up in a few years wishing they'd started earlier.

Future-Proof Your Family is the result: the book Todd wishes ordinary families already had. It doesn't claim to predict the future. It tries to do something more useful, which is to help readers understand the change that's under way, protect what matters to them, and put themselves and the people they love in the best possible position for whatever comes next.

Todd doesn't present himself as the world's leading AI expert. His view is simpler than that. He saw something important happening, decided not to look away, spent two years learning everything he reasonably could, and then wrote it down as clearly and honestly as he knew how.

He's still learning. If this book gets you started, he'd say that's the best thing it could do.